



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

-  Topic 01: Azure Landing Zone
-  Topic 02: Azure Compute
-  Topic 03: Azure Storage
-  Topic 04: Azure Network & Security
-  Topic 05: Azure Database offering

© Copyright Kodakoid

## Agenda

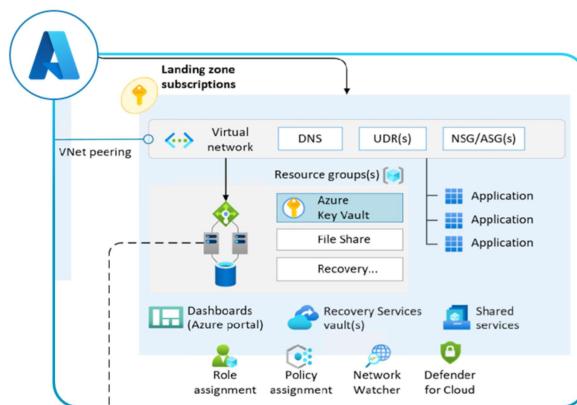
# Azure Landing Zone

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Have you ever written a program but didn't have an application server to run it on?  
Maybe you're a developer and you've written code that executes a basic task.

## Azure Landing Zone



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure landing zone is a Cloud management system that makes deployment and configuration of cloud applications simple, repeatable, consistent and compliant to your organizational standards and requirements.

No two organizations are same and hence their implementation of landing zones will be unique to meet their requirements. As a Kubernetes engineer or Architect, you might not be required to design the Landing Zone but it will be valuable to know these concepts in general and how your organization has implemented this in particular.

## ➤ Azure Landing Zone



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

One of the common analogies for understanding this scaffolding concept, is to compare it to foundational services upon which your suburb or city depends on.



## Azure Landing Zone



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Before a new housing complex is built and people move into them essentials services must be deployed. Similar to networking on Azure, a well planed road infrastructure is necessary for people to move around. Similarly, water and sewer facilities need to be in place. power and gas lines must be available and always-on along with the ability to charge back.

➤ Azure Landing Zone



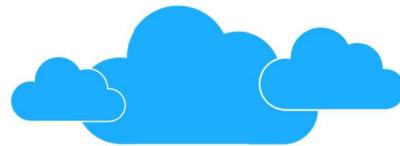
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You can think of Azure landing zone as the essential services required for the different types of applications you will host on Azure. Different applications like AKS will require different services, and policies, and therefore will live in different landing zones.

Reference: <https://github.com/Azure/AKS-Landing-Zone-Accelerator>

## Azure Landing Zone

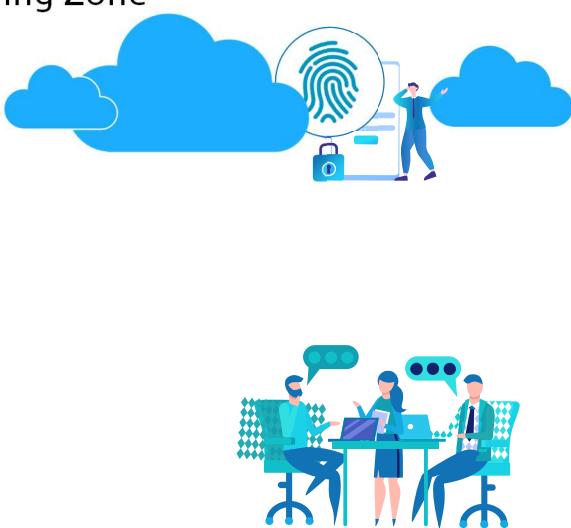


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

The primary purpose of the landing zone is to ensure when your application lands on Azure the required plumbing is in place providing greater agility and compliance with enterprise security and governance requirements.

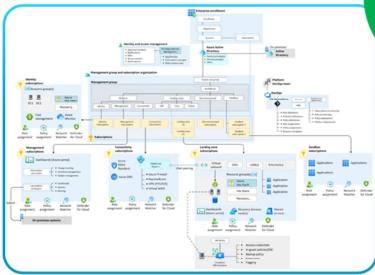
► Azure Landing Zone



Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

© Copyright KodeKloud

## Azure Landing Zone



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure landing zone is a Cloud management system that makes deployment and configuration of cloud applications simple, repeatable, consistent and compliant to your organizational standards and requirements.

No two organizations are same and hence their implementation of landing zones will be unique to meet their requirements. As a Kubernetes engineer or Architect, you might not be required to design the Landing Zone but it will be valuable to know these concepts in general and how your organization has implemented this in particular.

► Azure Landing Zone

The slide features three circular icons. The first icon shows a road network with a car and a person working on it. The second icon shows a worker carrying two large blue water jugs. The third icon shows a fuel pump with a person at the nozzle. In the center, there is a larger illustration of two people shaking hands in front of a house with a 'SALE' sign, symbolizing the completion of a project or the start of a new one.

© Copyright KodeKloud

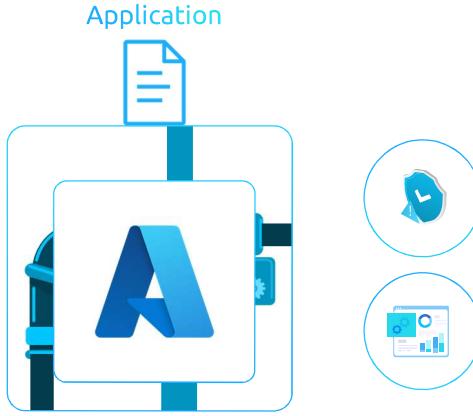
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

One of the common analogies for understanding this scaffolding concept, is to compare it to foundational services upon which your suburb or city depends on. Before a new housing complex is built and people move into them essentials services must be deployed. Similar to networking on Azure, a well planned road infrastructure is necessary for people to move around. Similarly, water and sewer facilities need to be in place. power and gas lines must be available and always-on along with the ability to charge back. You can think of Azure landing zone as the essential services required for the different types of applications you will host on Azure.

Different applications like AKS will require different services, and policies, and therefore will live in different landing zones.

Reference: <https://github.com/Azure/AKS-Landing-Zone-Accelerator>

## Azure Landing Zone



The diagram illustrates the Azure Landing Zone architecture. At the center is a hub containing a large blue letter 'A' and a smaller document icon above it. Above the hub is the word 'Application'. To the right of the hub are two circular icons: one showing a clock and another showing a bar chart.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

The primary purpose of the landing zone is to ensure when your application lands on Azure the required plumbing is in place providing greater agility and compliance with enterprise security and governance requirements.

## Azure Landing Zone

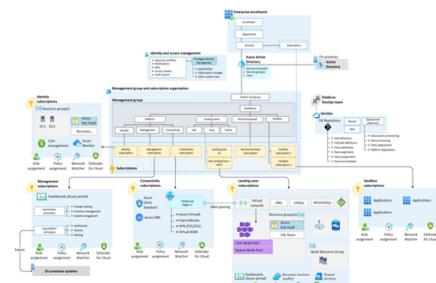
ENVIRONMENT	COMPLIANCE
 Azure billing & Active Directory tenant	 Security
 Identity & access management	 Governance
 Resource organization	 Management
 Network topology & connectivity	 Platform automation & DevOps

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

While an in-depth discussion of Landing zone is out of scope for this course, at high-level. These 8 areas needs to be covered for a well rounded landing Zone

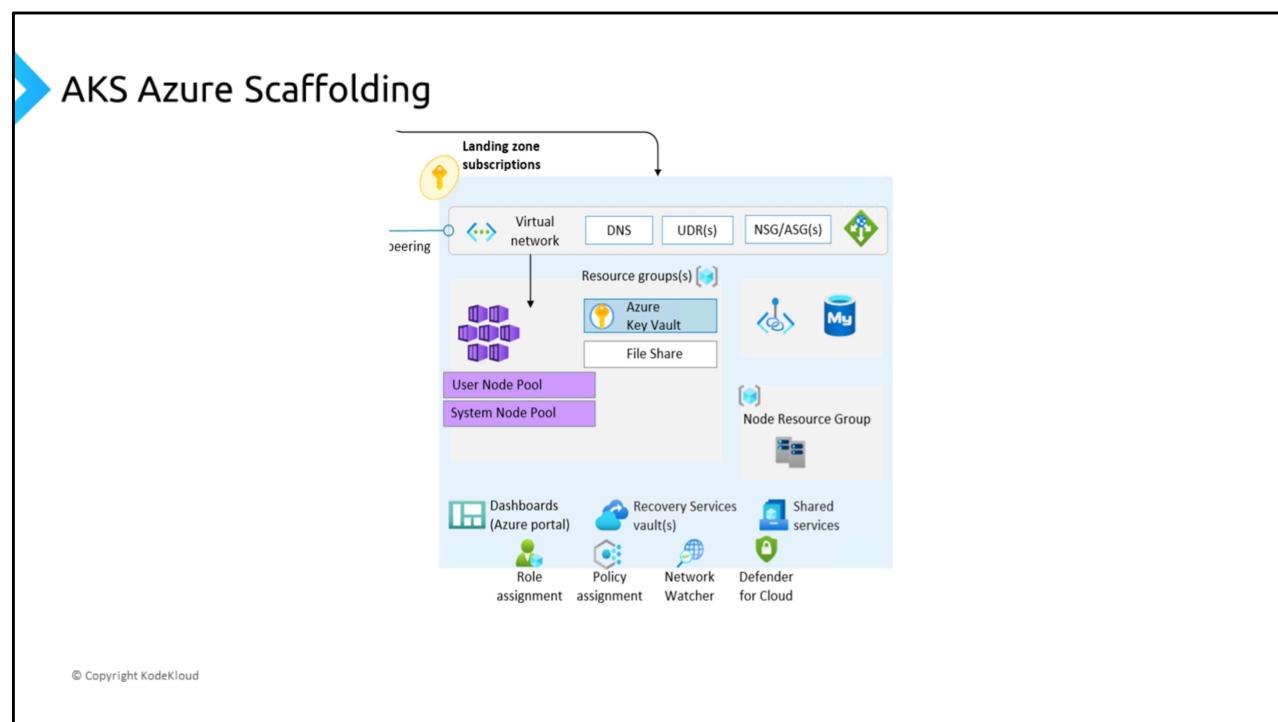
## ➤ AKS Azure Scaffolding



© Copyright KodeKloud

Now that you understand the basic concept of a Landing zone, when you need to deploy an AKS based application, you will have a AKS landing zone similar to the highlighted area here. Your application can take advantage of all the foundational services deployed and can connect to other Azure services in a predictable manner.

Reference: <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/scenarios/app-platform/aks/landing-zone-accelerator>



Reference: <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/scenarios/app-platform/aks/landing-zone-accelerator>

# Azure Compute

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Have you ever written a program but didn't have an application server to run it on?  
Maybe you're a developer and you've written code that executes a basic task.



## Azure Compute

Virtual Machines

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure offers Virtual Machine types and sizes that can address almost every type of workload

Azure Compute

Virtual Machines

Virtual Machines

Virtual Machines

Virtual Machines

Virtual Machines

© Copyright Kodekloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

For example, dev/test workloads, burstable, mission-critical production workloads, small scale to large systems and off course almost all of these can be used for AKS clusters. I have listed a few machine types per category, you can check this URL for full list of VMs and their specs.

Azure Compute

**Virtual Machines**

- Development Workload
- Test Workload
- Burstable
- Small Scale Systems
- Mission-critical Production
- Large Systems

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure offers Virtual Machine types and sizes that can address almost every type of workload – For example, dev/test workloads, burstable, mission-critical production workloads, small scale to large systems and off course almost all of these can be used for AKS clusters. I have listed a few machine types per category, you can check this URL for full list of VMs and their specs.



## Azure Compute

<https://learn.microsoft.com/en-us/azure/virtual-machines/vm-naming-conventions>

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You should also consider getting yourself familiar with Azure Compute naming convention. You will come across names like: Standard\_E64-32ads\_v5  
Every character has a meaning here, you can refer to the link of the screen to demystify the naming convention: <https://learn.microsoft.com/en-us/azure/virtual-machines/vm-naming-conventions>

## Azure Compute

### Virtual Machines

The diagram illustrates four categories of Azure Virtual Machines, each represented by a dashed circle containing a specific icon:

- Compute intensive:** Represented by a dashed circle containing binary code: 010101, 101010, 010101.
- Memory optimized:** Represented by a dashed circle containing a microchip or memory chip icon.
- Storage Optimized:** Represented by a dashed circle containing a cylinder icon.
- GPU Accelerated:** Represented by a dashed circle containing a GPU icon with a gear symbol.

Below each category, the corresponding text label is displayed:

- Compute intensive
- Memory optimized
- Storage Optimized
- GPU Accelerated

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

If you are new in Azure, you would need to understand what these series and naming conventions are. There are VM series optimized for compute, memory, disk, GPUs,

Azure Compute

## Virtual Machines

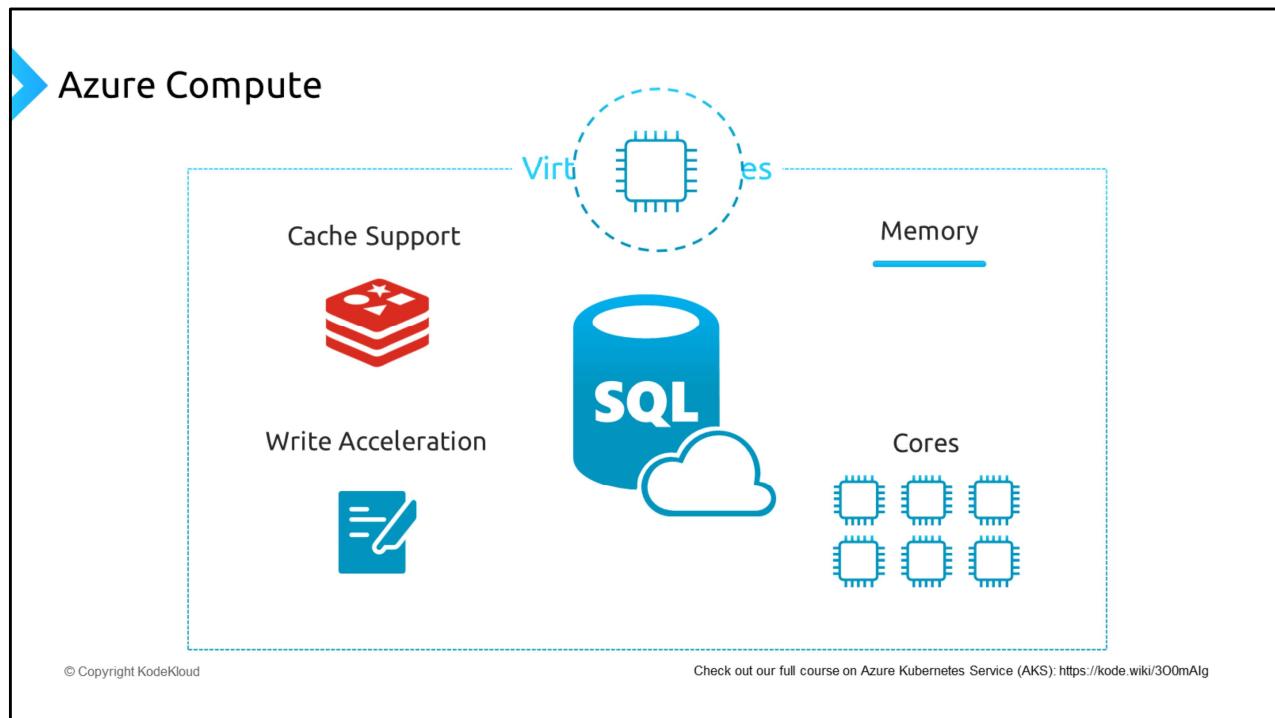
The diagram illustrates four categories of Azure Virtual Machines, each represented by a blue dashed circle containing a specific icon:

- Burstable**: Represented by a bar chart icon.
- Entry Level**: Represented by a monitor icon.
- High-performance Computing**: Represented by a lightning bolt icon.
- General Purpose and Confidential Computing**: Represented by a server rack icon.

© Copyright KodeKloud

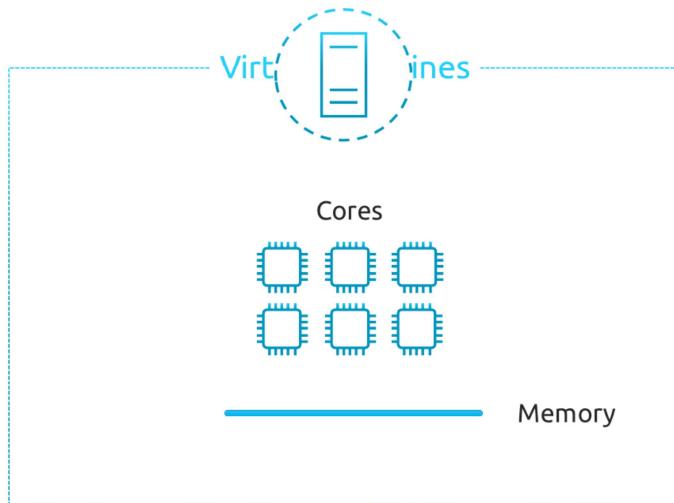
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

and more and as you could imagine this portfolio is growing constantly.



Just to give you a flavor, let me mention couple of these machine types and what criteria you can use to make your selection. Let us say you want to host a critical SQL database on a VM. In general these databases require a lot of memory but comparatively less number of CPU Cores. Then Memory optimized VM sizes should be considered because they offers large amounts of memory, high memory-to-core ratios (which means large amount of memory per core), premium disk and cache support, and write acceleration capabilities.

## Azure Compute



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

On the other hand, General purpose VM sizes provide balanced, but lower memory-to-core ratio. This might be good for smaller database servers.

Azure Compute

The diagram illustrates Azure Compute resources. At the top, a blue arrow points right towards a dashed blue rectangle labeled "Azure Compute". Inside this rectangle, five icons represent different compute services: a virtual machine icon, a storage icon, a function icon, a container instance icon, and another storage icon. Above the rectangle, the word "Virtual" is written in blue, with a dashed circle highlighting the "V" and "u".

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Now let us say, you need high disk throughput and IO for Big Data, ETL,

**Operational Data Store (ODS)**, and  
data warehousing processing, then you should consider storage optimized  
series. This series can offer very high throughput & low latency storage with directly  
mapped local NVMe storage.

➤ Azure Compute

The diagram illustrates the Azure Compute ecosystem. At the top is the Azure logo (a blue 'A' inside a square). A vertical line descends from it to a horizontal line that branches into four segments. Each segment ends with a logo: Intel (blue oval), Ampere Altra (red 'A' with a green checkmark circle), Linux Tux (the penguin mascot), and AMD (black and green logo).

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

While Azure has both Intel and AMD based machines, they have recently launched Ampere® Altra® Arm-based processors. This would be an excellent choice for Linux based containers as they have the best performance to cost ratio and have the lowest environmental footprint.



## Azure Compute

# Standard\_E64-32ads\_v5

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You should also consider getting yourself familiar with Azure Compute naming convention. You will come across names like: Standard\_E64-32ads\_v5  
Every character has a meaning here, you can refer to the link of the screen to demystify the naming convention: <https://learn.microsoft.com/en-us/azure/virtual-machines/vm-naming-conventions>



## Azure Compute

<https://learn.microsoft.com/en-us/azure/virtual-machines/vm-naming-conventions>

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You should also consider getting yourself familiar with Azure Compute naming convention. You will come across names like: Standard\_E64-32ads\_v5  
Every character has a meaning here, you can refer to the link of the screen to demystify the naming convention: <https://learn.microsoft.com/en-us/azure/virtual-machines/vm-naming-conventions>



Next Lesson

## Azure Storage

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Now lets move on to storage.

# Azure Storage

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Have you ever written a program but didn't have an application server to run it on?  
Maybe you're a developer and you've written code that executes a basic task.

## Azure Storage Categories

The diagram illustrates the three main categories of Azure Storage:

- Block Storage:** Represented by a blue square icon with four smaller squares inside, labeled "Azure Disk Storage".
- Object Storage:** Represented by a blue icon of a stack of boxes, labeled "Azure Blob and Data Lake Storage".
- File Storage:** Represented by a blue folder icon, labeled "Azure Files" and "Azure NetApp Files".

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

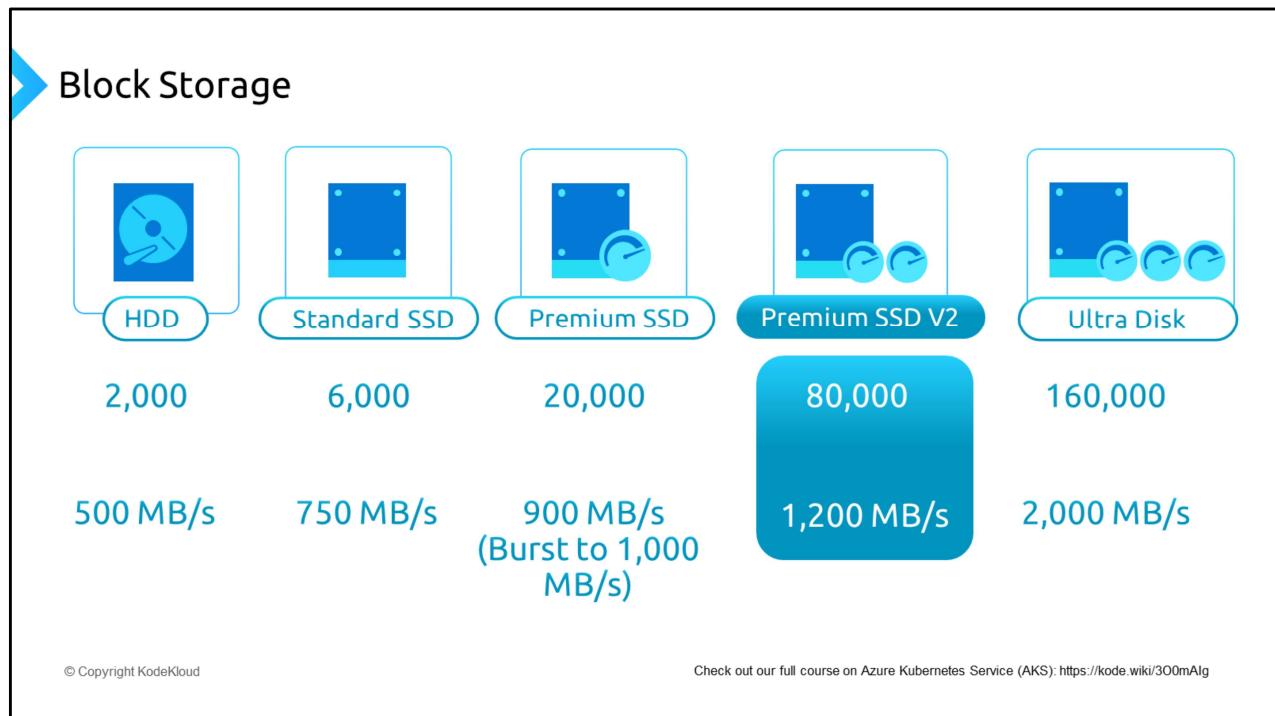
**The three most common** Azure Storage platform services are Block storage which includes capabilities like disks, Object Storage with Blob and Data lake and File storage with capability to deploy SMB and NFS fileshares.

Block Storage				
	HDD	Standard SSD	Premium SSD	Ultra Disk
IOPS	2,000	6,000	20,000	160,000
Throughput	500 MB/s	750 MB/s	900 MB/s (Burst to 1,000 MB/s)	2,000 MB/s

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You can use block storage options also called Azure Disks to create a Kubernetes DataDisk resource. Depending on your IOPS and throughput requirements you can select from a number of SKUs. As you can imagine cost profile of these disks increases from left to right as you move from Standard spinning HDDs to NVMe based Ultra disks.



Microsoft has recently added another offering called Premium SSD V2 which provides performance specs between Premium SSD and Ultra disks. At the time of this recording, they do not provide the ability to use this disk as OS disk.



Block Storage

Mountable Node

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

While we will talk more about AKS and Storage in a later section, it is worth noting that because Azure Disks are mounted as ReadWriteOnce, they're only available to a single node. For storage volumes that can be accessed by pods on multiple nodes simultaneously, use Azure Files, which I will describe in the next slide.

The diagram is titled "File Storage" with a blue arrow icon. It compares two storage SKUs: "Standard" and "Premium".  
**Standard SKU:**  
- Speed: 300 MBps  
- Latency: 10s of ms latency  
- Billing: Pay for storage used  
**Premium SKU:**  
- Speed: 10 GBps  
- Latency: Sensitively low latency  
- Billing: Pay for entire storage

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

File storage service in azure provides serverless file shares functionality with the option to choose either standard SKU with spinning disks or Premium SKU with SSD disks.

As you can see standard files are slower with higher latency while premium files are extremely fast.

The diagram illustrates two Azure File Storage SKUs: Standard and Premium. The Standard SKU is shown with a blue rounded rectangle containing the text "Standard", "300 MBps", and "10s of ms latency". The Premium SKU is shown with a blue rounded rectangle containing the text "Premium", "10 GBps", "Sensitively low latency", and "Pay for storage used". To the right, a blue speech bubble contains "100 GB", a horizontal bar spans from "100 GB" to "1 TB", and a yellow box contains a dollar sign "\$", indicating the cost per TB.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

As you start to build your AKS based solutions on Azure, you will need to know which option to pick and why. Cost implications are always critical in such a selection. While I have documented some of the common use cases for these file storage types, it is worth mentioning that while in standard SKU you pay for storage used in a Premium file type, you pay for entire storage provisioned. For example, if you have provisioned a 1 TB of storage used only 100 Gig, in case of standard storage you pay for only 100 gig while in case of premium files you pay for entire 1 TB.

If your performance needs are higher than what Premium files has to offer, you can explore Netapp Files. This is a first party solution in Azure but managed by their vendor i.e Netapp.

In AKS world, if multiple pods need concurrent access to the same storage volume, you use Azure Files to connect using the [Server Message Block \(SMB\) protocol](#).

## File Storage



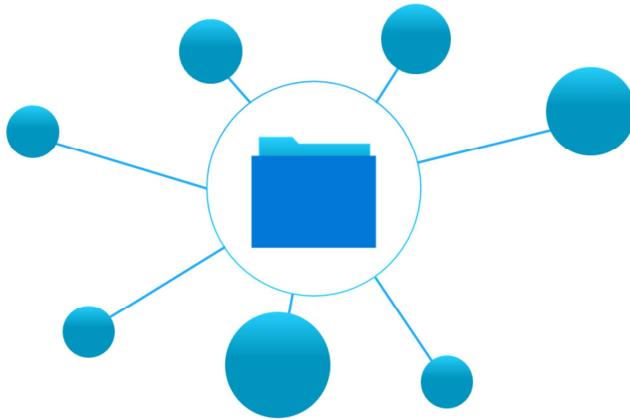
**NetApp™**

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

If your performance need are higher than what Premium files has to offer, you can explore Netapp Files. This is a first party solution in Azure but managed by their vendor i.e Netapp.

## File Storage



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

In AKS world, if multiple pods need concurrent access to the same storage volume, you we use Azure Files to connect using the [Server Message Block \(SMB\) protocol.](#)



## Object Storage

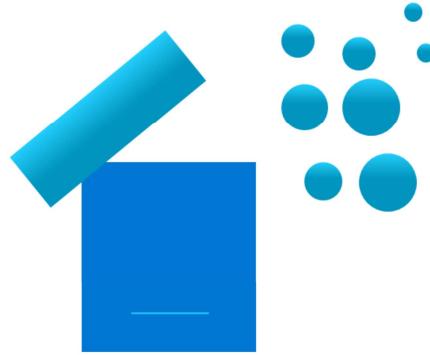


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure object storage is also called Blob storage and is optimized for storing massive amounts of unstructured data. Similar to file storage, azure blobs also have the ability for you to select between Standard or Premium performance tiers.

## Object Storage



© Copyright KodeKloud

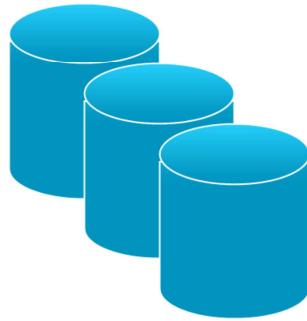
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

If this is the first time you are getting exposed to all the storage options, it can be intimidating. Trust me it get a little better with time

To make matters a little bit more confusing, to have the ability to enable hierarchical namespace which is also called ADLSCGen2. It enables a set of capabilities that are dedicated to big data analytics.

## Object Storage Durability

### Local Redundant



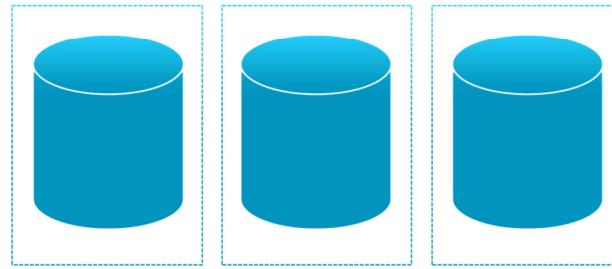
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

You also have the choice to select one of the four Durability or Redundancy options of Local Redundant, Zone Redundant, Read-Access Geo-redundant or Geo-Zone redundant.

## Object Storage Durability

### Zone Redundant



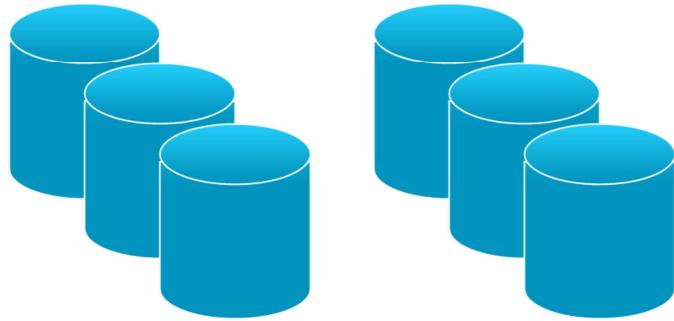
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Without making it complicated, as you move from left to right, Azure can make multiple copies of your data in Availability zones, regions or both.

➤ Object Storage Durability

## Read-access Geo-redundant

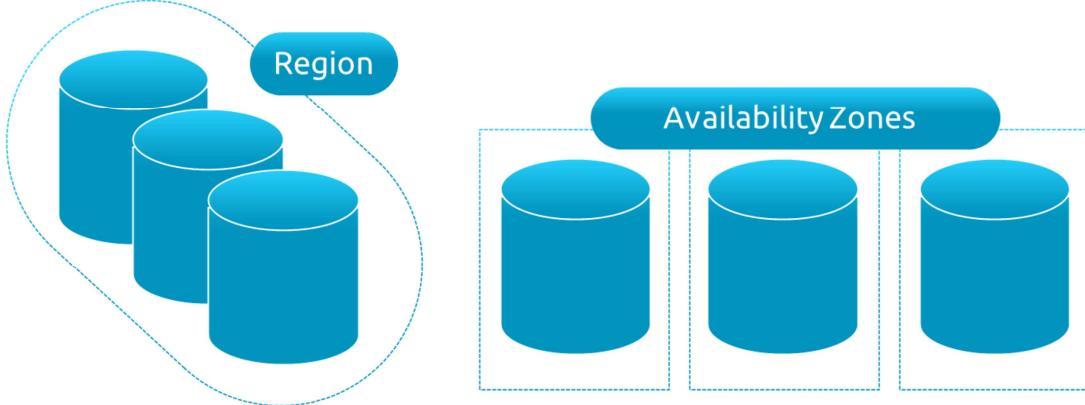


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Object Storage Durability

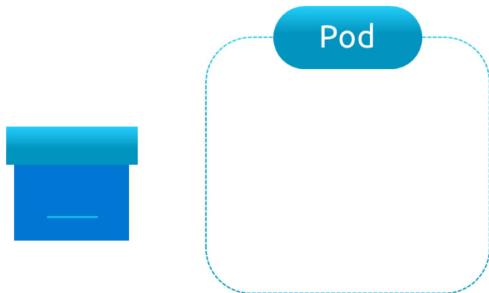
### Geo-zone Redundant



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Object Storage

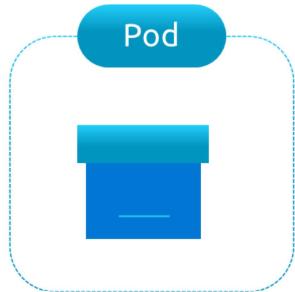


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

In terms of how you can use Blob Storage in AKS – you can mount Blob storage (or object storage) as a file system into a container or pod.

## Object Storage

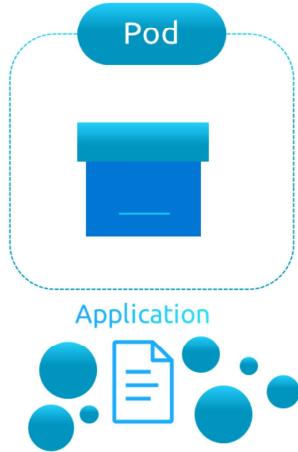


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Using Blob storage enables your cluster to support applications that work with large unstructured datasets like log file data, images or documents

## Object Storage

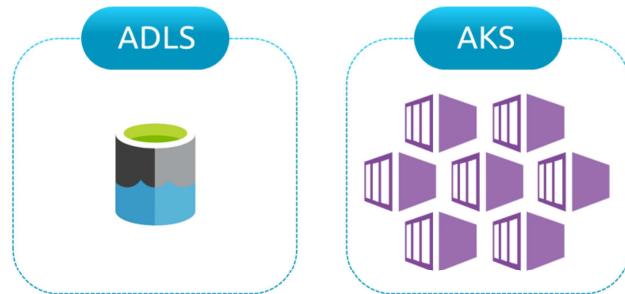


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Additionally, if you ingest data into Azure Data Lake storage i.e ADLS Gen2, you can directly mount and use it in AKS without configuring another interim filesystem.

## Object Storage



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

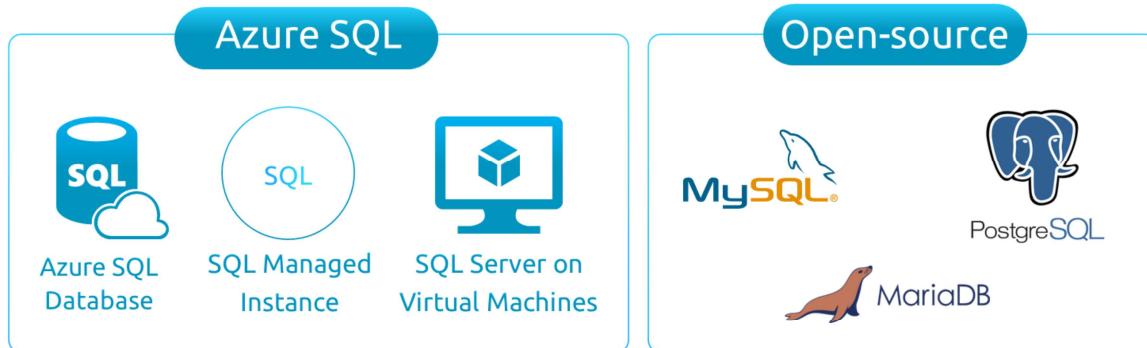
# Azure Database Offerings

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Have you ever written a program but didn't have an application server to run it on?  
Maybe you're a developer and you've written code that executes a basic task.

## Azure Database Offerings



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Most of your Kubernetes applications will interact with one or more databases. Azure has a rich of data offerings both in relational and non-relational databases. I have tabulated some of the common use case for these database offerings.  
Some of the most common Database offering used on Azure are:

## Non-relational Database Offerings

The logo features a stylized Earth with blue continents and green oceans, surrounded by a blue elliptical ring. Four blue rounded rectangular boxes are positioned around the Earth: 'JSON' at the top left, 'Graphs' at the bottom left, 'Column Families' at the top right, and 'Key-value Pairs' at the bottom right. Below the Earth, the text 'Cosmos DB' is written in a blue, sans-serif font.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

- 1.) Azure SQL: There are 3 way to deploy SQL databases based on the Microsoft SQL Server database engine – Azure SQL Database which is a PaaS solution, SQL Managed instance which is a hosted instance of SQL server and a traditional SQL on VM option
- 2.) There are then Azure Database for open-source relational databases like MySQL, MariaDB and PostgreSQL
- 3.) When it comes to non-relational databases, Azure's no #1 option is Cosmos DB which is a global-scale *NoSQL* database system that supports multiple application programming interfaces (APIs), enabling you to store and manage data as JSON documents, key-value pairs, column-families, and graphs.

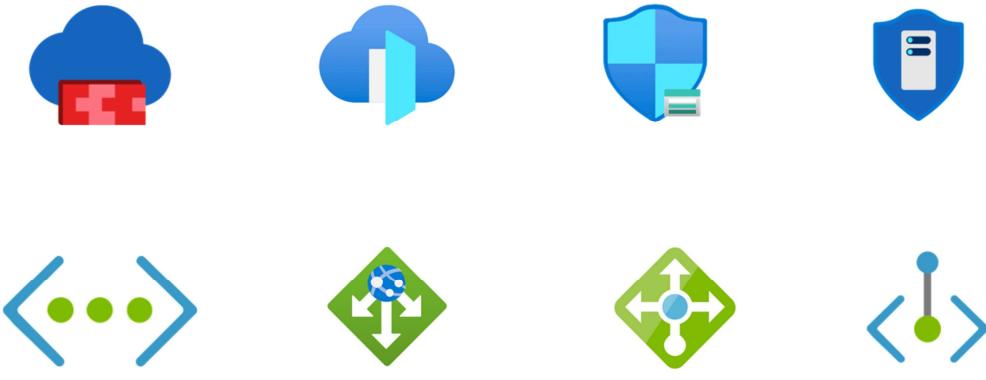
# Azure Network and Security

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Have you ever written a program but didn't have an application server to run it on?  
Maybe you're a developer and you've written code that executes a basic task.

## Azure Networking and Security



The image displays a grid of eight Azure service icons, each representing a different aspect of networking and security:

- Top-left: A blue cloud icon containing a red puzzle piece.
- Top-middle: A blue cloud icon containing a white bar chart.
- Top-right: A blue shield icon containing a white server tower.
- Bottom-left: A blue double-angle bracket icon containing three green dots.
- Bottom-middle-left: A green diamond icon containing a white globe with arrows pointing in various directions.
- Bottom-middle-right: A green diamond icon containing a white central node with arrows pointing outwards.
- Bottom-right: A blue double-angle bracket icon containing a single green dot.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure has a strong and extensive set of services when it comes to networking and security.

## Azure Networking and Security



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

It also provides you with an extensive list of 3<sup>rd</sup> party services so you can get additional functionality if missing in native products or extend your existing products from on-premises or another cloud. Let us cover some basic connectivity and application delivery services.

## Azure Virtual Network

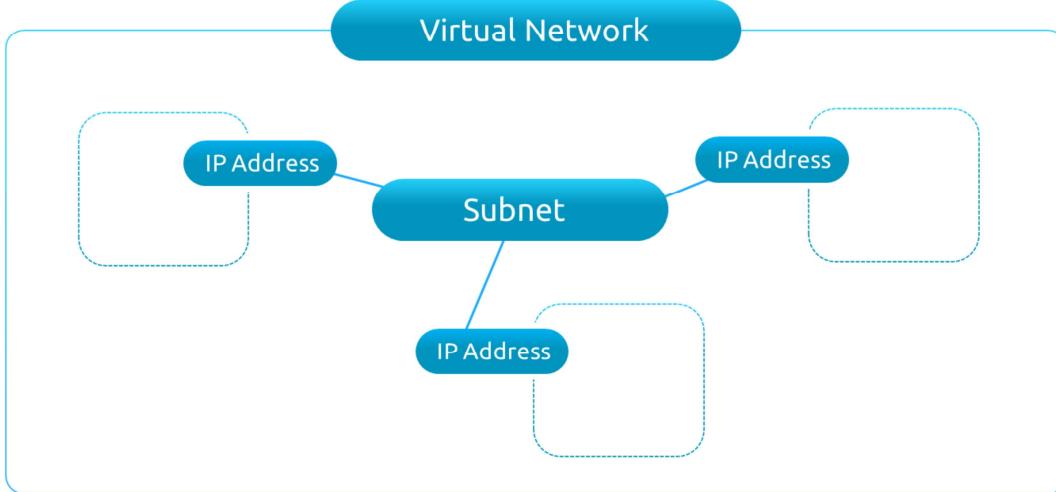
The diagram illustrates the relationship between a central service and its associated IP addresses. A large blue square containing the letter 'A' is connected by lines to three separate dashed boxes, each containing a blue rounded rectangle labeled 'IP Address'. This visualizes how a single service (A) is assigned multiple IP addresses.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Most services on Azure will require an IP address to function and an Azure Virtual Network or VNet is that fundamental building block in Azure that IP addressing depends on.

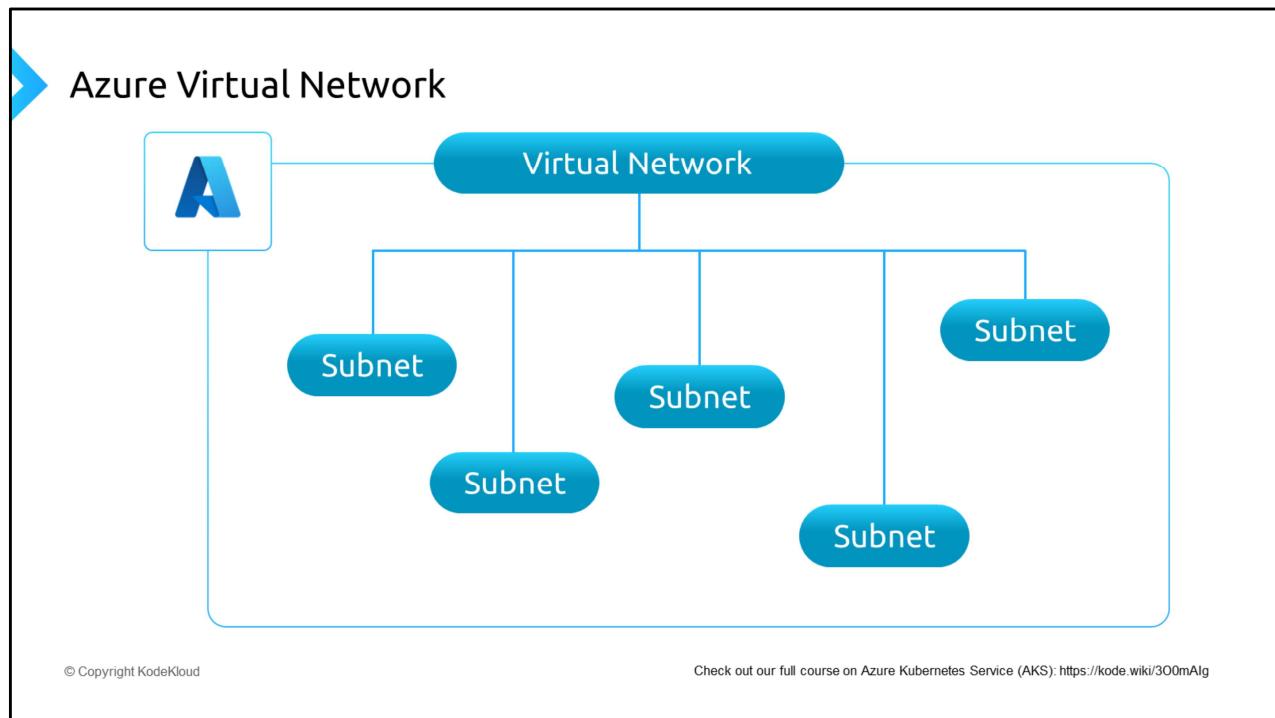
## Azure Virtual Network



© Copyright Kodekloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

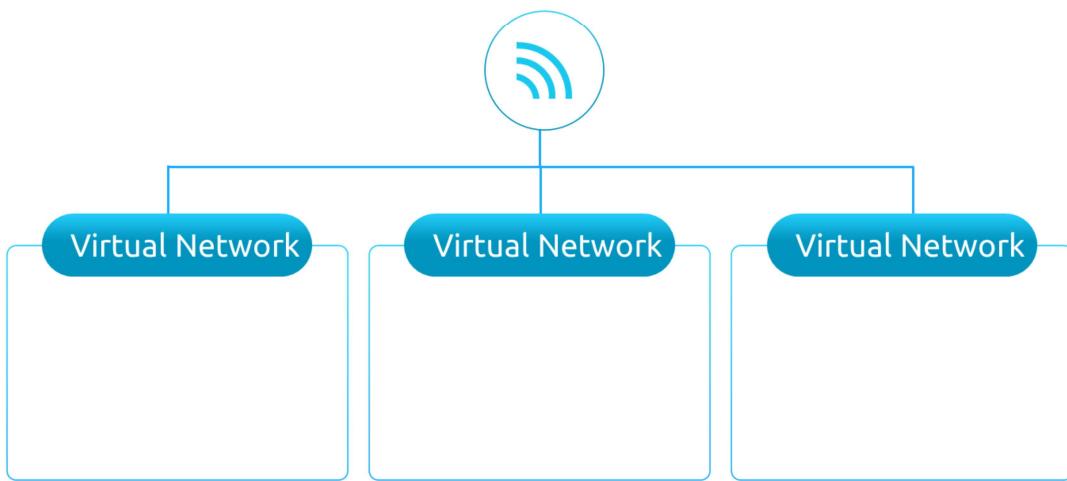
IP addresses are part of a subnet which in turn are part of a vnet.  
Azure routes traffic between subnets, connected virtual networks, on-premises networks, and the Internet, by default



You can protect resources within your subnet via Network security group by adding inbound and outbound rules to allow or deny based on IP addresses and port combination.

By default, a deny all rule is applied as the least priority rule and you will need to explicitly allow traffic to enable communication.

## Azure Virtual Network

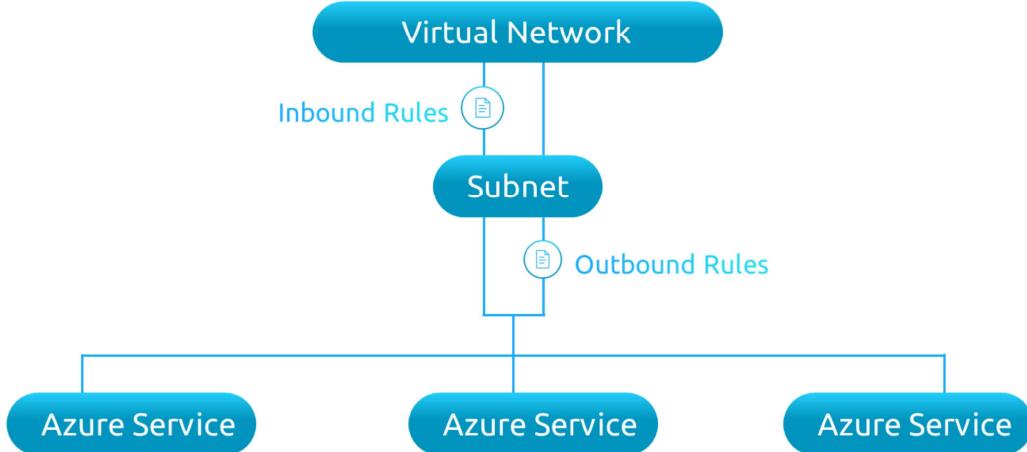


© Copyright Kodekloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

If you are coming from AWS background, it would be worth noting that subnets in Azure can span availability zones, and the same vnet can have private and public IPs.

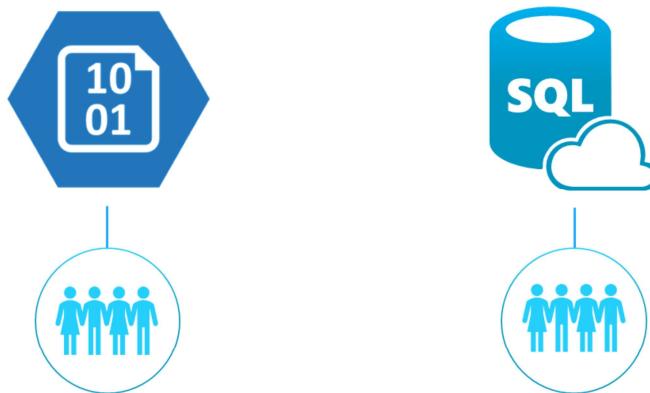
## Azure Virtual Network



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

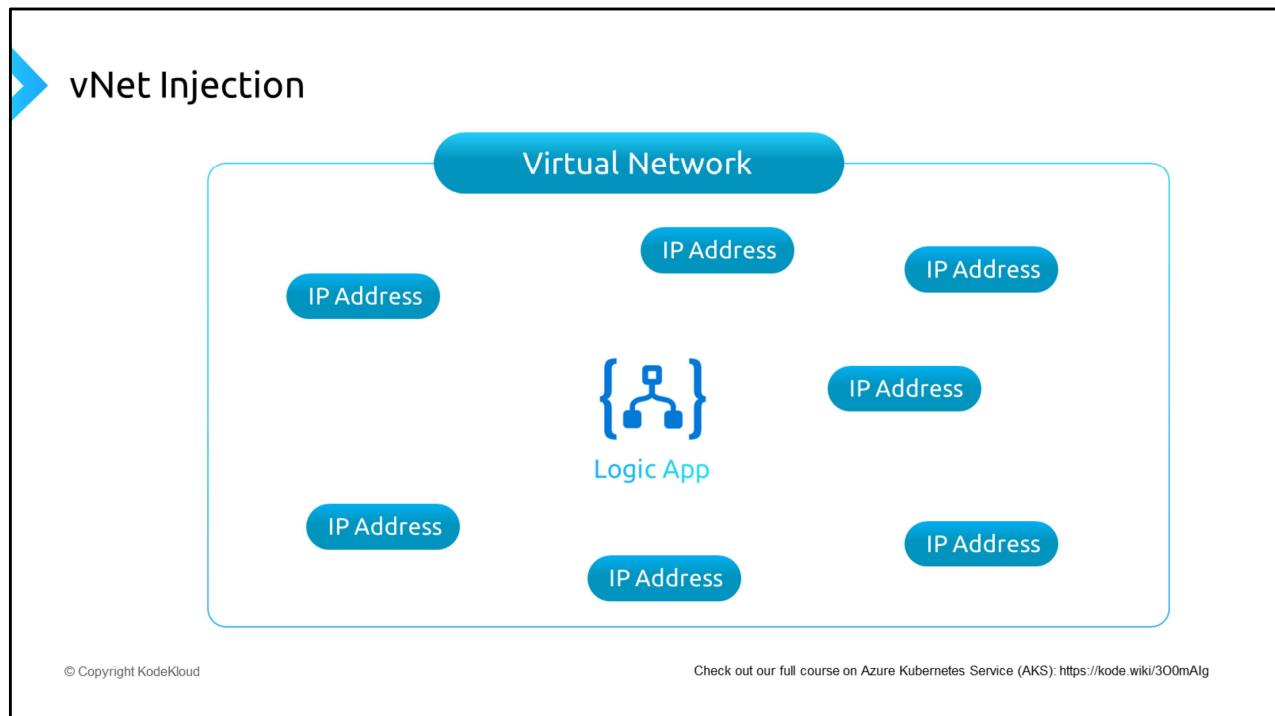
## PaaS Service Connectivity



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure multi-tenant PaaS Services like Azure Storage/SQL, etc. have public endpoints. In order to deploy and access these services privately and securely, there are three options.



vNet Injection : in which an instance of the service like Logic App is deployed in your private vNet. Because the entire instance is dedicated to you, you can assign a private IP from within you range. Not all PaaS services are vNet injection enabled and where you have the ability to do it, it is a very expensive option.

## vNet Injection

The diagram illustrates the concept of vNet Injection. It features a large blue rounded rectangle labeled "Virtual Network". Inside this rectangle, there are several blue rounded rectangles, each labeled "IP Address". A central node, represented by a blue icon with three lines and a dollar sign (\$), is connected to four of these IP address boxes. One IP address box is located above the central node, one to its left, one below it, and one to its right.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

The second method is Service Endpoint where you access multi-tenant PaaS Service like storage, SQL or Azure Container Registry using Microsoft backbone and have the ability to deny all access from public internet. This option is cost effective as the service remains multi-tenanted but you can secure it by denying public access.

➤ Service Endpoint

The diagram illustrates three Microsoft service endpoints. On the left, a blue hexagonal icon contains a white document icon with the text '10 01'. In the center, a blue cylinder icon contains a white 'SQL' text, with a small white cloud icon below it. On the right, a blue cloud icon contains a white server rack icon. Each icon is accompanied by a blue speech bubble containing the word 'Microsoft'.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

The last option is to use private link, where you can create a private IP within your IP ranges and link it to an instance of PaaS service that you deploy.

➤ Service Endpoint

© Copyright Kodekloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

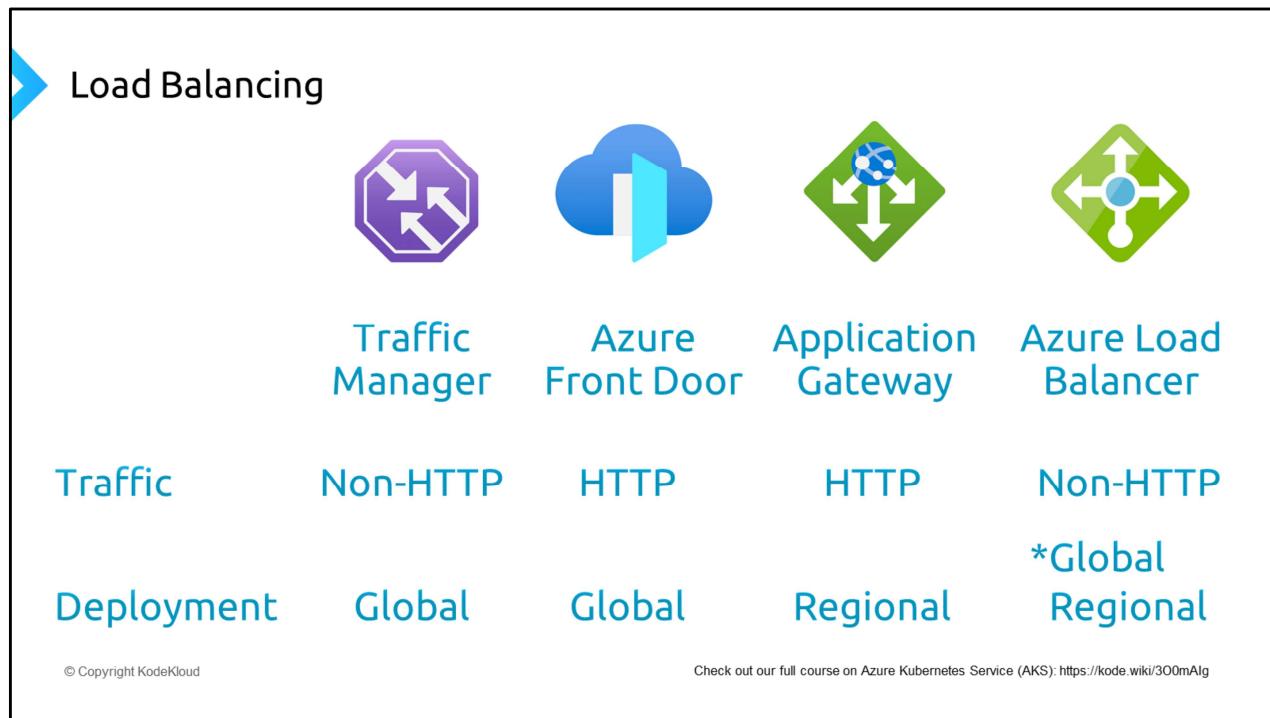
For example, you can create a private endpoint for AKS i.e Private cluster, or Azure SQL and access it via a private IP. This is by far the most secure way to access your PaaS Services

► Private Link

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

The last option it to use private link, where you can create a private IP within your IP ranges and link it to an instance of PaaS service that you deploy. For example, you can create a private endpoint for AKS i.e Private cluster, or Azure SQL and access it via a private IP. This is by far the most secure way to access your PaaS Services



There are four native load balancing options on Azure and all of them can distribute traffic to backend compute resources based on a defined criteria. What you end up using depends on the traffic type ie. Layer 7 or HTTP traffic or non-http traffic. For example, if you need to make routing decisions based on an HTTP request, you cannot use an Azure Load balancer for that.

The other important criteria is whether you need to load balance across azure regions or within a region. For example, if you need to host an active-active application with nodes across multiple azure regions, and Azure front door could be used.

At the time for this recording, Azure loadbalancer has a global option in preview and can be used for global non-HTTP traffic.

## Azure Kubernetes Services (AKS) - Summary

-  Azure Landing Zone
-  Azure Compute
-  Azure Storage
-  Azure Network & Security
-  Azure Database offering

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>



## Building and Containerizing Sample Application

-  Simple Web Application in C#
-  Containerized an Application
-  Local Docker Cluster for hosting

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

-  Topic 01: Kubernetes Overview
-  Topic 02: Deploying an AKS Cluster
-  Topic 03: Scaling the Deployment using kubectl
-  Topic 05: Pushing Image to ACR
-  Topic 06: Upgrading your application
-  Topic 07: Azure Kubernetes Fleet

© Copyright KodeKloud

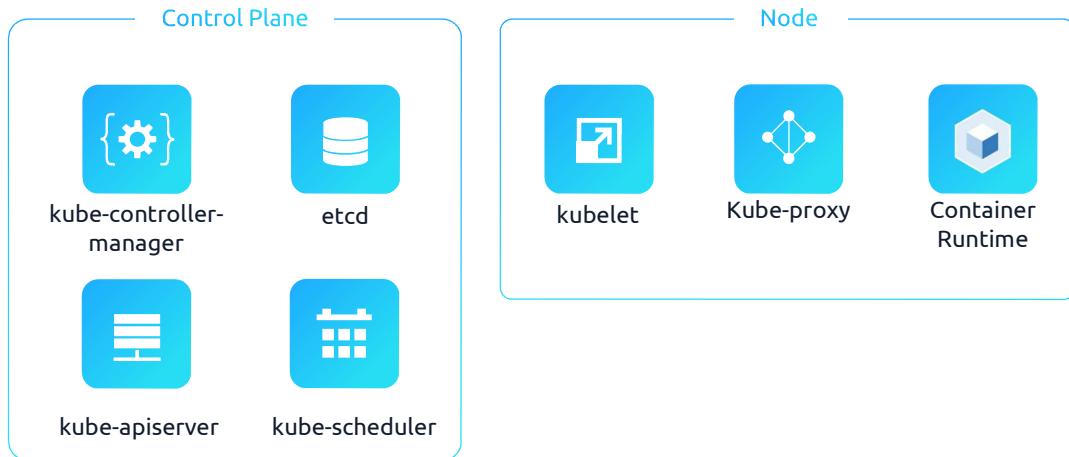
## Agenda

# Kubernetes Overview

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## Kubernetes Components



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

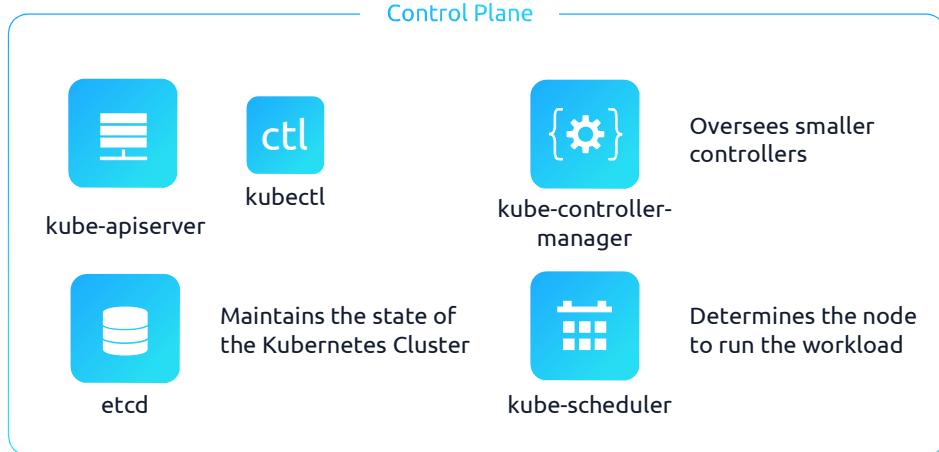
While this course is not designed to make you a Kubernetes expert, I'll cover some Kubernetes fundamentals so you can follow along the rest of the course. We, at KodeKloud, have some of the best Kubernetes courses across any platform. So do check them out later if you want to learn more about Kubernetes.

K8s architecture can be broken down into two parts: The control plane components which include the Controller manager, API Server, etcd, and scheduler  
And the node components which include Kubelet, the kube proxy, container runtime

The control plane components provides the core K8s services and orchestration of the application workloads.

While the actual application workloads run on the node component. In Azure K8s service control plain is automatically created, configured and managed by azure.

# Kubernetes Components



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

There are 4 main services running in the control plane.

kube-apiserver is how the underlying K8S APIs are exposed. This component provides the interaction for management tools, such as kubectl.

Etcd : is used to maintain the state of your K8s cluster and configuration. This highly available service is a key value store within K8s.

kube-controller-manager : oversees a number of smaller controllers that perform actions such as replicating pods and handling node operations.

When you create or scale applications, the kube-scheduler determines what nodes can run the workload and starts them

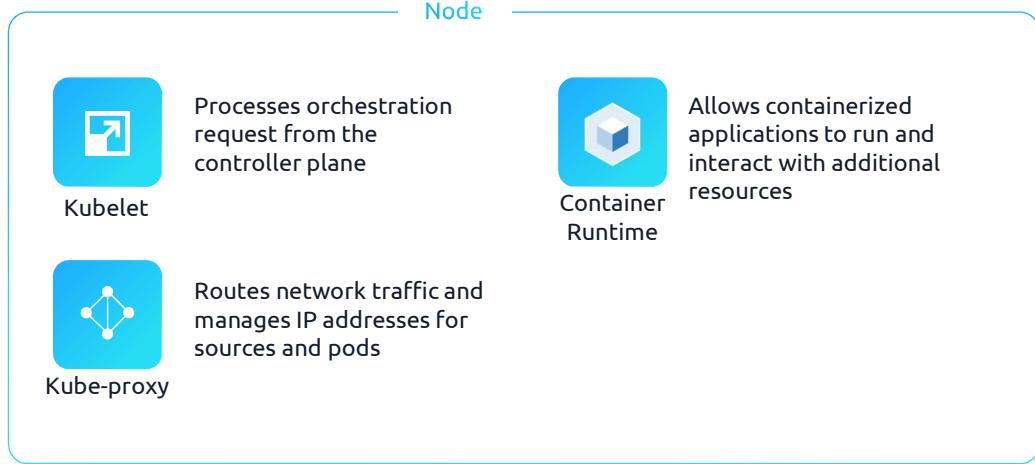
On the agent node components side, we have

K8s architecture can be broken down into two parts: The control plane components which include the Controller manager, API Server, etcd, and scheduler  
And the node components which include Kubelet, the kube proxy, container runtime

The control plane components provides the core K8s services and orchestration of the application workloads.

While the actual ..application workloads run on the node component. In Azure K8s service control plain is automatically created, configured and managed by azure.

## Kubernetes Components



© Copyright KodeKloud

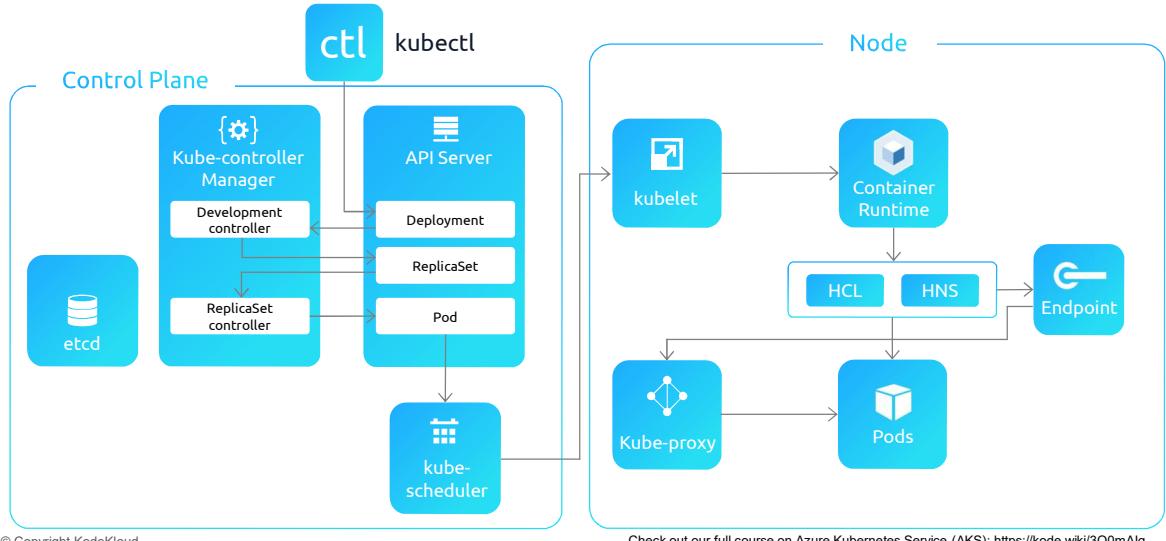
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Kubelet : which is a k8s agent that processes the orchestration request from the control plane and schedule the requested containers

Virtual networking is handled by the kube-proxy on each node. The proxy routes network traffic and manages IP addressing for services and pods.

The container runtime is the component that allows containerized applications to run and interact with additional resources such as virtual network and storage.

# Kubernetes Components



Now that you know some of the components and their high-level functionality, let us see how these components talk to each other when you issue a command to create a new Kubernetes deployment.

Do not worry if this ends up looking very complicated, you do not need to remember all of this. Just remember there are multiple subcomponents with a well-defined function for each of them.

- 1: KubeCTL gives a POD Deployment manifest to the API Server.
- 2: API Server Creates Deployment Resource and that info is saved in etcd.
- 3: Deployment Controller watches for any new deployment resource and creates the replicaset resource.
- 4: Replicaset controller watches for any new replicaset resource.
- 5: Replicaset controller creates the POD resource based on how many replicaset and actual pods are present.
- 6/7: Kube Scheduler watches for unbound POD resource and schedules them on to a node.
- 8: Kubelet calls Container Runtime Interface (CRI) and (Container Network Interface) to create the POD/Containers with networking.
- 9&10: CRI calls Host Compute Layer (HCS) to create the Container and Host

Networking Service (HNS) to create the Endpoint.

11&12: Kubeproxy watches for any new Endpoint and programs HNS with Load Balancer and Access Control List (LB and ACL) Policies.

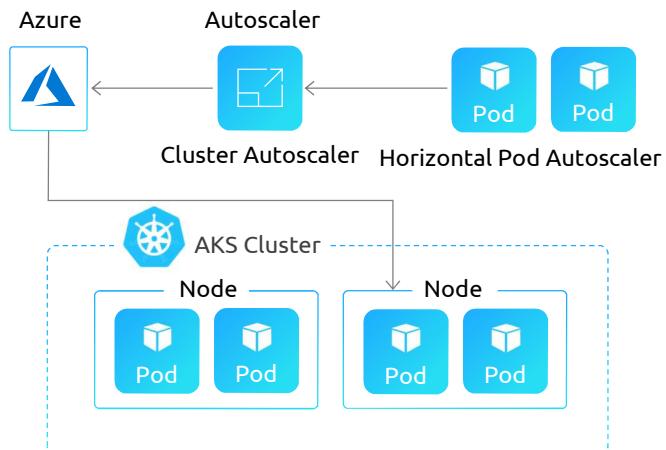
Now that we have covered some k8s fundamentals, let us start deploying a cluster in Azure.

# Scaling the Nodes using Azure CLI

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ▶ Cluster Autoscaler



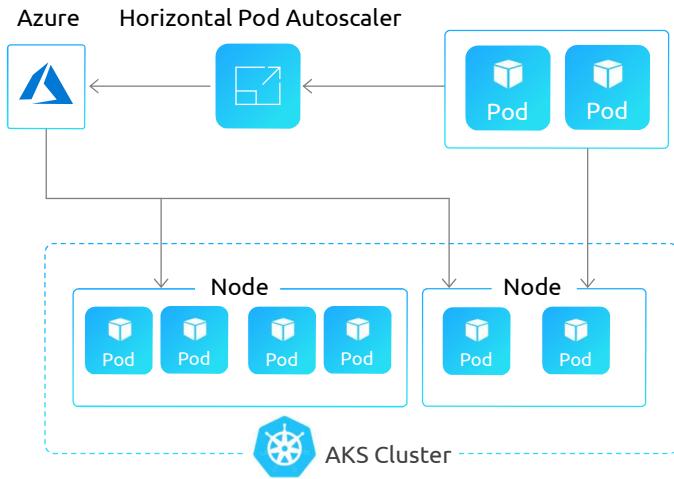
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

*We ended the last video with a cluster that has been configured to run 30 pods and is running only 16.*

# Cluster Autoscaler

1. Horizontal Pod Autoscaler (HPA) obtains resource metrics and compares them to user-specified threshold
2. HPA evaluates whether user-specified threshold is met or not
3. HPA increases/decreases the replicas based on the specified threshold
4. Deployment controller adjusts the deployment based on increase/decrease in replicas



© Copyright KodeKloud

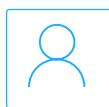
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

As I mentioned in the previous slide, Kubernetes uses the horizontal pod autoscaler (HPA) to monitor the resource demand and automatically scale the number of pods. HPA is updated every 60 seconds with key metrics like CPU utilization to take scale out or scale in decisions.

This is 4 step process.. In step 1..

## ► Azure Container Registry (ACR)

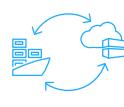
Manage a Docker private registry



Integrated  
with AAD



Manage images for all  
types of containers



Use familiar, open-  
source Docker CLI tools



Azure Container  
Registry geo-replication

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

A container *registry* like Docker hub is a service that stores and distributes container images and related artifacts. In the previous module we containerized an application and pushed it to Docker hub.

Azure Container Registry is a private registry and provides users with direct control of their container content, with integrated authentication, geo-replication, distribution and, virtual network configuration with Private Link, and many other enhanced features.

In addition to Docker-compatible container images, Azure Container Registry supports a range of content artifacts including Helm charts and Open Container Initiative (OCI) image formats.

The *tag* for an image or other artifact specifies its version. A single artifact within a repository can be assigned one or many tags and may also be "untagged." That is, you

can delete all tags from an image, while the image's data ( its layers) remain in the registry.

## Azure Container Registry (ACR) Tiers

Resource	Basic	Standard	Premium
Included storage (GiB)	10	100	500
Storage limit (TiB)	20	20	20
Maximum image layer size (GiB)	200	200	200
Maximum manifest size (MiB)	4	4	4
ReadOps per min	1000	3000	10000
WriteOps per min	100	500	2000
Download Bandwidth (Mbps)	30	60	100
Upload Bandwidth (Mbps)	10	20	50
Private endpoints	n/a	n/a	200
Geo-replication	n/a	n/a	Supported

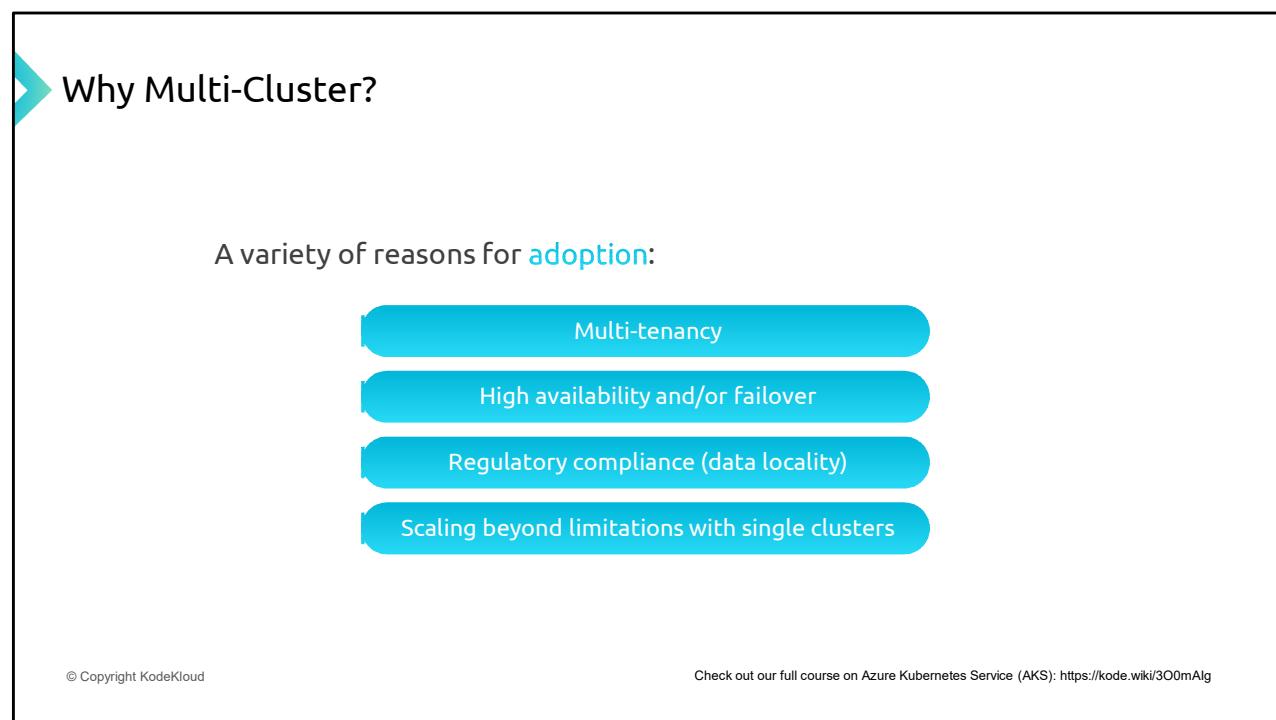
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Azure Container Registry is available in multiple SKUs. I have added some differences in the table here. This should help you choose a tier that can work for you. In case you want to upgrade, there is a command to upgrade. You can move freely, both up and down, between tiers as long as the tier you're switching to has the required maximum storage capacity. For example, if you are using 200 GB in a premier tier and run a command to downgrade ACR to standard, that change will fail as the maximum storage included in Standard tier is 100 GB.

If you recall, we had created a ACR when creating AKS cluster. Now lets push our

image to that registry.



## Why Multi-Cluster?

A variety of reasons for adoption:

- Multi-tenancy
- High availability and/or failover
- Regulatory compliance (data locality)
- Scaling beyond limitations with single clusters

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Until now, we have seen how to manage a single Kubernetes cluster in AKS. In real life, you will be dealing with 100s if not 1000s of AKS clusters.

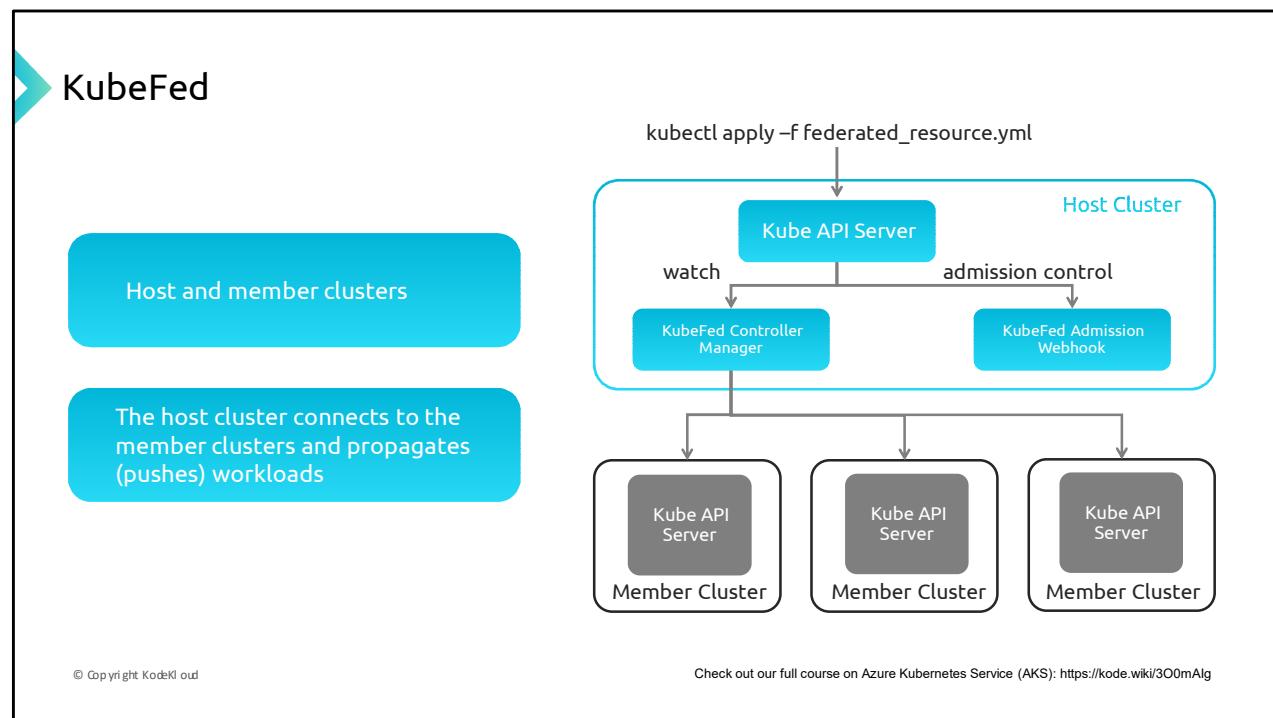
Multi cluster system, by itself is not a new concept as you can imagine many large organisations would have hundreds of Kubernetes clusters spread across the world to host their applications.

A logical question to ask here, is why do I need multi cluster systems in the first place ? Can I not host all my applications in a single cluster across Pods and namespaces ?

Multi-cluster pattern are used for variety of reasons – Some use it to achieve multi tenancy, other use it as it provides a convenient way to achieve high availability or failover, making sure that a failure of a single cluster region or even cloud provider will not impact your application's availability.

If there is nothing else, you can always depend on compliance as a reason to complicate matters, haha but honestly there are real compliance reasons requiring your application to be hosted on separate Kubernetes clusters across geographic boundaries. For example, in US if you are hosting a betting app, by law you are required to host that app in each state boundary.

If your application is huge, you can also use multi cluster system to scale beyond prescribed limits.



So let's look some of the open source solutions for multicluster management.

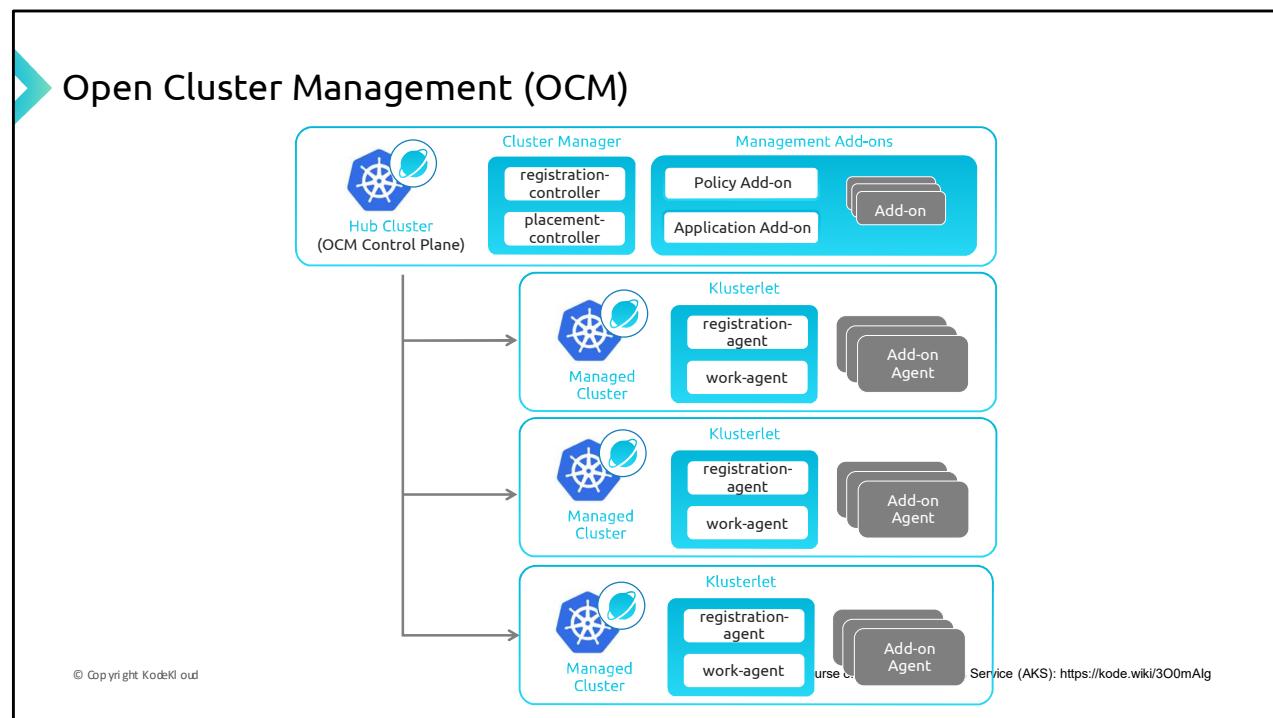
The first one is Kubernetes Cluster Federation aka KubeFed, the official multicluster solution sponsored by Multicluster Special Interest Group.

It is built on top of the concept of Kubernetes federation, the core idea of which is to provide a way to coordinate services among a group of Kubernetes clusters. This project has a long and complicated history and has two major versions so far – Kubefed v1 and kubefed V2. V1 is already deprecated.

Multi-cluster patterns are used for variety of reasons – Some use it to achieve multi-tenancy, others use it as it provides a convenient way to achieve high availability or failover, making sure that a failure of a single cluster region or even cloud provider will not impact your application's availability.

If there is nothing else, you can always depend on compliance as a reason to complicate matters, haha but honestly there are real compliance reasons requiring your application to be hosted on separate Kubernetes clusters across geographic boundaries. For example, in US if you are hosting a betting app, by law you are required to host that app in each state boundary.

If your application is huge, you can also use multi cluster system to scale beyond prescribed limits.



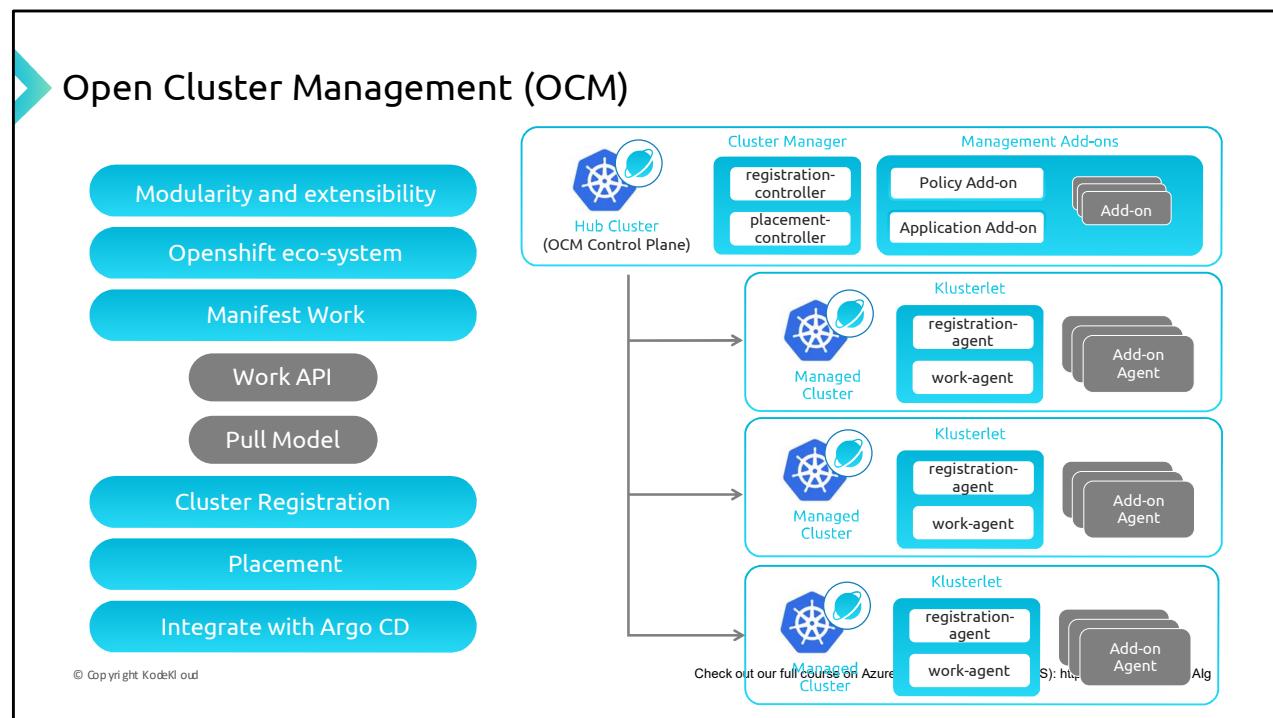
## Open Cluster Management (OCM)

is a community driven project focusing on multicluster and multicloud deployments. Interesting fact - this project is backed by redhat and you will see them actively backing this project.

<https://open-cluster-management.io/concepts/>

The orchestration in OCM is similar to K8s, in a sense that they have a hub cluster that mimics K8s control plane and spoke clusters have an agent called Klusterlet similar to kubelet in K8s. Klusterlet in turn have registration and work agents.

If you want to know more about this system, you can access their repo via the link on the screen. Key concepts to understand before you can use OCM is their Cluster Registration process that you see on the screen now and how they use placement and scheduling. If you or your organisation is using OCM I would love to get your feedback on how its working in AKS. Please share your comments in this module.



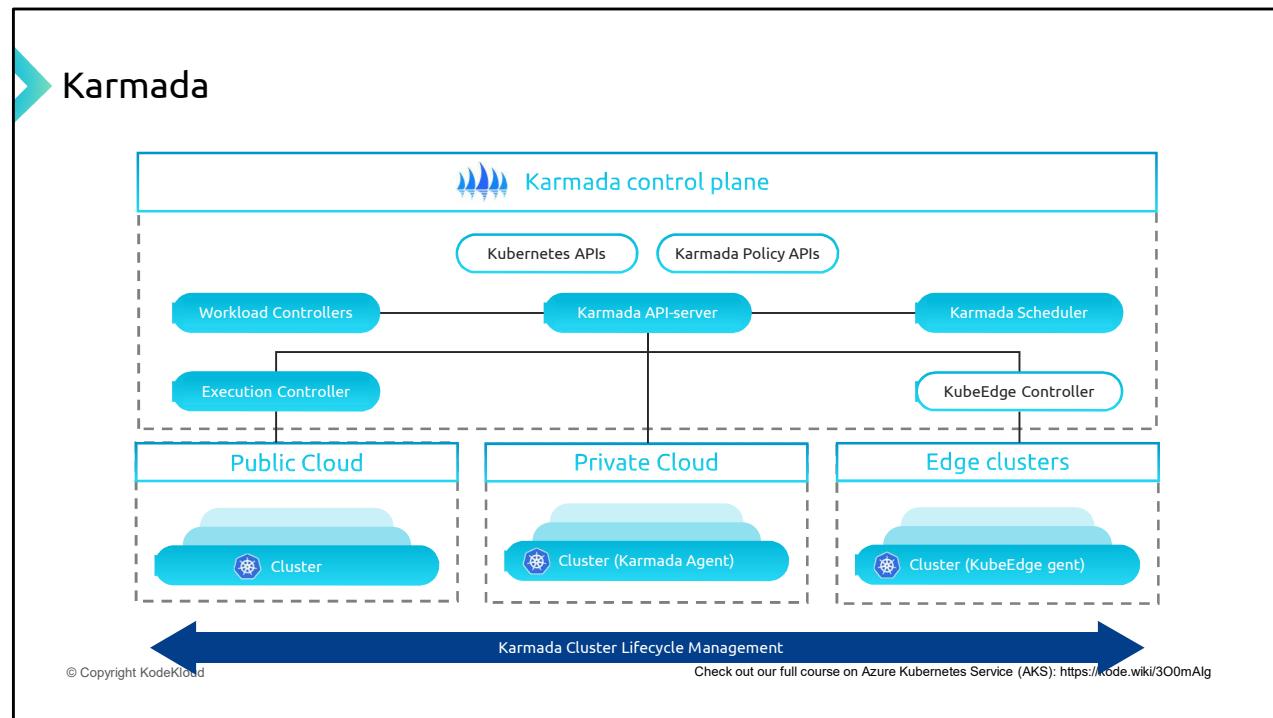
## Open Cluster Management (OCM)

is a community driven project focusing on multicluster and multicloud deployments. Interesting fact - this project is backed by redhat and you will see them actively backing this project.

<https://open-cluster-management.io/concepts/>

The orchestration in OCM is similar to K8s, in a sense that they have a hub cluster that mimics K8s control plane and spoke clusters have an agent called Klusterlet similar to kubelet in K8s. Klusterlet in turn have registration and work agents.

If you want to know more about this system, you can access their repo via the link on the screen. Key concepts to understand before you can use OCM is their Cluster Registration process that you see on the screen now and how they use placement and scheduling. If you or your organisation is using OCM I would love to get your feedback on how its working in AKS. Please share your comments in this module.

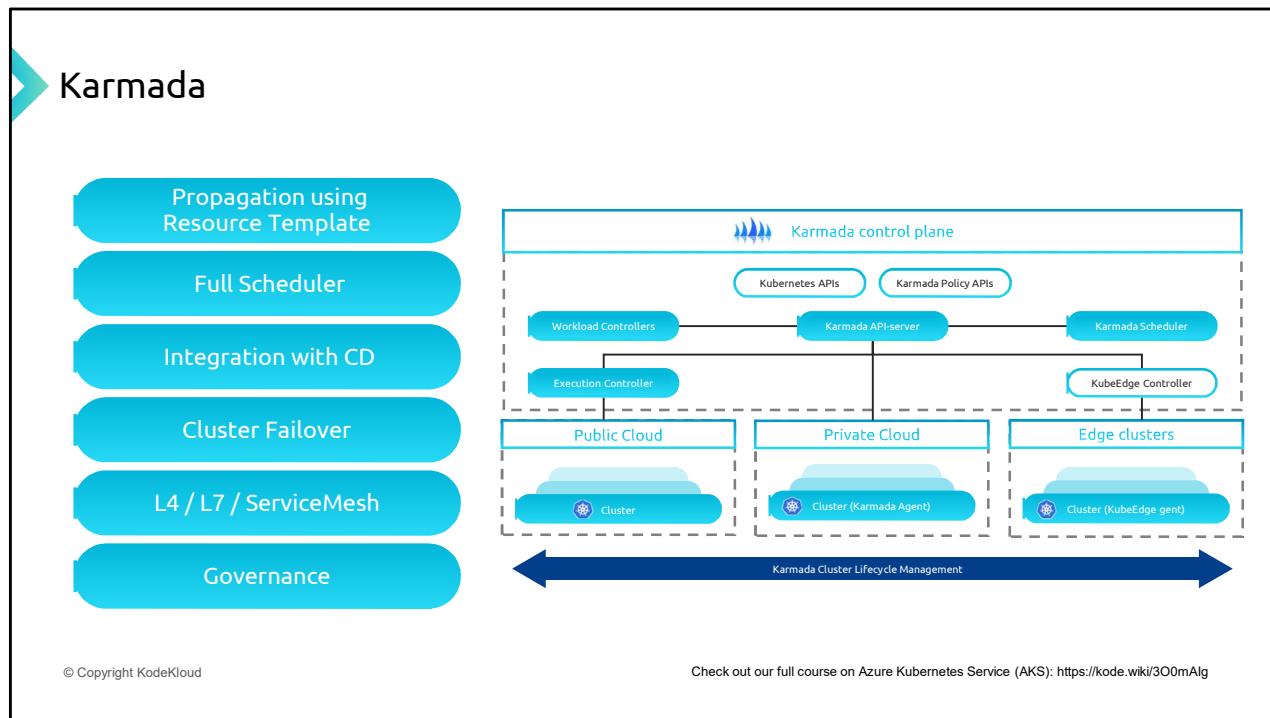


The next opensource solution for multi cluster management that we will touch upon is Karmada.

On the screen, you can see Karmada's design that is taken from their website and some aspirations and features they have. They have couple of really interesting features: if you look at the karmada control plane they use their API server. This two tier approach is different than OCM in that Karmada has another API server instead of using K8s API server implementation. You also need to install another K8s control plane in the clusters that join Karmada which is the karmada agent.

You can also observe that Karmada Workload controllers and scheduler talks to Karmada API-server and not to K8S native API servers.

If you want to know more about this implementation, check out the link of the screen: <https://karmada.io/docs/>



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

The next opensource solution for multi cluster management that we will touch upon is Karmada.

On the screen, you can see Karmada's design that is taken from their website and some aspirations and features they have. They have couple of really interesting features: if you look at the karmada control plane they use their API server. This two tier approach is different than OCMIn that Karmada has another API server instead of using K8s API server implementation. You also need to install another K8s control plane in the clusters that join Karmada which is the karmada agent.

You can also observe that Karmada workload controllers and scheduler talks to Karmada API-server and not to K8S native API servers.

If you want to know more about this implementation, check out the link of the screen: <https://karmada.io/docs/>

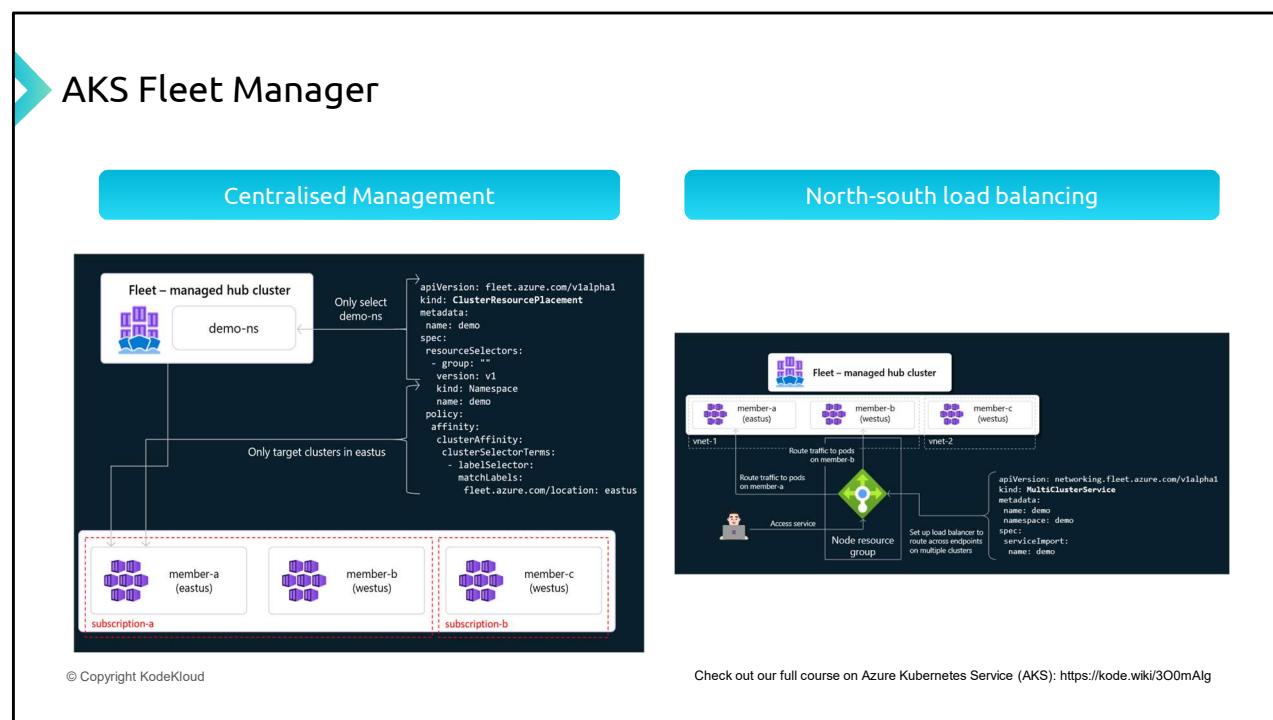
## Enterprise Options



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

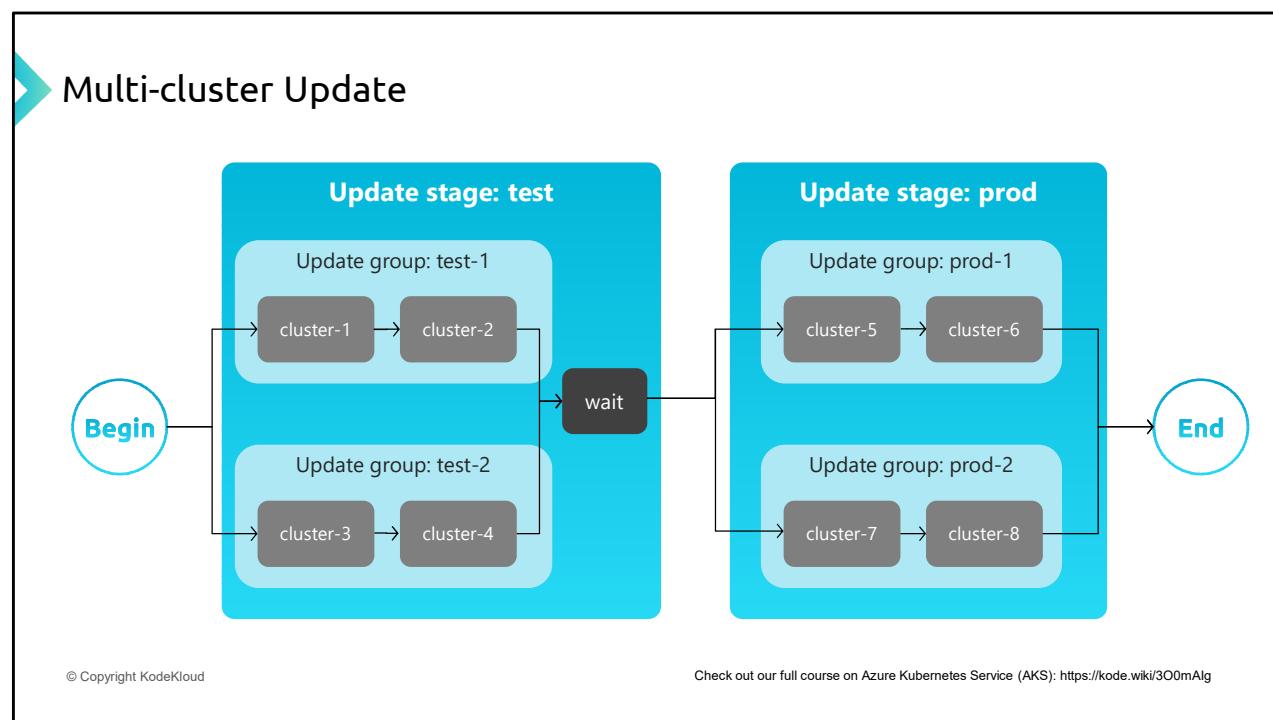
When it comes to enterprises, there are few of notable Multi-cluster management tools that you will come across and can see on your screen. Azure AKS Fleet manager is a new entrant and



Azure Kubernetes Fleet Manager (Fleet) is a managed service that helps you centrally manage multiple Kubernetes clusters at scale. Fleet provides a single pane of glass for managing your clusters, and it can automate tasks such as cluster provisioning, upgrades, and configuration management.

There are 4 key features of Azure Kubernetes Fleet Manager:

- **Centralized management:** Fleet provides a single pane of glass for managing your Kubernetes clusters. You can view all of your clusters in one place, and you can perform operations on them in bulk.
- **Automated tasks:** Fleet can automate tasks such as cluster provisioning, upgrades, and configuration management. This can help you reduce operational overhead and improve the reliability of your Kubernetes deployments.
- **Policy-based management:** Fleet supports policy-based management. This means that you can define policies that govern how your Kubernetes clusters are managed. For example, you can define policies that specify the minimum version of Kubernetes that your clusters must run, or the maximum number of pods that a cluster can have.
- **North-south load balancing:** Fleet can provide north-south load balancing for your Kubernetes clusters. This means that you can distribute traffic across multiple clusters, which can improve the performance and availability of your applications.



You can also use Kubernetes Fleet Manager to orchestrate multi-cluster updates. This means that you can update your clusters in a planned manner across multiple clusters and multiple environments to achieve a consistent environment across the set of clusters. This can also help test the updates in lower environments before deploying to production. This can be fully automated using Update groups.

## Summary

-  Fundamental Concepts and Benefits
-  How to deploy AKS Cluster
-  Deploying application to AKS Cluster
-  Scaling
-  Image Management and Pushing Image to ACR
-  Upgrading Application
-  Azure Kubernetes Fleet

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

# Network Security

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In this networking module, we will cover Networking options in AKS including configuration options like Kubenet, Azure CNI and networking policies. When we deployed AKS cluster in the previous module, we did not change the default networking options. It is now time to look a little deeper on available options and how you can use them to design your deployments.

There are three videos on this topic, in the first video I'll talk about some basics like virtual network, Subnets, NSGs and UDRs

I'll follow that up with details around Kubenet and

CNI.

The final video is around Network policies.. Lets get started

-  Topic 01: Networking Options
-  Topic 02: Configuration Options
-  Topic 03: Networking Policies

© Copyright KodeKloud

## Agenda

## Networking Security

- Virtual Networks, Subnets, NSGs, and UDRs
- Kubenet and Azure CNI
- Network Policies

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Ok, we have now reached the culmination of this video course on Azure Kubernetes Service (AKS). Our journey began by delving into the fundamental aspects of Azure, encompassing essential concepts such as Compute, Storage, and Networking. With a some foundation in place, we transitioned towards the creation and containerization of a basic C# web application. Throughout that module, we focused on leveraging Docker desktop and crafting Dockerfiles to facilitate containerization.

Subsequently, we navigated to the Azure portal to establish an AKS cluster to deploy our world class application onto that cluster. Harnessing the power of kubectl, we effectively executed various imperative actions, including scaling of deployments. we explored the deployment of container images within the Azure Container Registry (ACR), unearthing its potential for application enhancement and version management. By harnessing the image hosted on ACR, we facilitated a seamless application upgrade, only to effortlessly revert to a previous version when needed. Concluding this module, we bestowed upon you a overview of the Kubernetes Fleet, empowering you with a holistic understanding of this powerful toolset.

The subsequent module, entirely dedicated to AKS networking options and policies,

showcased the indispensable importance of robust network management within AKS deployments. Furthermore, we delved into security-centric topics, delving into the intricacies of Azure AD integration, defender for AKS, and Azure policy for AKS. Empowered with this knowledge, you are poised to fortify your AKS infrastructure with optimal security measures and governance practices.

As we progressed, we ventured into the realm of CI/CD patterns within AKS, looking at the art of managing AKS using declarative techniques. By adopting these practices, you can streamline the continuous integration and delivery pipelines within your AKS environment.

An indispensable component of any infrastructure, observability received some attention in module seven. We explored observability within AKS, equipping you with the necessary tools and techniques to monitor and gain valuable insights into your AKS deployments.

Finally, before we brought this course to a close, we presented you with an overview of alternative platforms that embrace the power of containers. By expanding your horizons and exploring these platforms, you can unlock a world of possibilities and further enhance your containerization endeavors.

We sincerely hope that your journey throughout this course has been as fulfilling as our commitment to its production. It has been an absolute pleasure to guide you through the intricacies of AKS, empowering you with the knowledge and skills needed to excel in your containerization journey. This is your host Pranav, signing off for now.

# Azure Networking Fundamentals

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## VNet



© Copyright KodeKloud

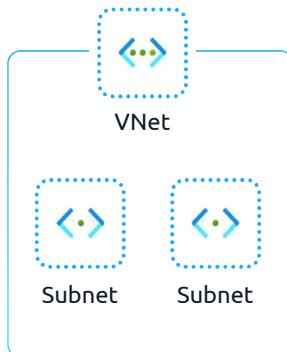
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Ok, before we go any further, let us familiarize ourselves with some Azure networking nomenclature.

A Virtual network or vnet is an Azure construct which defines one or more private IP Address prefixes that you intent to use in your environment. For example, (Subnet screen), I can define a vnet CIDR for vnet1 as 10.2.0.0/16, which can allow me to use about 65534 IP addresses. CIDR, by the way, stands for Classless Inter-Domain Routing. You can use this or many IP address management portals to plan your IP ranges. In simple terms, we define Network address and Mask in this format to define what range of Ips we can use.

Azure Vnets support both IPv4 and IPv6 address spaces, for simplicity we will use IPv4 here.

## ► Subnet

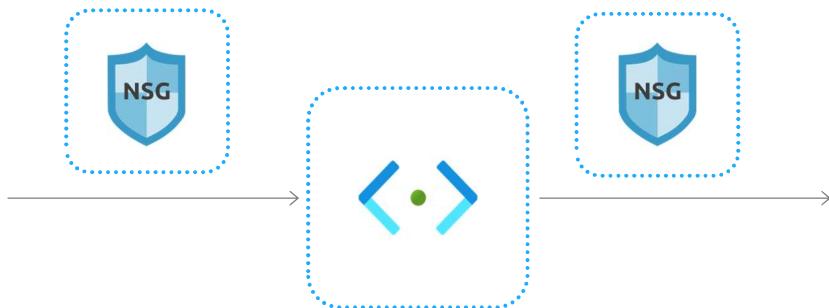


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Before you can allocate any of the IP addresses from a vnet to a container, Mor any other resource, you need to define atleast one subnet, again in CIDR format. Let us define two. They offcourse, need to be a subset of what we defined in the vnet and cannot overlap.

## ► Network Security Group (NSG)



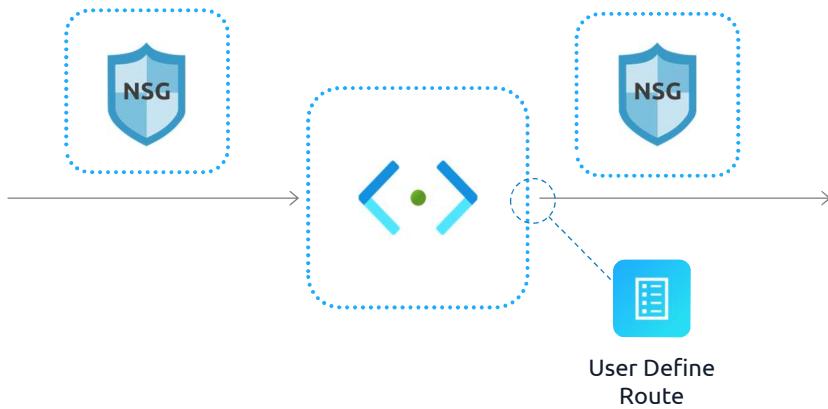
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Next constructs to understand are Network Security group or NSG and User Defined route or UDR.

NSGs are used to filter network traffic between Azure resources that traverse a subnet. A NSG contains security rules that allow or deny inbound or outbound network traffic. For each rule, you can specify source and destination, port, and protocol. Think of this as a lightweight firewall.

## ► User Define Route (UDR)

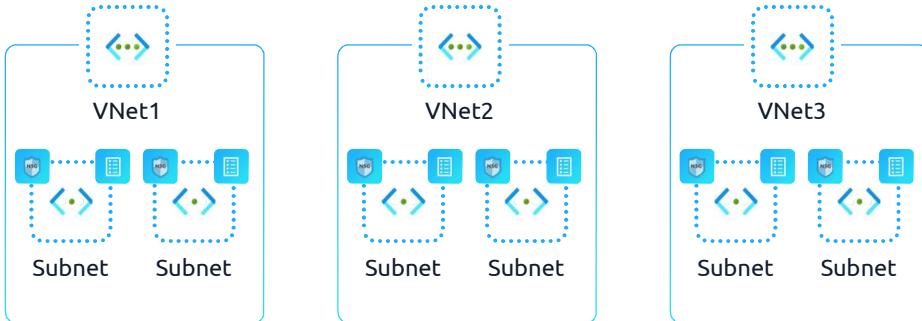


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

UDR defines how traffic initiated in that subnet is routed. If your pod needs to talk to another pod hosted in a node in another subnet or vnet, azure needs to know where to move that traffic. A route table is a collection of individual routes used to decide where to forward packets based on the destination IP address. There are three ways to populate route tables: User defined routes, that you see in the screen, BGP routes based on data received via Border gateway protocol and System routes.

## ▶ Subnet

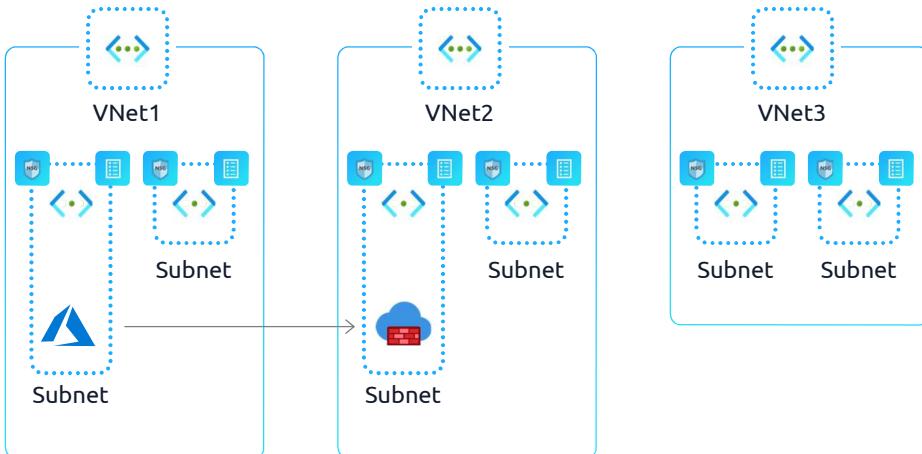


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Once you define your NSGs and UDRs, both of which are optional, you start deploying resources in your subnet.

## ▶ Subnet



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

And as you deploy more services, you will have multiple vnets and a defined method to route the traffic. For example, you might decide to push all traffic from your frontend subnet in vnet1 to a firewall in vnet 3 before it goes anywhere else.

If you have more than one vnet, you will need to peer them if you need resources in one vnet to talk to resources in another vnet.

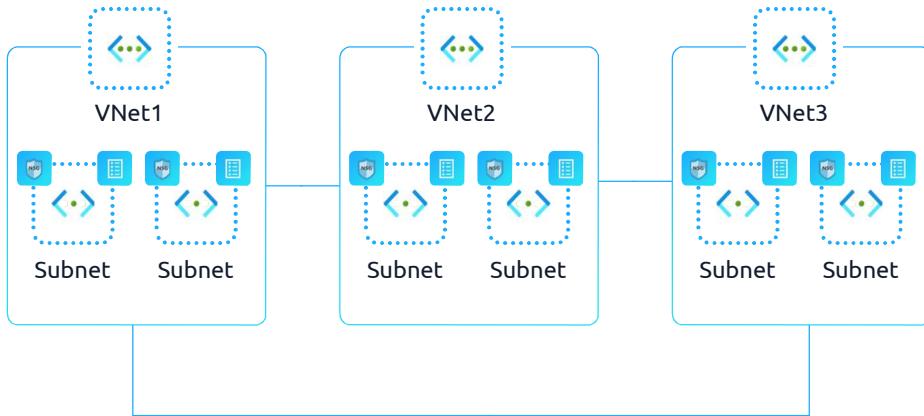
Now that you know these Azure networking concepts lets jump to AKS networking.

# AKS Networking Options

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ▶ Subnet



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

There are two networking options to chose from when you deploy K8s – Kubnet and Azure CNI. these used to be called basic and advanced options respectively.

Depending on the option you choose, that plugin gets deployed on every VM of the AKS cluster providing networking for the cluster and to the nodes and pods inside them.

With *kubenet*, nodes get an IP address from the Azure subnet defined by you while Pods receive an IP address from a logically different address space to the subnet of the nodes. Network address translation (NAT) is then configured so that the pods can reach resources on the Azure virtual network.

By default, AKS assigns a /24 address space to each VM, you can optionally specify what that address pool to use.

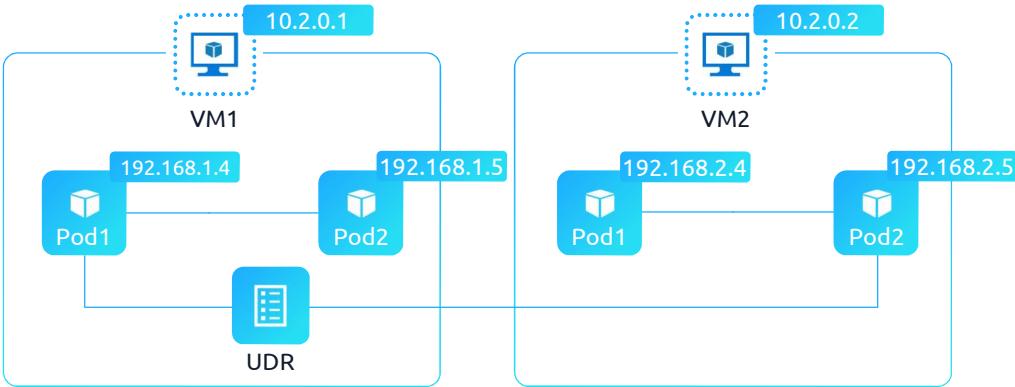
Pods on the same VM can communicate with each other via a bridge. For Pods to communicate across VMs, aks uses UDR to route traffic. You know what UDRs are already, so if Pod 1 from VM1 need to talk to Pod1 from VM2, there will be a UDR on

subnet1 that will force traffic to VM2 so it can reach the destination pod.

If your Pod needs to communicate to another resource outside the aks cluster, AKS uses SNAT i.e Source Network Address Translation.

Lets say our first pod with IP 192.168.1.4 needs to talk to a VM on another Vnet with IP 10.10.1.1, AKS will use the VMs IP where the pod is hosted to SNAT the traffic to the destination VM. Since SNAT is being used the destination will use Source VM IP and pre designated port combination to respond back to any requests.

## ▶ Networking Options for AKS - KubeNet

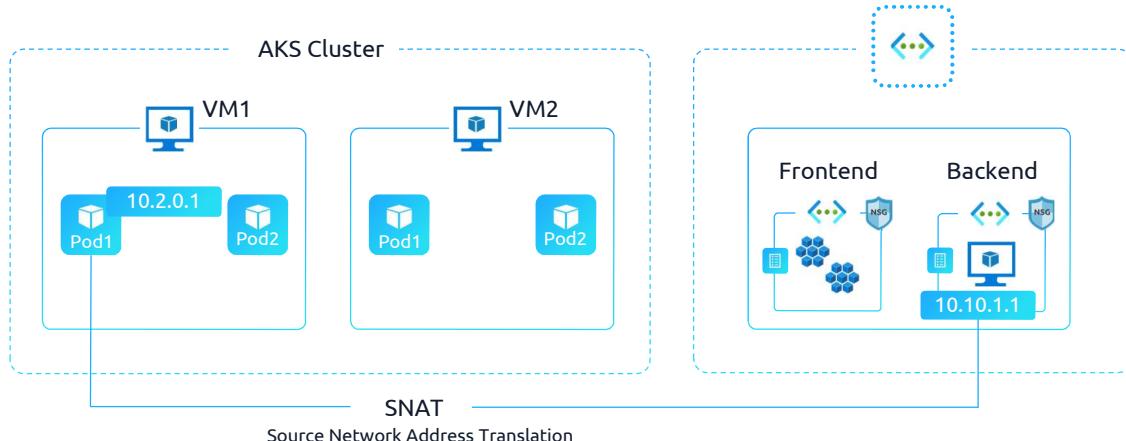


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Pods on the same VM can communicate with each other via a bridge. For Pods to communicate across VMs, aks uses UDR to route traffic. You know what UDRs are already, so if Pod 1 from VM1 need to talk to Pod1 from VM2, there will be a UDR on subnet1 that will force traffic to VM2 so it can reach the destination pod.

## ➤ Networking Options for AKS - KubeNet



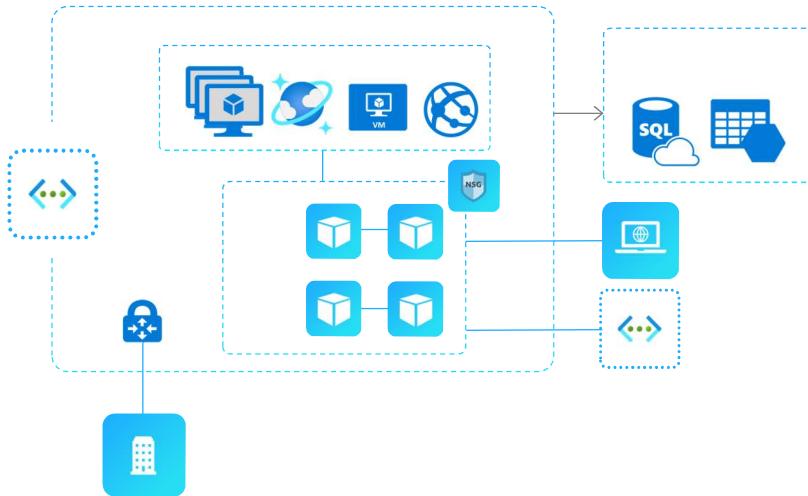
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

If your Pod needs to communicate to another resource outside the aks cluster, AKS uses SNAT i.e Source Network Address Translation.

Lets say our first pod with IP 192.168.1.4 needs to talk to a VM on another Vnet with IP 10.10.1.1, AKS will use the VMs IP where the pod is hosted to SNAT the traffic to the destination VM. Since SNAT is being used the destination will use Source VM IP and pre designated port combination to respond back to any requests.

## Azure CNI



© Copyright KodeKloud

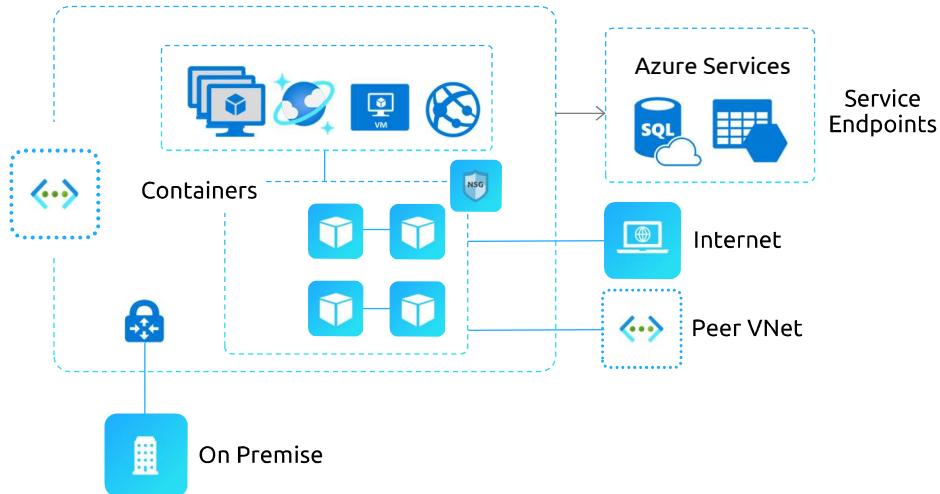
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

As I mentioned before if you choose Azure CNI, the CNI plugin gets deployed on every VM and unlike Kubenet, this can assign an IP from your vnet to your pods.

When you deploy a cluster with Azure CNI, you specify which of your vnet and subnet you want to deploy the cluster in. When you do that, not only the VMs in the cluster get IPs from your subnet, the pods inside the cluster also get IPs from subnet controlled by you.

This makes the pod, directly routable from any other system, be it another AKS pod or a VM or in fact from your on-premise system. This provides seamless and directly connectivity to your pods. And this also means that your pod can also talk to any other service on azure and on-premise as long as firewalls allow that to happen..

## Azure CNI



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

NSGs and UDRs that you would have defined for your subnets would automatically apply to the pods in your cluster. In the next video, we will talk about network policies, I'll have more to talk about NSGs and UDRs when it comes to aks pods.

Kubernetes also have lot of networking capabilities that it offers natively like ingress controllers, DNS etc so all these will also work seamlessly when you use it via CNI plugin.

Can you think of a reason why Azure CNI wil have much better network performance than Kubenet ? I'll answer this question at the end of this video.

## Azure CNI

Why do Azure CNI have better Network Performance than Kubenet?

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Can you think of a reason why Azure CNI will have much better network performance than Kubenet? I'll answer this question at the end of this video.

## ➤ Azure CNI

An implementation of Container Networking Interface that is in line with Cloud Native Computing Foundation (CNCF) Specification.

**Calico**

 flannel

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Lets take a deeper look at what CNI is doing. CNI is an implementation of container networking interface that inline with CNCF specification. These are many vendors which have deployed a flavor of this specification. like Calico, flannel etc.

## CNI Modules

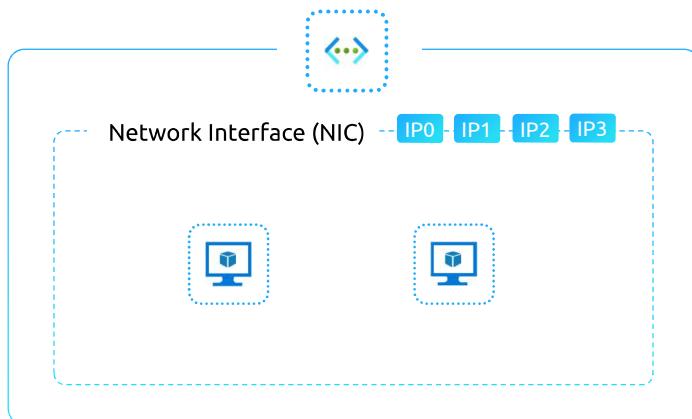


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Any CNI have two modules to it – one is networking that is when a container is initiated it get a NIC card assigned to it so it has connectivity to talk to other containers and resources. The other function is IP Address management or IPAM. After a NIC card is assigned, it needs an IP address from your subnet. The function to allocated an IP from a pool and deallocated it when the container is terminated is done by this IPAM module.

## ▶ IP Management

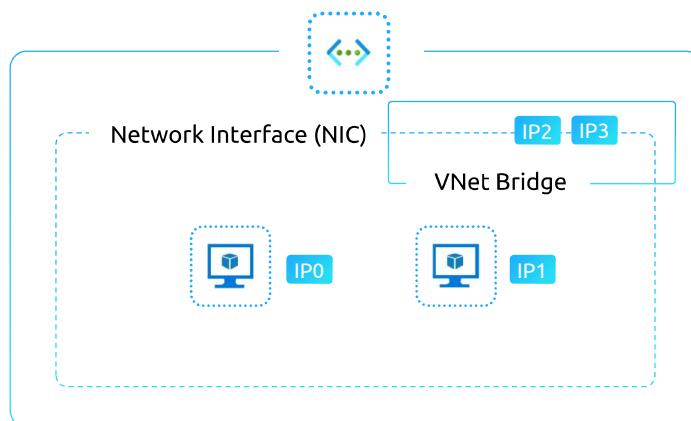


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

The way the IP address management works in Azure CNI is when you provision your cluster, Azure assigns multiple set of IPs to the NIC. Let us say I have one cluster with 2 VMs and each of those VMs have a single NIC card. Every NIC has a primary IP, which is the address of the VM. Apart from that IP, Azure adds another set secondary of IPs to that NIC. The number of IPs assigned is the same that you would have configured as max pods that you want to deploy per VM. When you are deploying a cluster via CLI, you have the option to specify how many maximum pods you want to deploy. Based on that, number of secondary IPs are assigned to the NIC.

## ▶ IP Management

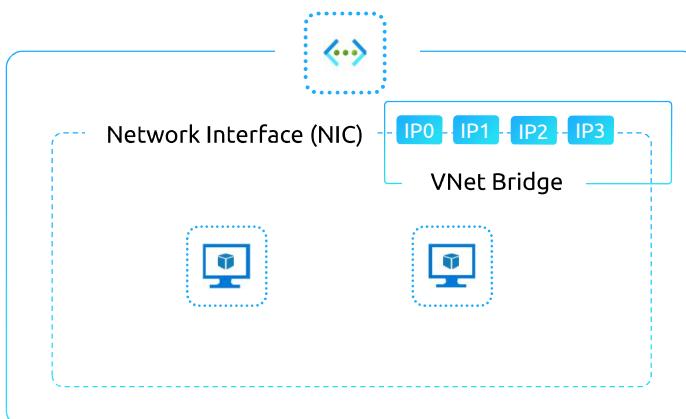


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

As your containers get deployed, the CNI on that node will draw an IP from that pool of IPs and assign it to the container. When the container terminates CNI puts it back to the pool and available IPs. Vnet Bridge is used by CNI to manage these interactions.

## ▶ IP Management



Azure Supports:

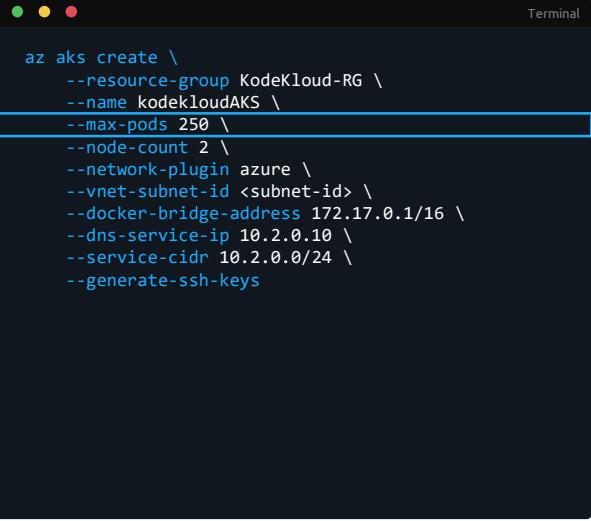
250 Containers / pod / VM  
64,000 Ips / VNet

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

At the time of this recording, Azure support 250 containers or Pods per Vm and 64K IP per vnet i.e a /16 range.

## Code



```
az aks create \
--resource-group KodeKloud-RG \
--name kodekloudAKS \
--max-pods 250 \
--node-count 2 \
--network-plugin azure \
--vnet-subnet-id <subnet-id> \
--docker-bridge-address 172.17.0.1/16 \
--dns-service-ip 10.2.0.10 \
--service-cidr 10.2.0.0/24 \
--generate-ssh-keys
```

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

We can use labels to select the type of pods and then there is a separate section in the yaml file that lets you specify what the policies are which are essentially filtering rules that are defined based on labels on names or namespaces or IPBlocks as shown in this screen.

## Kubenet vs. CNI

Capability	Kubenet	CNI
Pod-VM connectivity; VM in the same VNet	Works when initiated by Pod	Works both ways
Pod-VM connectivity; VM in peered VNet	Works when initiated by Pod	Works both ways
On-premises access, over VPN/ER	Works when initiated by Pod	Works both ways

© Copyright Kodekloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/summa/ig>

Just to contrast both the options, while CNI is giving you direct connectivity to your vnet, it can also potentially use up lot of Ips. I have seen organization where this can be a challenge especially if you intent to deploy lots and lots of containers. Kubenet on the other hand use very few of your Ips at the cost of not giving you direct connectivity. As highlighted on this screen, if you are using kubenet, there is no native way for a resource on another vnet to access your pod directly they can only respond back if Kubenet based pod initiates a connection.

There are couple of other differences in that Kubenet does not support Virtual endpoints that is bursting nodes on Azure Container instances ACIs and you cannot have multiple clusters sharing the same subnet when using Kubenet.

If you use Kubenet you will not be able to use Azure network policies. We will discuss network policies in the next video.

## ➤ Use Kubenet when:

-  Limited Vnet IP space
-  Most communication is within the cluster
-  Integrating with VNets is not a priority

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In summary, use Kubenet when you have limited VNET space and when resources outside the cluster do not need to initiate traffic to the pods running on Kubenet. Also recall Kubenet uses SNAT to communicate out of the cluster. This can have a performance overhead if the quantum of traffic going out of the cluster is large. And hence Kubenet is a slower option of the two.

## ➤ Use CNI when:

-  Need VNet integration
-  Have VNet IP space
-  Lot of communication to resources outside the cluster
-  Don't want to manage UDRs
-  Using Windows Server node pools

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

If you have enough IPs to current and future AKS clusters, it's best to use CNI. Also if you happen to use windows server node pools, CNI is the only available option.

# Network Policies

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ▶ Why use Network Policies?

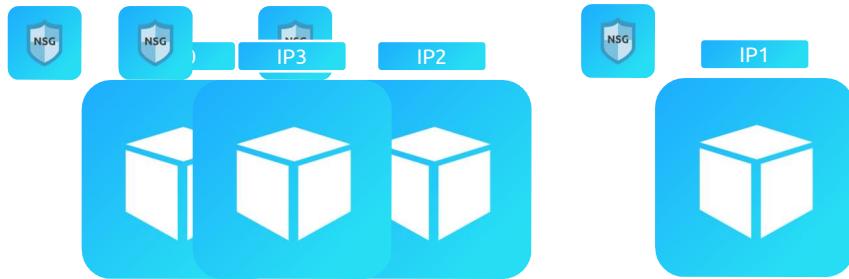


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In the previous videos, we talked about network security groups or NSGs to filter traffic that are applied to a subnet so you can filter both incoming and outgoing traffic. You can offcourse use NSGs to secure your AKS pods when using CNI. But this is not a best solution, and the reason for that is Pods are dynamic and they would come and go on the fly so if you use an IP specific NSG the utility of that will expire the moment a pod is terminated and when a new one comes up its likely to have another IP.

## ▶ Why use Network Policies?

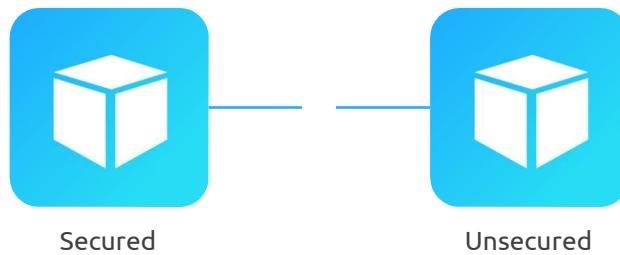


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

so if you use an IP specific NSG the utility of that will expire the moment a pod is terminated and when a new one comes up its likely to have another IP.

## ➤ Why use Network Policies?

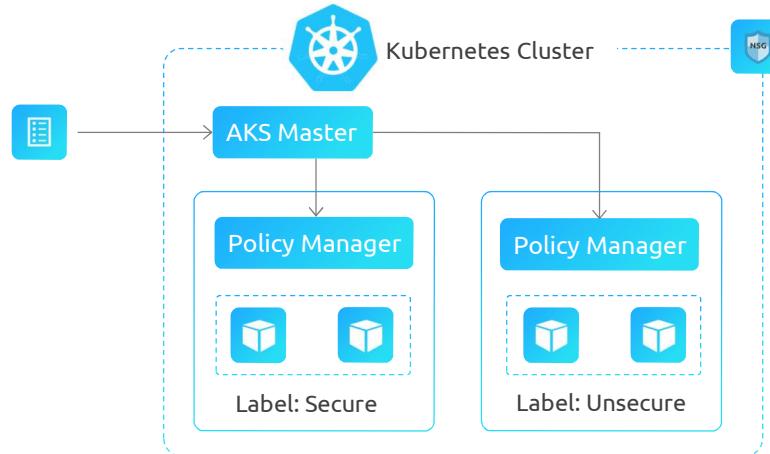


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

There will also be scenarios where you would like to define rules based on Kubernetes constructs like pod labels. For example, you might like to block traffic from pod labeled Secure to another pod labeled unsecured. That is not something an NSG can address. As you can imagine, we need a fine-grained method to define these policies.

## Kubernetes Network Policies



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

This is where Kubernetes network policies come in. This is a specification of Kubernetes that is used to filter traffic between pods and specification defines a YAML file and the YAML file has a format that lets us choose labels and help define policies around those labels.

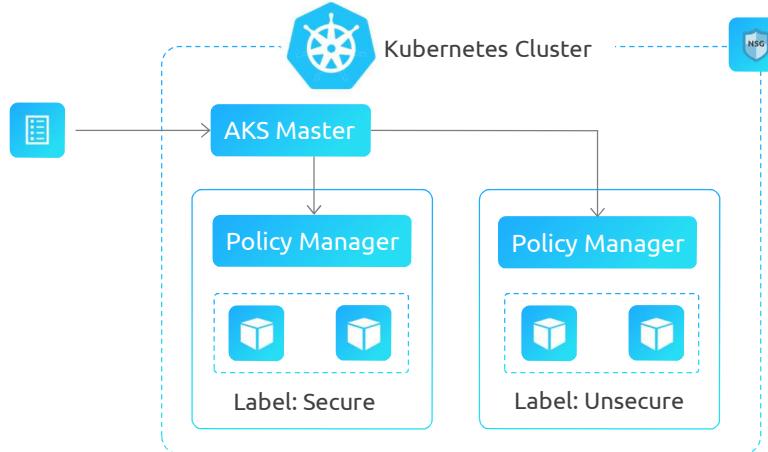
## Code

```
Terminal
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: demo-policy
  namespace: demo
spec:
  podSelector:
    matchLabels:
      role: server
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          app: frontend
  egress:
  - to: 80
  - ipBlock:
      cidr: 10.2.0.0/22

© Copyright KodeKloud Check out our full course on Azure Kubernetes Service (AKS): https://kode.wiki/300mAlg
```

We can use labels to select the type of pods and then there is a separate section in the yaml file that lets you specify what your network policies are, which are essentially filtering rules that are defined based on either labels or names or namespaces. You can also use IPBlocks as is shown in this screen.

## Azure Network Policies

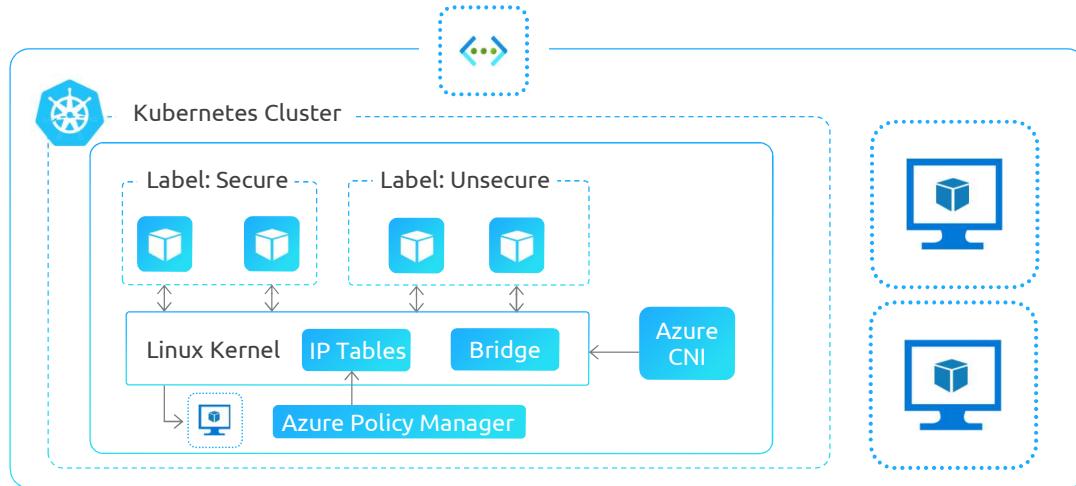


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

So Azure has deployed their own version of Network policy specification that deploys a demonset which goes and sits on every node in the cluster and anytime you go and apply a yaml file through the k8s API the policy engine which is running on the vms can intercept the file and enforce the policies defined by you. The most common way to enforce these rules is via IPTable rules or other filtering capabilities that are part of the linux . If you are using windows-based container aks leverages *Host Network Service (HNS)* supported *ACLPolicies* which is part of windows kernel

## Azure Network Policies



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

When you deploy the yaml file, Azure Network policy engine will translate the policy file into IP filtering rules provisioned into linux kernel in the form of IPTable rules. This policy engine works in collaboration with Azure CNI, which in turn depends on linux bridge functionality.

This is offcourse the default policy option in AKS

You also use this engine outside Azure if you want a consistent policy experience or if you decide to built K8S cluster in VMs on azure in DIY style.

Please note that Azure network policies in available both for linux and windows hosts but is in preview for windows. I am sure it will be generally available in due course.

## Azure vs. Calico Policies

Capability	Azure	Calico
Supported platforms	Linux, Windows 2022	Linux, Windows 2019, 2022
Supported networking options	Azure CNI	Azure CNI (Windows & Linux), KubeNet (Linux)
Compliance with Kubernetes specification	All policy types supported	All policy types supported
Additional features	None	Extended policy model consisting of Global Network Policy, Global Network Set, and Host Endpoint.
Support	Supported by Azure support and Engineering team	Calico community support. For more information on additional paid support, see Project Calico support options.
Logging	Logs available with <code>kubectl log -n kube-system &lt;network-policy-pod&gt;</code> command	Information on component logging available at Calico component logs page.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Calico network policy is another option available as first party that you can deploy on AKS clusters. Calico an open-source network and network security solution founded by [Tigera](#). You can see their website for implementation details.

A quick comparison for both Azure Network policy manager and Calico network policy are listed on your screen. If you are planning to use Kubenet then Calico would be your goto option. Also note that if you choose Calico, Microsoft support engineers will not help troubleshoot if they determine that Calico is a potential cause of any incidents.

This brings to the end of this section on networking and policies. Hope you found this useful and you can now jump on to the associated labs to get your hands dirty.



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

# Open Service Mesh

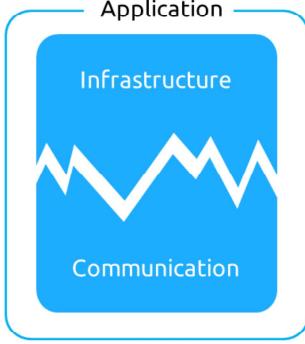
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

# Service Mesh



Service Mesh



Application

Infrastructure

Communication

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Service mesh, In Kubernetes, provides a configurable communication layer that decouples infrastructure and communication concerns from the application itself.

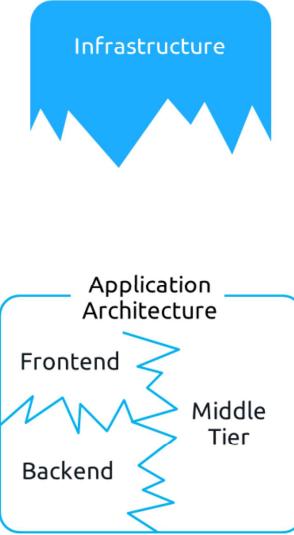
# Service Mesh

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

It achieves this by moving the intricacies of infrastructure plumbing to a separate proxy container or process, commonly known as a sidecar. This sidecar proxy is injected alongside each application service, intercepting and managing the communication between them.

# Service Mesh



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

This application service can be broken down based on your application architecture, you can call it Service, A,B or C or Frontend, Middle tier and bankend. Names does not matter here.

Importantly, the sidecar proxy remains loosely-coupled to the application service, meaning it is created together with the service, shares its lifecycle, and can be managed independently.

This architecture allows for seamless integration of essential service mesh functionalities while keeping the application codebase clean and free from communication logic.

## Service Mesh

```
graph LR; Sidecar[Sidecar] --- Application[Application]; Application --- Service[Service]
```

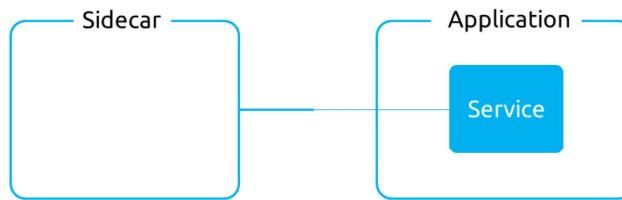
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

By adopting this sidecar approach, Kubernetes service mesh empowers developers to focus on building resilient and scalable applications, while the infrastructure layer handles the complexities of service communication.



## Service Mesh



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>



## Service Mesh



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

By adopting this sidecar approach, Kubernetes service mesh empowers developers to focus on building resilient and scalable applications, while the infrastructure layer handles the complexities of service communication.

# Service Mesh

The diagram illustrates the components of a Service Mesh. On the left, three icons represent developer tools: a person icon labeled 'Developers', a hand icon labeled 'Resilient', and a gear icon labeled 'Scalable'. To the right of these is a central graphic showing a 3D grid structure with blue and orange blocks, labeled 'Service Mesh'. Further to the right is a blue speech bubble icon containing the word 'Infrastructure', and below it is a blue wireframe sphere icon.

Developers

Resilient

Scalable

Service Mesh

Infrastructure

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Service Mesh Components

The diagram illustrates the Service Mesh Components. On the left, a large blue box labeled "Data Plane" contains three service instances: "Instance A", "Instance B", and "Instance C". Each instance is associated with a "Sidecar Proxy" component, which is shown as a dashed box around the main instance. To the right of the Data Plane, there are three icons with corresponding descriptions: 1. A circular arrow icon labeled "Handles actual routing of network traffic". 2. An up-and-down arrow icon labeled "Manages and handles traffic". 3. A double-headed arrow icon labeled "Establishes secure channels". At the bottom left of the slide, there is a copyright notice: "© Copyright KodeKloud". At the bottom right, there is a link: "Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>".

The Kubernetes service mesh consists of two essential components: the Data Plane and the Control Plane.

The Data Plane handles the actual routing of network traffic to service instances within the mesh. It performs critical functions such as service discovery, which enables services to dynamically locate and communicate with each other. The Data Plane also manages traffic, ensuring that requests are efficiently routed to the appropriate service instances. It handles traffic resiliency features like circuit breakers and retries, which help maintain service availability and reliability. Additionally, the Data Plane establishes secure channels for communication, managing identity and access management (IAM) as well as encryption to ensure secure communication between services within the mesh. Overall, the Data Plane focuses on the low-level networking and traffic management aspects, enabling seamless and secure communication between services.

# Service Mesh Components

The diagram illustrates the Service Mesh Control Plane, which is a central component of the Service Mesh. It is represented by a blue rounded rectangle containing the text "Service Mesh Control Plane". To the left of this main box is a smaller box labeled "Control Plane". Below the main box, there are three blue icons with corresponding descriptions: 1) An eye icon representing "Management and monitoring of service mesh". 2) A plus sign icon representing "Monitors the health of service instances". 3) A document icon representing "Applies application-wide policies". To the right of the main box, there is another description: "Streamlines operational aspects of Service Mesh".

Control Plane

Service Mesh Control Plane

Management and monitoring of service mesh

Monitors the health of service instances

Applies application-wide policies

Streamlines operational aspects of Service Mesh

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

On the other hand, the Control Plane is responsible for the management and monitoring of the service mesh. It takes care of provisioning new instances of services as needed, ensuring that the mesh can scale and handle increased traffic demands. The Control Plane constantly monitors the health of service instances, probing and terminating unhealthy instances to maintain overall system reliability. It also applies application-wide policies, such as rate limiting or access control, to enforce consistent behavior across the mesh. By providing centralized management and monitoring capabilities, the Control Plane streamlines the operational aspects of the service mesh, allowing operators to efficiently manage and maintain the overall health and performance of the application.

# Service Mesh Landscape

The diagram illustrates the Service Mesh Landscape with the following components:

- Istio:** Represented by a blue sailboat icon.
- Linkerd:** Represented by a green and blue striped icon.
- Open Service Mesh:** Represented by a blue and green grid-based icon.
- Microsoft Azure:** Represented by a blue 'A' icon.
- Microsoft:** Represented by the Microsoft logo (four colored squares).
- Lyft:** Represented by the Lyft logo.
- IBM:** Represented by the IBM logo.
- Google:** Represented by the Google logo.
- Twitter:** Represented by the Twitter logo.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

There are many Service mesh implementations that you can use on K8s. I'll mention three most popular ones.

Istio is widely regarded as a full-featured, customizable, and extensible service mesh. It has gained immense popularity and is backed by a collaboration between Lyft, IBM, and Google. Written in Go, Istio is designed to be application-agnostic. It offers a comprehensive set of features and provides robust capabilities for traffic management, security, observability, and more.

Linkerd, on the other hand, is known for its ease of use and lightweight nature. It is feature-rich and cross-platform, making it suitable for various environments. Linkerd builds upon early work from Twitter and is written in Scala, running on the JVM. It focuses on simplicity and ease of adoption, making it an attractive option for teams seeking a straightforward service mesh solution.

Open Service Mesh (OSM) is a cloud-native service mesh developed by Microsoft. It is designed to be a lighter-weight alternative to Istio. OSM prioritizes simplicity, aiming to provide a service mesh that is easy to understand, install, maintain, and operate. It

leverages the Service Mesh Interface (SMI) specification, allowing for standardization and easier integration with existing Kubernetes environments.

Microsoft Azure supports Open Service Mesh as its developed by them. Historically that was the default option on AKS, at the time of this recording Microsoft has started to offer native Istio service as preview. It remains to be seen how this will change their recommendation.

We will look into the OSM implementation in the rest of this video.



## Open Service Mesh Capabilities



Open Service Mesh

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

So lets look at some of the features of OSMas is implemented in Azure



## Open Service Mesh Capabilities



Open Service Mesh



Microsoft Azure

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Open Service Mesh (OSM) is a powerful service mesh solution that seamlessly integrates with Microsoft Azure, providing numerous benefits for managing microservices-based applications.

## Open Service Mesh Capabilities



Simplified operator experience through the use of Service Mesh Interface (SMI)



Leverages the CNCF-compliant Envoy proxy as lightweight sidecar containers



Provides standard service mesh features



Fine-grained Security



Use of Open-source Tools



Supports External Certificate Management

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

With OSM, Azure users experience a simplified operator experience through the use of Service Mesh Interface (SMI), which offers standardized APIs for easy mesh configuration. OSM leverages the CNCF-compliant Envoy proxy as lightweight sidecar containers, ensuring efficient and secure communication between services. It provides standard service mesh features such as traffic access control, traffic splitting, and traffic metrics. Additionally, OSM enables fine-grained security with mutual TLS (mTLS) encryption for secure service-to-service communication.

## Open Service Mesh Capabilities

The diagram illustrates the three main capabilities of Open Service Mesh (OSM). At the top left is a logo featuring a stylized blue and orange structure on a grid. To its right is a large blue square icon containing a white user profile silhouette. Below these are three smaller blue square icons: one with a white telephone handset, one with a white padlock, and one with a white eye. Below each of these three icons is a corresponding label: 'Communication', 'Security', and 'Observability' respectively. At the bottom left of the slide is a copyright notice: '© Copyright KodeKloud'. At the bottom right is a link: 'Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>'.

Monitoring and tracing capabilities are facilitated through open-source tooling, allowing users to gain insights into application behavior. Furthermore, OSM supports external certificate management, making it convenient to handle security certificates for enhanced protection. Overall, OSM's implementation on Azure empowers developers with a robust service mesh solution that streamlines communication, security, and observability within their applications.

## ➤ Demo



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

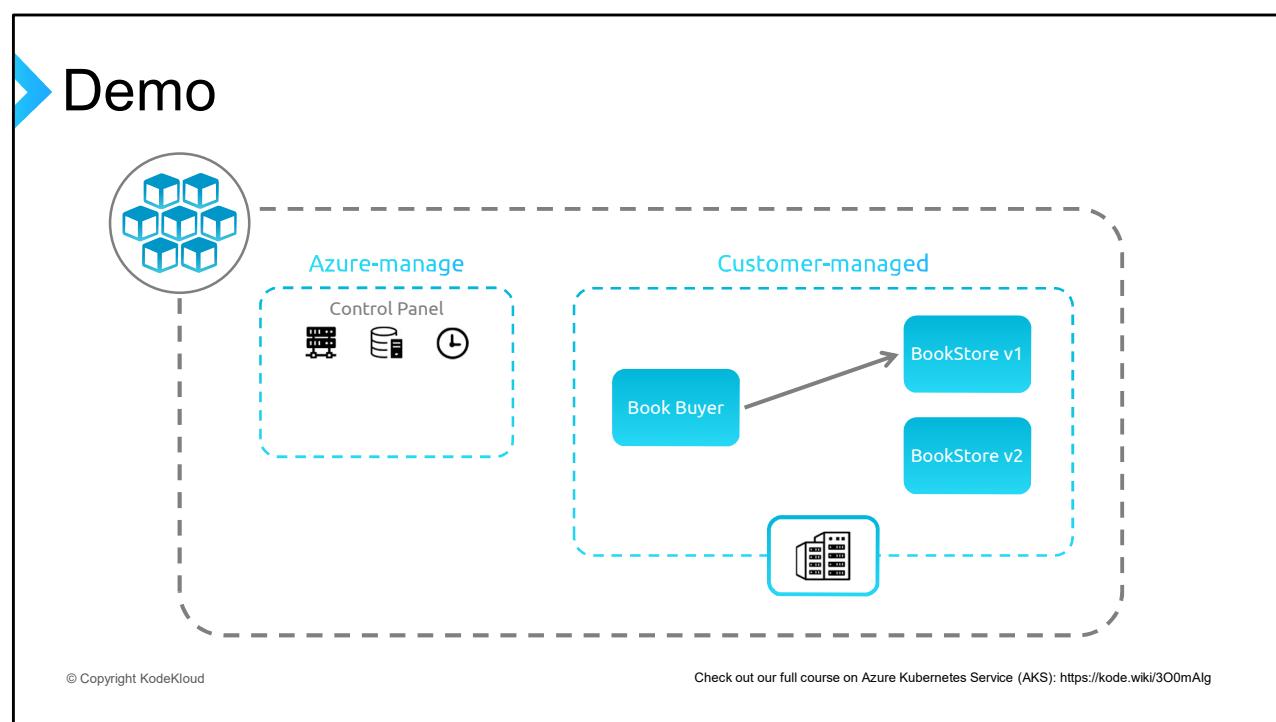
## Demo

The diagram illustrates the connection between a cluster icon (containing three blue cubes) and two management options. A dashed line connects the cluster to the 'Azure-managed' box, which contains icons for a control panel, a database, and a clock. Another dashed line connects the cluster to the 'Customer-managed' box, which contains a server icon.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In this demo, we will showcase the power of Open Service Mesh (OSM) by demonstrating its key capabilities. To begin, we'll deploy a sample bookstore application on your AKS cluster. This application comprises two services that will be automatically deployed within the cluster.



The two services included in the sample bookstore application are the book Buyer and bookstore v1. Initially, the book Buyer service will establish communication with the bookstore v1 service. Later in the demo, we will deploy bookstore v2 to demonstrate the functionality of traffic splitting within OSM.

## Demo

The diagram illustrates the two modes of managing an AKS cluster:

- Azure-managed:** Represented by a dashed blue box containing a "Control Panel" icon (server, database, clock) and the text "Permissive mode TRUE".
- Customer-managed:** Represented by a dashed grey box containing a "Book Buyer" icon pointing to "BookStore v1" and "BookStore v2", and a "Pod" icon.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Next, we will enable the OSMadd-on for AKS cluster, which will default to the permissive traffic mode. In this mode, all traffic will be allowed by default, serving as our starting point.

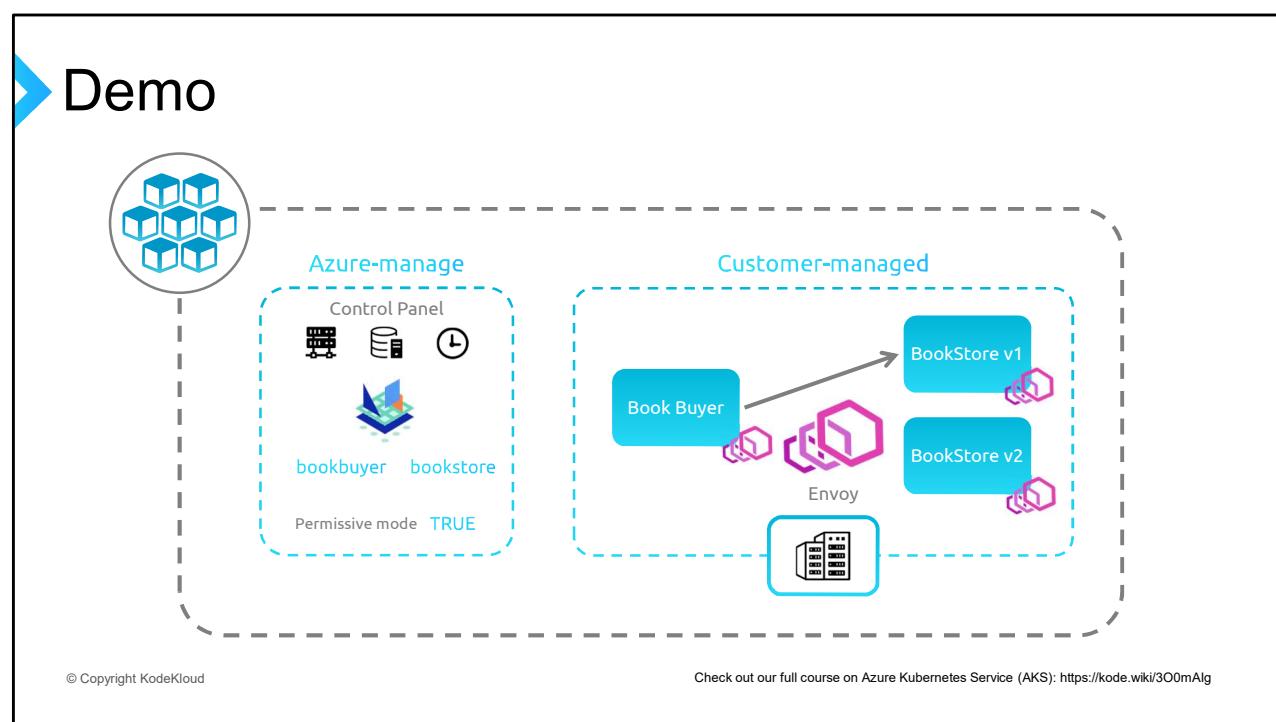
## Demo

The diagram illustrates the demo setup for managing namespaces in Azure Kubernetes Service (AKS). It is divided into two main sections by dashed lines:

- Azure-manage:** This section contains a Control Panel icon with three icons below it: a server, a database, and a clock. Below the icons are the labels "bookbuyer" and "bookstore". A text box indicates "Permissive mode TRUE".
- Customer-managed:** This section shows a "Book Buyer" service connected to two BookStore services, "v1" and "v2". A separate icon representing a Kubernetes cluster is also shown.

At the bottom left, there is a copyright notice: "© Copyright KodeKloud". At the bottom right, there is a link: "Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>".

Afterward, we will proceed to onboard the bookbuyer and bookstore namespaces to be managed by OSM. This will involve configuring OSM to handle and govern the services within these namespaces.



This process will enable the injection of Envoy sidecar proxies alongside the application containers within the pod.

## Demo

The diagram illustrates the configuration of a service named 'bookbuyer' within a Kubernetes cluster. It is divided into two main sections: 'Azure-manage' and 'Customer-managed'.

- Azure-manage:** This section contains a 'Control Panel' icon with icons for a server, database, and clock. Below it are two service icons: 'bookbuyer' and 'bookstore'. A status line indicates 'Permissive mode TRUE'.
- Customer-managed:** This section shows the 'bookbuyer' service interacting with two versions of the 'BookStore' service: 'BookStore v1' and 'BookStore v2'. An arrow points from 'Book Buyer' to 'BookStore v1', which is crossed out with a large red circle containing a diagonal slash. A separate arrow points from 'Book Buyer' to 'BookStore v2'. Below these services is a blue rectangular box representing a container or pod.

At the bottom left, there is a copyright notice: "© Copyright KodeKloud". At the bottom right, a call to action reads: "Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>".

At this stage, the book buyer service should retain its ability to communicate with the bookstore v1 service. Once we confirm this communication, we will proceed to change the permissive mode of OSM to "false." This change will result in the cessation of communication between the book buyer and bookstore v1 services.

## Demo

The diagram illustrates the deployment of traffic policies between services. It is divided into two main sections: **Azure-manage** and **Customer-managed**.

**Azure-manage:** This section contains a Control Panel icon and icons for **bookbuyer** and **bookstore**. Below these, it shows **Permissive mode TRUE**.

**Customer-managed:** This section shows a **Book Buyer** service interacting with two BookStore services, **v1** and **v2**. A green checkmark icon indicates a **Traffic Policy** is applied. A separate icon represents a **Cloud Controller Manager**.

At the bottom left, there is a copyright notice: © Copyright KodeKloud. At the bottom right, a link to a full course on Azure Kubernetes Service (AKS) is provided: Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Next, we will deploy a traffic policy to enable traffic between the services, thereby establishing communication and allowing the desired interaction.

## Demo

The diagram illustrates the deployment of a traffic split policy in OSM, divided into two main sections: **Azure-manage** and **Customer-managed**.

**Azure-manage:** This section contains a Control Panel icon and two service icons: **bookbuyer** and **bookstore**. Below these, it shows "Permissive mode TRUE".

**Customer-managed:** This section shows a **Book Buyer** service connected to two instances of **BookStore**: **v1** and **v2**. A central **Traffic Policy** box indicates a 50% split between the two stores.

At the bottom, there is a note: "Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>".

Lastly, we will deploy a traffic split policy to evenly distribute the traffic between both versions of the service. This policy will ensure a balanced distribution of requests, allowing us to showcase the capabilities of traffic splitting within OSM.

# AAD and AKS

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ▶ Identity in Azure AD



User Identity



Application Identity



Service Principal Identity



Managed Identity

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## Identity in Azure AD



User Identity

### Used for:

- Authentication
- Access Control



Application Identity



Service Principal Identity



Managed Identity

## ▶ Identity in Azure AD



Application Identity

### Enables:

- Applications to authenticate
- Access resources



# Identity in Azure AD



## Acts

As an identity of Credentials for the application to interact with Azure Services

## Allows:

- Application all services to authenticate
- Access resources in Azure

## Often Use:

### Non-interactive Scenarios

- Automations
- Service-to-service Communications



## Identity in Azure AD



Managed Identity

**Simplifies**  
authentication process for resources

**Eliminates**  
the need of managed explicit credentials



## Identity in Azure AD



Service Principal Identity



Managed Identity

Use either a client ID and client Secret or a certificate to authenticate and obtain access tokens

**Note:**

- AKS clusters that use service principal eventually expires.
- Service principal must be renewed to keep the cluster working.

© Copyright KodeKloud

Obtain tokens directly from azure AD, eliminating the need for explicit authentication details like secrets or services

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

# Kubernetes RBAC



Identities



Access Control

Providers

Users, Groups

Service Accounts

Namespace scoped Roles and RoleBindings

ClusterRoles and ClusterRoleBindings

## Kubernetes Identity Providers

### X509 Certificates



Manually Issued  
Configurable trusted CAs

### Kubernetes Service Accounts



Automatically generated by a built-in [Admission Controller](#)

### OpenID Connect (OAuth2)



AAD (and many others)

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## Sample Role



```
---  
kind: ClusterRole  
apiVersion: rbac.authorization.k8s.io/v1  
metadata:  
  name: kodekloud-clusterrole  
rules:  
- apiGroups: [apps]  
  resources: [daemonsets, deployments]  
  verbs: [get, patch]  
- apiGroups: [""]  
  resources: [configmaps]  
  verbs: [get]
```

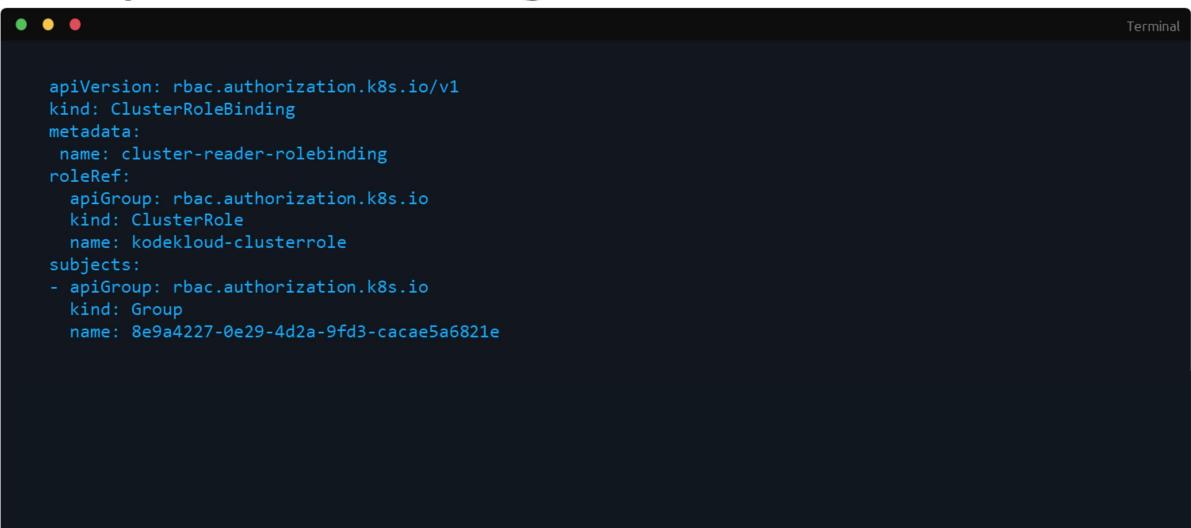
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

This how a K8s role looks like.

The role named "kodekloud-clusterrole" is Clusterwide Role that grants permissions to perform get and patch operations on daemonsets and deployments resources within the **apps** API group. It also allows get operations on configmaps resources in the core Kubernetes API group.

## Sample Role Binding



```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: cluster-reader-rolebinding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: kodekloud-clusterrole
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: Group
  name: 8e9a4227-0e29-4d2a-9fd3-cacae5a6821e
```

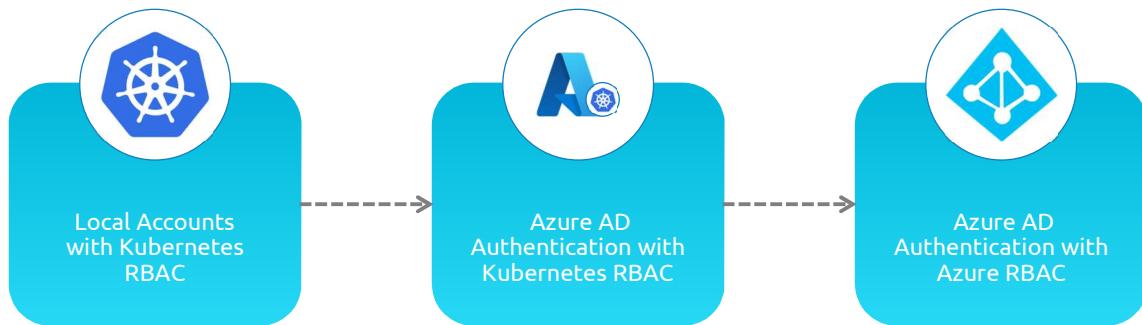
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Once the role is created we can use Rolebinding to assign that role .

In this sample, we have a clusterrolebinding that binds the ClusterRole showed in the previous slide, to a specific group of subjects identified by their unique identifier. This binding provides the subjects with the permissions and access defined by the ClusterRole shown before.

## ► Authentication and Authorization in AKS



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

# AKS Defender

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ► Microsoft Defender for Containers



Hardening



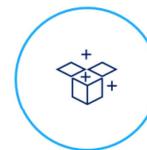
Vulnerability management



Advanced threat detection



Multi-cloud support



Deployment and configuration

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Continuously assess and improve the security posture of your containerized environments and workloads

Identify runtime threats with prioritized, container-specific alerts – using powerful insights from Microsoft Threat Intelligence

Single container security solution for Kubernetes clusters, across Azure, AWS, GCP and on-premise

Hybrid security approach with agentless capabilities

Frictionless at scale deployment with easy onboarding and support for standard Kubernetes monitoring tools

## Native integration in AKS



Defender Profile



Azure Policy add-on

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

What do we need for the defender plan, we are looking at 2 key components:

- The defender profile
- The azure policy add on

## ➤ The Defender Profile

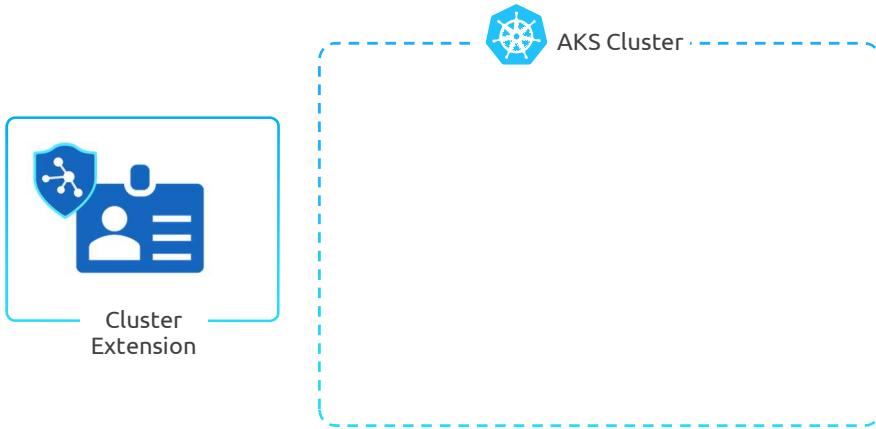


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

You will see both of these components are natively integrated to Azure Kubernetes services.

## ► The Defender Profile



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

- The **defender profile**, when you are working in Azure it is integrated as azure security profile, and when you are working outside of azure this is integrated as a cluster extension.

## ➤ The Defender Profile



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

This allows us to:

- Provide run time detection
- And get host level visibility to the cluster
- The defender profile as you can see is deployed for each worker node and with the defender profile, we collecting security data events and inventory and sending it to the defender for cloud backend.

## ➤ Azure Policy Add-on



Azure Policy add-on

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

- The second component is the **azure policy add on**, this component is extending the gate keeper, to apply Kubernetes data plane policies enforcements and hardening capabilities, essentially every API request to the Kubernetes API server is been monitoring against best practices, we use this for different recommendations you can see in portal.

## ➤ Azure Policy Add-on



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

- To enable friction less deployment, The deployment options are:
  - From the MDC portal – auto provisioning
  - From the MDC portal – recommendations quick fix
  - With Rest API
  - From Arm (azure resource manager)
  - From CLI

## ➤ Azure Policy Add-on



Policy  
Enforcement

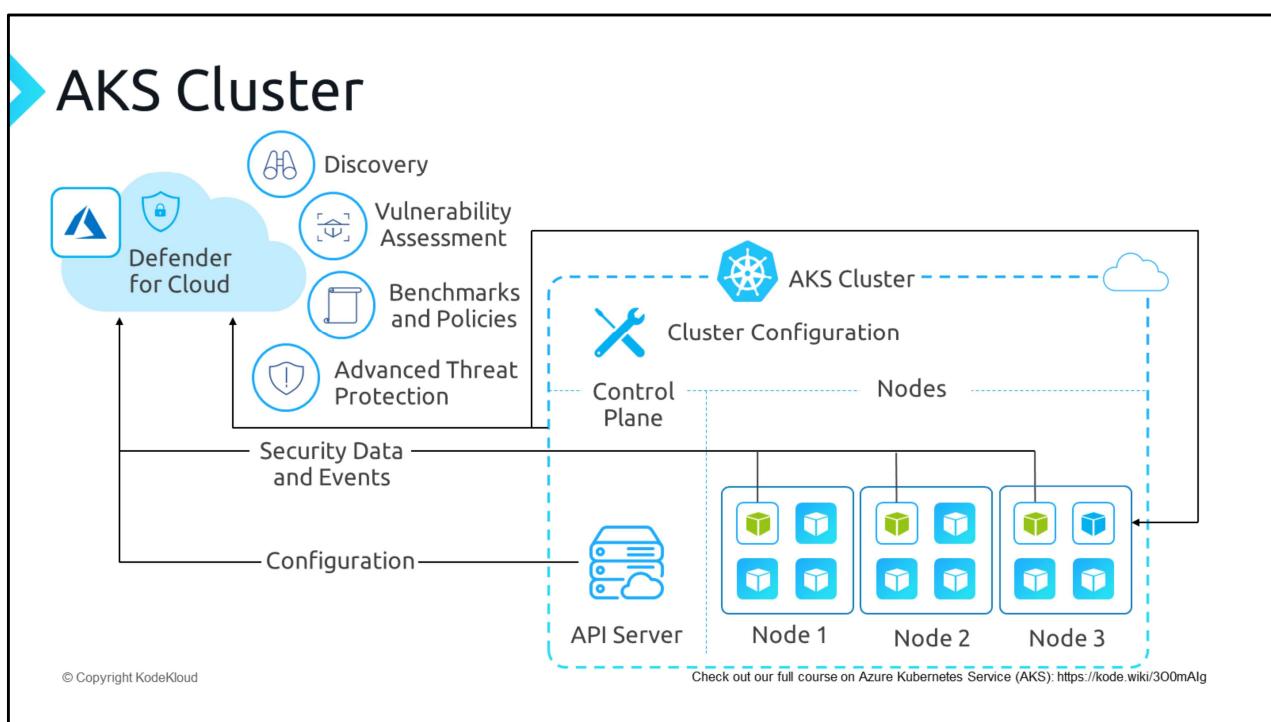


Hardening  
Capabilities



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>



## The Defender Profile

Pod Name	Namespace	Kind	Description

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

As we said the defender profile is deployed as a daemon set, so it gets to be deployed on every single node in the cluster, as a pod.

## Azure Container Registry(ACR) Scanning

The screenshot shows the Azure Container Registry (ACR) scanning interface. At the top, there's a table with columns for 'Name ↑', 'Max score ↑↓', 'Current score ↑↓', and 'Potential score increase ↑↓'. A dropdown menu 'Remediate vulnerabilities' is open, showing a score of 6, a current score of 0.56 (represented by a progress bar), and a potential increase of +8%. Below the table, two remediation steps are listed: 'Container registry images should have vulnerability findings resolved' and 'Running container images should have vulnerability findings resolved'.

Below the interface, four circular icons represent different triggers for image scans:

- On Push:** An icon showing a single image with an arrow pointing to it.
- Recently Pulled:** An icon showing a single image with an arrow pointing away from it.
- On Import:** An icon showing three stacked images with a downward arrow.
- Continuous Scan:** An icon showing three stacked images.

On the left side of the slide, there is a copyright notice: "© Copyright KodeKloud". On the right side, there is a link to a full course: "Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>".

Triggers for an image scan:

On push: whenever an image is pushed to the registry

Recently pulled: any image that been pulled within the last 30 days is scanned on a weekly basis

On import: when an image is imported into ACR

Continous scan

Every 7 days after an image is pulled. Does not require the security profile

Continous scan for running images. Every 7 days as long as the image runs.

Runs instead of the above when Defender profile is enabled

# Advanced threat detection

The screenshot shows a cloud-based security platform interface. On the left, a dashboard for a Docker image named '2e7c9245e5fd' displays the following information:

- Image:** 2e7c9245e5fd
- Total vulnerabilities:** 66
- Vulnerabilities by severity:**
  - High: 8
  - Medium: 58
  - Low: 0
- CI/CD Scan Findings:** Image Not Scanned

Below the dashboard, there's a table of findings:

Severity	ID	Security Check	Category	Patch Available	Package type	Pack
High	178369	Debian Security Update for tzd...	Debian	Yes	N/A	N/A
High	177998	Debian Security Update for mer...	Debian	Yes	N/A	N/A
High	372268	GNU Bash Privilege Escalation ...	Local	No	N/A	N/A
High	178391	Debian Security Update Multipl...	Debian	Yes	N/A	N/A
High	178944	Debian Security Update for lib...	Debian	Yes	N/A	N/A
High	178867	Debian Security Update for elf...	Debian	Yes	N/A	N/A
Medium	178601	Debian Security Update for sub...	Debian	Yes	N/A	N/A

On the right, a detailed view of a specific vulnerability is shown:

**178369-Debian Security Update for tzdata (...**

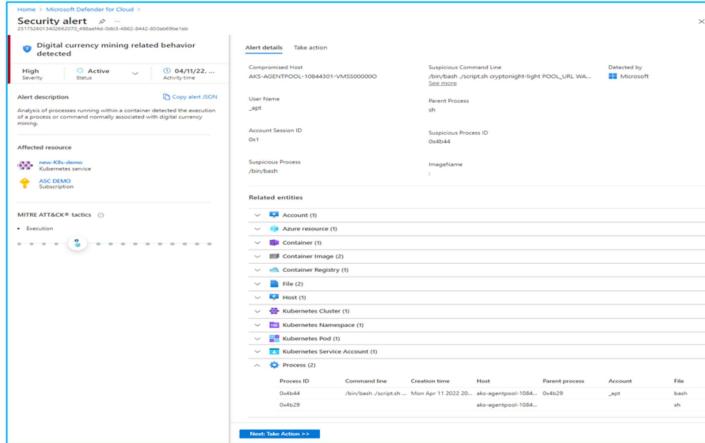
- Description:** Debian has released security update for tzdata to fix the vulnerabilities.
- General information:**
  - ID: 178369
  - Severity: High
  - Type: Vulnerability
  - Published: 1/28/2021, 2:47 PM GMT+2
  - Patchable: Yes
  - Cvss 3.0 base score: 5.3
- Remediation:** Refer to Debian security advisory DLA 2424-1 for patching details.
- Patch:** Following are links for downloading patches to fix the vulnerabilities:  
DLA 2424-1:debian
- Additional information:**
- Affected resources:**

Digest	OS	Repository	Registry
fe9f8e9122d9	Linux	dotnet/core/sdk	imageSci
2e7c9245e5fd	Linux	dotnet/core/sdk	ascdemo

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

# Advanced threat detection



The screenshot shows a Microsoft Defender for Cloud security alert for "Digital currency mining related behavior detected". The alert details section includes:

- Compressed Host:** AKS-AGENTPOOL-10844301-VMSS00000000
- User Name:** \_JET
- Suspicious Command Line:** /bin/bash ./script.sh crypotight-light POOL\_URL\_Wall...
- Detected by:** Microsoft
- Parent Process:** sh
- Suspicious Process ID:** 0x4044
- ImageName:** /bin/bash

**Related entities:**

- Account (1)
- Azure resource (1)
- Container (1)
- Container Image (2)
- Container Registry (1)
- File (2)
- Host (1)
- Kubernetes Cluster (1)
- Kubernetes Namespace (1)
- Kubernetes Pod (1)
- Kubernetes Service Account (1)
- Process (2)

Process ID	Command line	Creation time	Host	Parent process	Account	File
0x4044	/bin/bash ./script.sh ...	Mon Apr 11 2022 20...	aks-agentpool-1084...	0x4029	_JET	sh
0x4029						

**Next Take Action**

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## ➤ Advanced threat detection



Rich Detection Suite



Leading threat intelligence



Understand Risk and Context



Automate Response

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

### Rich detection suite

- Control plane and workload level detections
- Deterministic, AI, and anomaly-based alerts to identify threats

### Leading threat intelligence

- Microsoft's global threat intelligence with honeypot networks, research malware feeds, in addition to memory forensic techniques to identify fileless attacks

### Understand risk and context

- Prioritized alerts mapped to MITRE ATT&CK® tactics to easily understand the Kubernetes context effect across the attack lifecycle and to identify response action

### Automate response

- Automate actions with tools of your choice: SIEM

integration, email notifications, workflow automations and continuous export

# Azure Policy for AKS

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## ➤ Azure Policy



Business Critical  
Application

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Before we talk about Azure Policy for AKS, lets quickly see what is Azure Policy and why and how do you use it.

## Azure Policy

- Data Residency
- Sovereignty Requirements

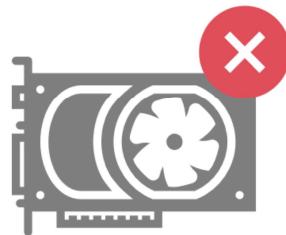


© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

As an organization, you will have many requirements or intents that you would like to enforce or atleast report against. For example, lets say you a business critical application and you would like to deploy your resources for that application only in Australia because of data residency and sovereignty requirements. Lets take another scenario where you would like to make sure that none of the AKS deployments use GPU based machines in order to control costs.

## Azure Policy



GPU-based Machines



Cost

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

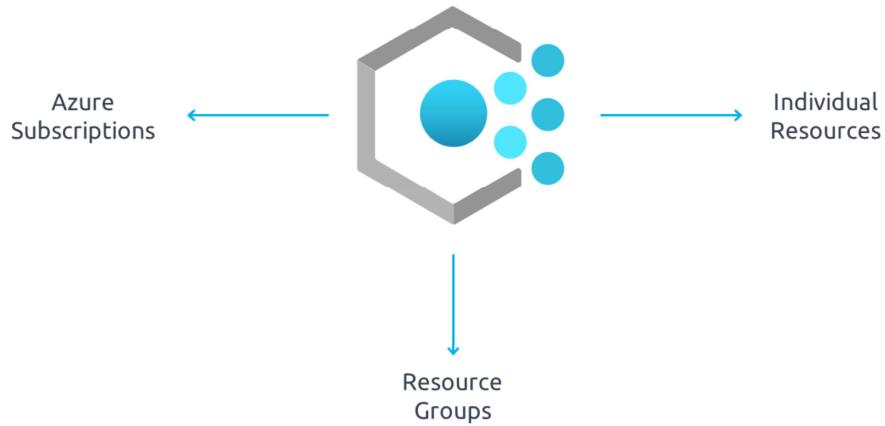
Azure Policy is a service that enables you to define and enforce governance and compliance rules for your resources. It provides a centralized and consistent way to apply and enforce policies across Azure subscriptions, resource groups, or individual resources. With Azure Policy, you can establish rules that govern resource configurations, access controls, and adherence to specific compliance requirements. These policies are defined using JSON-based policy definitions and can be customized to align with your organization's unique needs. By implementing Azure Policy, you can ensure that your Azure environment remains secure, compliant, and well-governed, promoting best practices and minimizing risks associated with misconfigurations or non-compliant resources.

## ➤ Azure Policy



A service that enables you to define  
and enforce governance and  
compliance rules for your resources

## Azure Policy



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Azure Policy

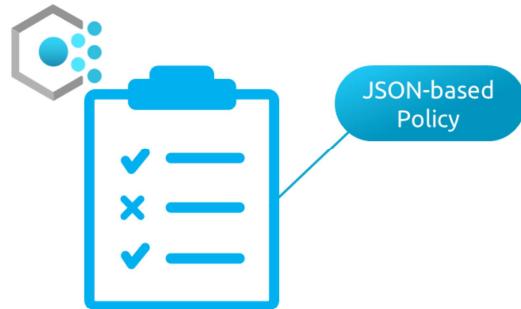


- Resource Configurations
- Access Controls
- Adherence to specific compliance requirements

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Azure Policy



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

## Azure Policy

The diagram features a central blue hexagonal icon with a white gear and three smaller teal circles. This is set against a large light blue circle labeled 'Azure Environment'. To the left of the circle is a teal square icon containing a white shield with a checkmark, labeled 'Secure'. To the right is another teal square icon containing a white clipboard with a checklist, labeled 'Compliant'. At the bottom is a teal square icon containing a white hand holding a gear, labeled 'Well-governed'.

Secure      Compliant      Well-governed

Azure Environment

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

© Copyright KodeKloud

Before we talk about Azure Policy for AKS, lets quickly see what is Azure Policy and why and how do you use it.

as an organization, you will have many requirements or intents that you would like to enforce or atleast report against. For example, lets say you a business critical application and you would like to deploy your resources for that application only in Australia because of data residency and sovereignty requirements. Lets take another scenario where you would like to make sure that none of the AKS deployments use GPU based machines in order to control costs.

Azure Policy is a service that enables you to define and enforce governance and compliance rules for your resources. It provides a centralized and consistent way to apply and enforce policies across Azure subscriptions, resource groups, or individual resources. With Azure Policy, you can establish rules that govern resource configurations, access controls, and adherence to specific compliance requirements. These policies are defined using JSON-based policy definitions and can be customized to align with your organization's unique needs. By implementing Azure Policy, you can ensure that your Azure environment remains secure, compliant, and well-governed,

promoting best practices and minimizing risks associated with misconfigurations or non-compliant resources.



## Demo

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Lets jump on Azure portal and deploy a built-in policy to make sure that none of the AKS admins can deploy resources outside Australia.

On the Azure portal, lets open Policy blade. Under Authoring, I'll go to Definitions. Here you can see 100s of pre-defined definitions that you can use out of box. There are two definition types – Policy and Initiative. Initiative is a bucket with multiple policies inside it. It is useful where you have a bunch of related policies. For example I can create a KodeKloud AKS policy initiate with all AKS polices clubbed together . As you can see there are a number of pre-defined initiates like this PCI DSS ones which I can use in case I am deploying a workload that need to report against PCI DSS standards.

Coming back to the task at hand, I'll looking for a location based policy, so lets search with the keyword location.. And there.. Allowed locations policy is what we need. Lets click on Assign. The scope is the level at which this policy will be assigned. It can either be at a subscription level or you select a resource group under a subscription. I'll select just my subscription. You can optionally change the Assignment Name and add any description. I'll click Next. Here I'll select regions that will be in my allowed list. Any location outside this list will be blocked.

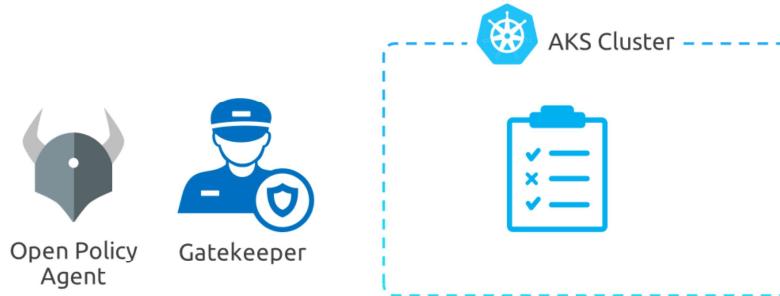
In the remediation tab, you can create a managed identity to move any resources created before this policy to an allowed region. If you need to apply this to your environment, be mindful of any impact this might cause. I'll click next. Lets add a message which admins will see if they run a non-compliant deployment.

Message from Kode Kloud Admin: You are allowed to deploy to only Australian regions for this subscription

Then lets click create.

In order to validate if our policy works as expected, lets try to create a new AKS cluster. I'll give the cluster a name and choose a location outside Australia. Since we do not need to make any other changes, click Review and create. Azure resource manager engine runs the validation to determine if the deployment will succeed as as expected the validation has failed. If we look at the details, we can see our super helpful message that we had added in the policy. Now that you some background on Azure policies, lets see how Policies work in AKS

## ▶ Open Policy Agent



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Native Azure policy does not work directly on AKS Clusters. This is where Open Policy Agent or OPA and Gatekeeper come into picture. OPA is an open-source policy engine that is used for policy enforcement

Gatekeeper extends the capabilities of OPA by providing a Kubernetes-native integration for policy enforcement.

Gatekeeper leverages the OPA engine to evaluate policies and enforce them as admission control policies in Kubernetes clusters. It integrates with the Kubernetes API server and acts as a validating webhook, intercepting and evaluating requests made to the cluster. By using Gatekeeper, you can define and enforce custom policies to validate and control various aspects of Kubernetes resources, such as deployments, pods, services, and more.

## ► Open Policy Agent



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Native Azure policy does not work directly on AKS Clusters. This is where Open Policy Agent or OPA and Gatekeeper come into picture. OPA is an open-source policy engine that is used for policy enforcement

## ▶ Open Policy Agent



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Gatekeeper extends the capabilities of OPA by providing a Kubernetes-native integration for policy enforcement.

Gatekeeper leverages the OPA engine to evaluate policies and enforce them as admission control policies in Kubernetes clusters. It integrates with the Kubernetes API server and acts as a validating webhook, intercepting and evaluating requests made to the cluster. By using Gatekeeper, you can define and enforce custom policies to validate and control various aspects of Kubernetes resources, such as deployments, pods, services, and more.



## Demo

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

So lets see the Azure Policy for AKS in action . I have created a brand new AKS cluster with default options and I have called it AKSPolicyDemo. If I click on the Policies blade in the cluster, you can see that I have an option to enable the add-on from here. This also indicated at the moment its disabled. Lets also verify this using command line.

Lets run Az aks show and query the addonprofiles

```
Az aks show -g AKSPolicyDemo_group -n AKSPolicyDemo --query  
addonProfiles.azurepolicy
```

And this shows that the azure policy addon is currently disabled.

So lets enable Azure policy for aks via command line using az aks enable-addons comand

```
az aks enable-addons --addons azure-policy --n AKSPolicyDemo --g  
AKSPolicyDemo
```

This will install `azure-policy` pod is in `kube-system` namespace and `gatekeeper` pod is installed in `gatekeeper-system` namespace

Lets quickly validate that too by running

```
kubectl get pods -n kube-system | grep policy
```

Here you can see the two components running – one is Azure policy and another one the associated webhook. These are orchestrators that sit behind the Azure policy addon that we just enabled. These pod synchronize the policy that we deploy from Azure policy to Kubernetes. It is worth noting that the add-on checks in with Azure Policy service for changes in policy assignments every 15 minutes.

And

```
kubectl get pods -n gatekeeper-system
```

The `gatekeeper` is an opensource component that runs inside the cluster to do the enforcement of the policies. So while we assign the policies via `azure policy` it is the `gatekeeper` inside AKS that is doing the hard work of making sure that those policies are applied inside aks.

Now that Policy add-on is installed, lets go to Azure policy and apply a simple policy to ensure that AKS can only be configured to listen to allowed ports.

I'll go to Policy, then click on Definitions

I'll then search for Search for Kubernetes Policy Definitions.

I have selected Policy Definition **Kubernetes cluster services should listen only on allowed ports** which allows us to restrict ports to be used for Cluster Services. You have been this process before, I'll Assign the Policy and complete the necessary fields, allowing only port 80 and 443.

Now lets this deploy this simple yaml file to our AKS cluster. Note that this will deploy nginx and use TCP port 80

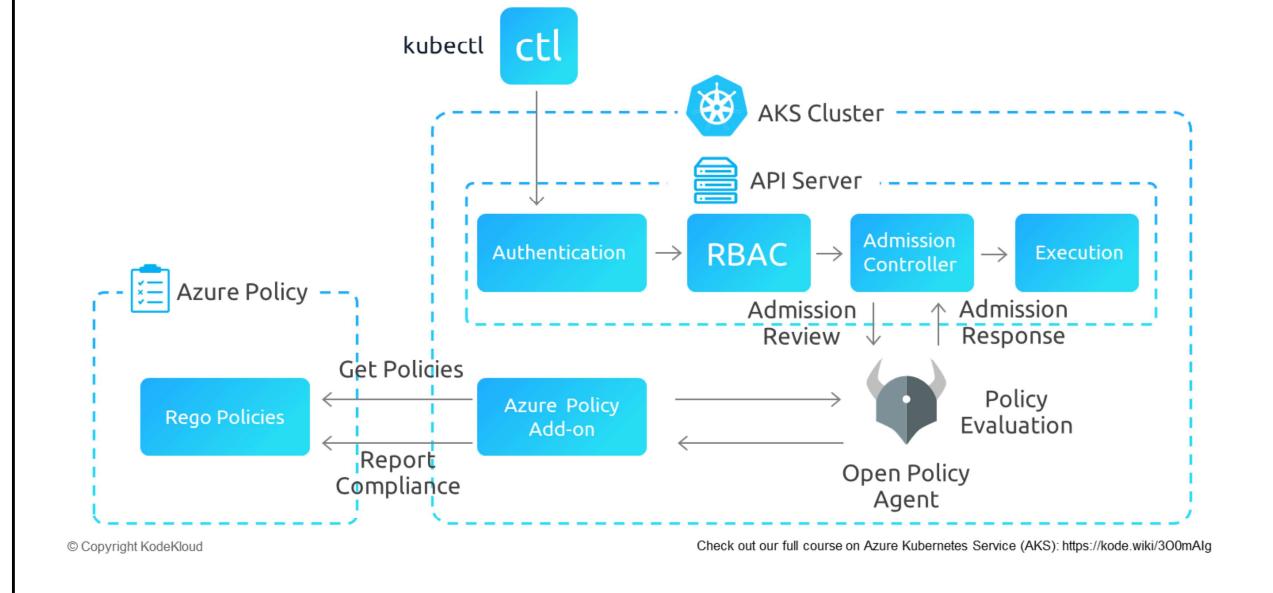
Do you think this gatekeeper will allow this deployment?

Yes offcourse, we have allowed ports 80 and 443

Lets now try to deploy ngnix but this time lets use port 8080 and see the result.

As you can see the deployment has failed with the custom error message we have set when we deployed the policy.

## AKS and OPA/Gatekeeper



Let's take a look at how these components work together. The Azure managed Kubernetes control plane uses OPA Gatekeeper to enforce policies on behalf of Azure Policy.

Inside the Kubernetes (K8s) cluster, we have the API server. We can use kubectl to interact with this component by sending requests to it. When we do that, the request goes through a pipeline of components. First, it goes through authentication, followed by authorization or RBAC. Next, it passes through optional admission controllers, which can be multiple in the cluster. Only after these three stages have passed, the command is executed.

Azure Policy works alongside Gatekeeper and integrates with the Admission Controller module. On the left side, we have Azure Policy. When we enable the policy add-on, Azure Policy and Gatekeeper are deployed in the cluster. When we assign a policy or initiative to a subscription or resource group where the AKS cluster is hosted, the Azure Policy add-on synchronizes the policy details from Azure Policy into Gatekeeper. Gatekeeper is configured to integrate with the admission controller. So, when a request comes into the API server, the admission controller sends the request to Gatekeeper. Gatekeeper then validates if the request complies with the synchronized policies. It sends

the response back to the admission controller, and based on the compliance status, the request will proceed accordingly.

This brings us to the end of the video on Azure policy for AKS.

## ► Azure Kubernetes Services (AKS) - Summary

-  Azure Networking
-  Virtual Networks, Subnets, Network Security Groups
-  Open Service Mesh
-  IAM in AKS
-  Azure Defender
-  Azure Policy

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Ok, we have now reached the culmination of this video course on Azure Kubernetes Service (AKS). Our journey began by delving into the fundamental aspects of Azure, encompassing essential concepts such as Compute, Storage, and Networking. With a some foundation in place, we transitioned towards the creation and containerization of a basic C# web application. Throughout that module, we focused on leveraging Docker desktop and crafting Dockerfiles to facilitate containerization.

Subsequently, we navigated to the Azure portal to establish an AKS cluster to deploy our world class application onto that cluster. Harnessing the power of kubectl, we effectively executed various imperative actions, including scaling of deployments. we explored the deployment of container images within the Azure Container Registry (ACR), unearthing its potential for application enhancement and version management. By harnessing the image hosted on ACR, we facilitated a seamless application upgrade, only to effortlessly revert to a previous version when needed. Concluding this module, we bestowed upon you a overview of the Kubernetes Fleet, empowering you with a holistic understanding of this powerful toolset.

The subsequent module, entirely dedicated to AKS networking options and policies,

showcased the indispensable importance of robust network management within AKS deployments. Furthermore, we delved into security-centric topics, delving into the intricacies of Azure AD integration, defender for AKS, and Azure policy for AKS. Empowered with this knowledge, you are poised to fortify your AKS infrastructure with optimal security measures and governance practices.

As we progressed, we ventured into the realm of CI/CD patterns within AKS, looking at the art of managing AKS using declarative techniques. By adopting these practices, you can streamline the continuous integration and delivery pipelines within your AKS environment.

An indispensable component of any infrastructure, observability received some attention in module seven. We explored observability within AKS, equipping you with the necessary tools and techniques to monitor and gain valuable insights into your AKS deployments.

Finally, before we brought this course to a close, we presented you with an overview of alternative platforms that embrace the power of containers. By expanding your horizons and exploring these platforms, you can unlock a world of possibilities and further enhance your containerization endeavors.



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

# CI/CD Workflow for AKS

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## Kubernetes Deployment

 Imperative	 Declarative
 Specific commands to create and manage Kubernetes Objects	 Defining the desired state in a YAML or JSON File called a "Manifest"
 Each step and action are defined	 Treats manifest as the source of truth, reconciling actual state with the desired state.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In Kubernetes, there are two primary ways to deploy applications: Imperative and Declarative.

Imperative deployment involves issuing specific commands to create and manage Kubernetes objects. With imperative commands, you define each step and action required to deploy and configure your application. For example, in module 4 we used imperative commands via kubectl run to create a deployment, then use kubectl expose to create a service, and so on.

## Kubernetes Deployment

➤ Imperative

Pros	Cons
 More fine-grained control	 More error-prone
 Allows immediate execution of commands	 Lack idempotency
 More flexible	
 Useful for ad-hoc or one-time operations	

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Now this is great for learning AKS and to get started with Kubernetes. In most environments, you will get exposed to Declarative way of deploying your applications on AKS.

Declarative deployment involves defining the desired state of your application in a YAML or JSON file called a manifest. You specify the desired configuration and characteristics of your application, and Kubernetes takes care of managing the resources to achieve that desired state. This approach treats the manifest as the source of truth, allowing Kubernetes to reconcile the actual state with the desired state.

## Kubernetes Deployment

 Declarative

Pros	Cons
 Promotes desired-state configuration management approach	 Requires careful attention and proper management
 Provides clear, version-controlled representation	
 Easier to track changes and collaborate with teams	
 Reusable, shareable, and can be stored in source control, enabling reproducible deployments	
 Allows easy scaling and updating	

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Each approach has its pros and cons.

Imperative deployment offers a more fine-grained control over the deployment process. It allows for immediate execution of commands and provides more flexibility in managing individual resources. It is useful for ad-hoc or one-time operations, such as troubleshooting or quick prototyping. However, imperative commands can be more error-prone, especially in complex deployments, and they lack idempotency, meaning that re-running the same command may result in different outcomes.

Declarative deployment, on the other hand, promotes a desired-state configuration management approach. It provides a clear, version-controlled representation of the application's configuration and makes it easier to track changes and collaborate with teams. Declarative manifests are reusable, shareable, and can be stored in source control, enabling reproducible deployments. It also allows for easy scaling and updating of deployments. However, declarative deployment requires careful attention to the manifest file and proper management of changes to avoid unintended consequences.

## Kubernetes Deployment

 Imperative	 Declarative
 Offers more control and flexibility for specific tasks	 Provides consistency, versioning and reproducibility

© Copyright KodeKloud

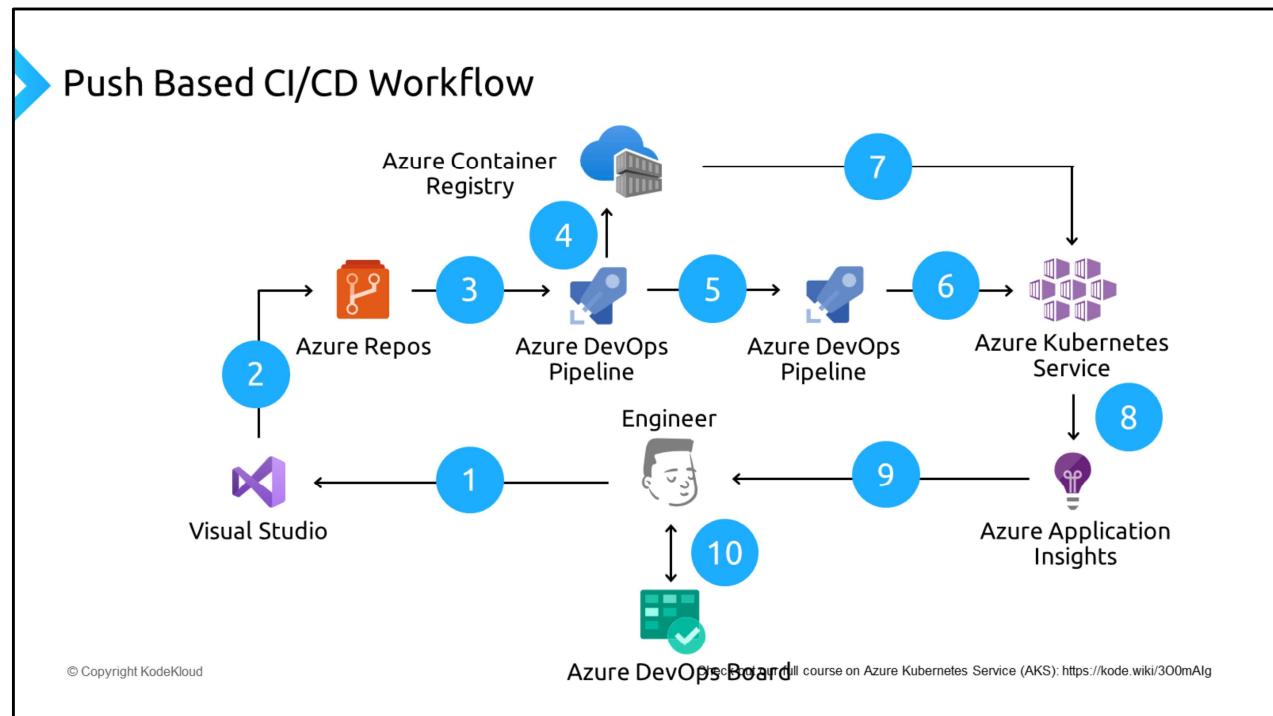
Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

In summary, imperative deployment offers more control and flexibility for specific tasks, while declarative deployment provides consistency, versioning, and reproducibility benefits. The choice between them depends on the specific use case, the complexity of the deployment, and the desired level of control and automation. In practice, a combination of both approaches is often used, leveraging the strengths of each for different stages of the application lifecycle.

# Push-based CI/CD Workflow

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>



While, we have taken shortcuts, a best practice deployment using Microsoft DevOps toolset wil look like this

1. Create the application source code, if you recall we created an ASP.net based application in Module 3.
2. We then Commit Application Code to azure repos
3. Continuous integration then triggers application build, container image is build and any unit tests that you might have configured will run. Because you are an awesome dev, this unit test will all succeed
4. In Step 4, Azure DevOps pipeline will push the Container image to Azure Container Registry.
5. Continuous deployment trigger orchestrates deployment of application artifacts with environment-specific parameters.
6. In Step 6, DevOps pipeline will Deployment to Azure Kubernetes Service

(AKS) cluster

7. Container is launched using Container Image from Azure Container Registry that we had saved in step 4
8. Optionally, Application Insights can collect and analyze health, performance, and usage data.
9. As an engineer, you can then Review health, performance, and usage information.
10. Finally, you can Update any backlog items.

You can offcourse, use any devops toolset, I have used Azure devops as an example. Now that you have seen a sample push based pipeline, lets see how the pull based pipeline is different.

Note Visio file for this image is : <https://arch-center.azureedge.net/cicd-for-containers.vsdx>

# Pull-based CI/CD Workflow

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

## What is GitOps?

```
graph TD; GIT[GIT] --- IDE[IDE]; GIT --- Build[Build]; GIT --- Test[Test]; subgraph CI [Continuous Integration]; IDE; Build; Test; end; subgraph K8sGitOps [Kubernetes GitOps]; Deployment["Deployment<br>(clusters, apps)"]; ML["Monitoring Logging<br>(Observability)"]; Management["Management<br>(operations)"]; end; Deployment --- K8sGitOps; ML --- K8sGitOps; Management --- K8sGitOps;
```

“ Immutability Firewall ”

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Before we discuss, pull based CICD, we need to understand what is GitOps. So here is the basic concept.

## What is GitOps?

Operational Framework



Kubernetes Cluster Management



Application delivery

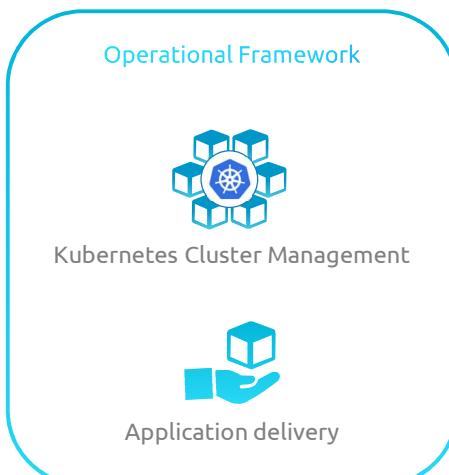
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

GitOps is a methodology for managing and automating application deployment. It involves storing application configurations and infrastructure as code in a Git repository.

## What is GitOps?

Operational Framework



Kubernetes Cluster Management

Application delivery

### Applies

Version Control

Collaboration

Compliance

CI/CD to Infrastructure

Everything-as-a-code

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Changes to the desired state are made through Git commits, as you would in Push based method as well but in this method the git commits triggers automated synchronization with the Kubernetes cluster using gitops tools such as Flux or Argo CD to ensure that the cluster state matches the desired state, enabling easy auditing and rollbacks.

## What is GitOps?

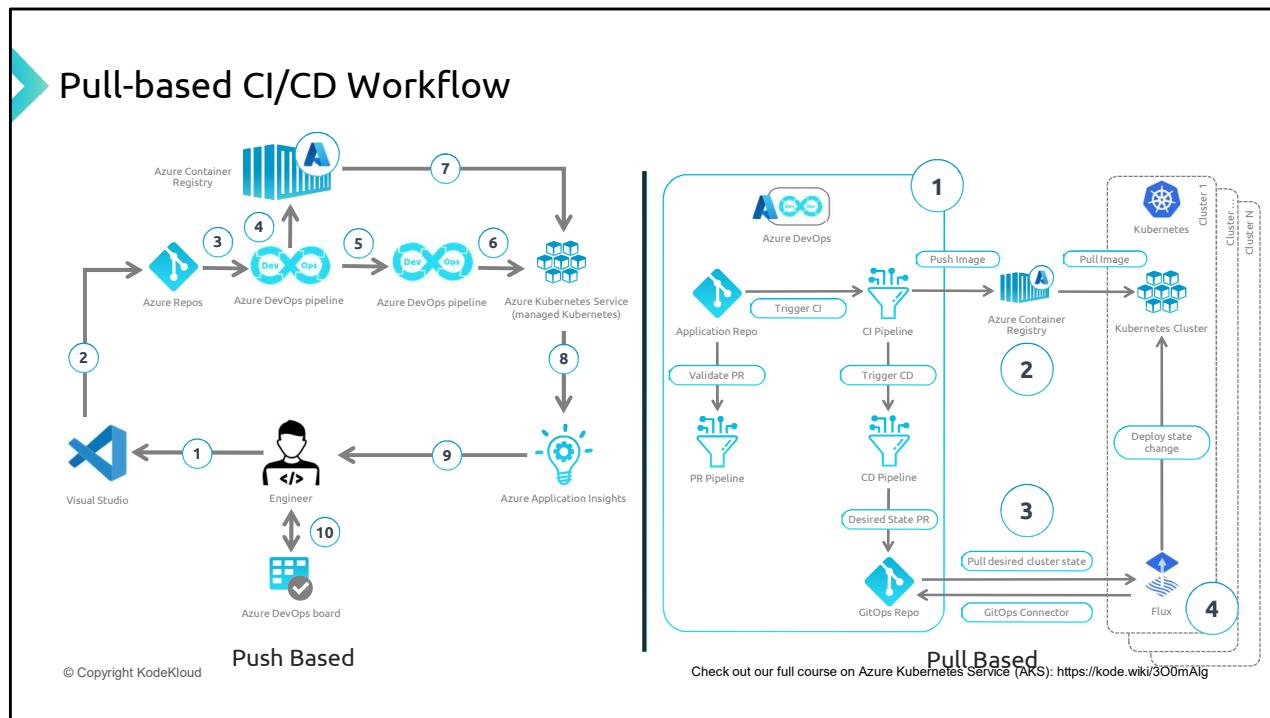
Enabled by **automation** and **sync** tools like **FluxCD** and **ArgoCD**

The diagram shows a central rounded rectangle labeled "Operational Framework". Inside, there are two sections: "Kubernetes Cluster Management" with a gear icon and "Application delivery" with a hand holding a cube icon.

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

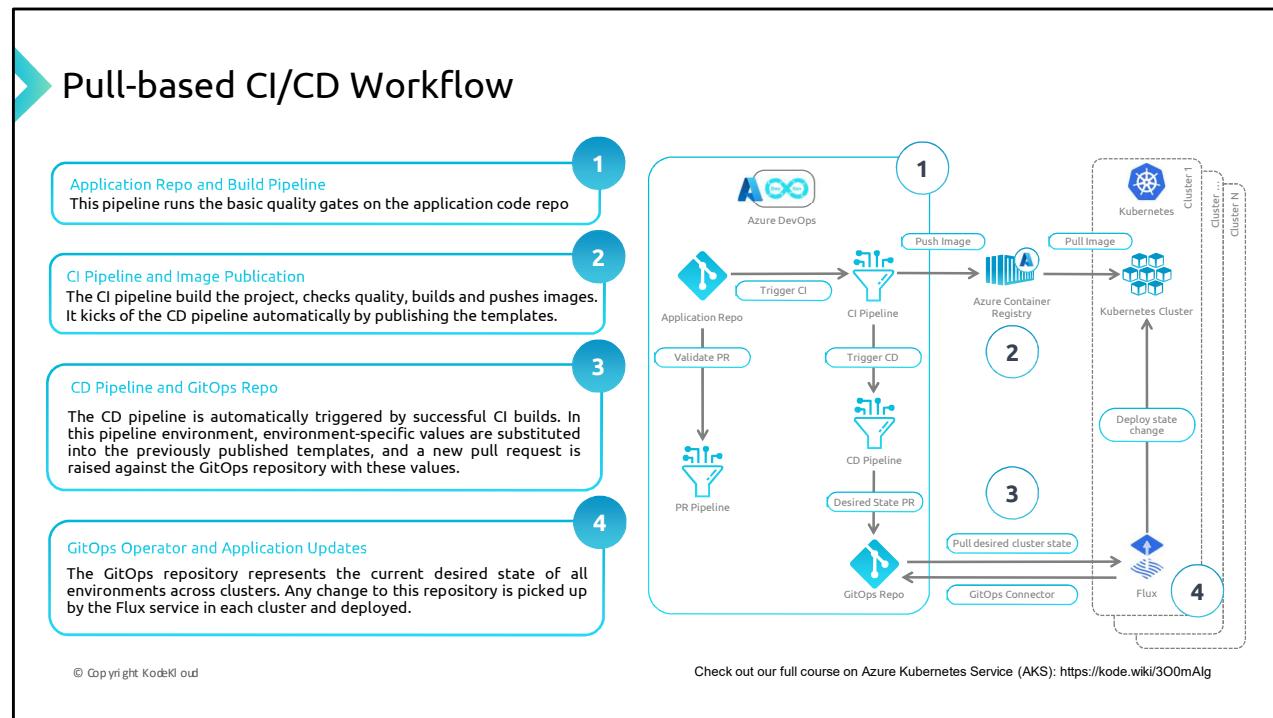
It promotes collaboration, continuous integration, and provides a reliable, version-controlled approach to managing Kubernetes applications



Lets now see who this differs to the Push based workflow we saw before. Two GitOps operators that you can use with AKS are [Flux](#) and [Argo CD](#). They're both [\(CNCF\)](#) projects and are widely used.

While the outcome remains the same, there are couple of differences I want to highlight here. In the Pull based workflow, the GitOps operator – flux in this diagram, takes care of monitoring and applying changes to the AKS cluster based on the Git repository's desired state. Developers simply make changes to the repository, and the operator ensures that the changes are propagated to the cluster. It provides a declarative and automated approach to managing and deploying applications on AKS, promoting consistency and traceability.

Overall, the push- based CI/CD workflow offers immediate feedback and deployment triggered by direct code pushes, while the pull- based ( GitOps) workflow promotes automation and declarative management of AKS clusters based on changes made to a Git repository. The choice between these workflows depends on factors such as deployment speed requirements, desired level of automation, and team preferences.



Lets talk a little bit more about the steps involved in pull based methos.

the Pull-based CI/CD workflow for AKS involves maintaining an application repository, executing build and CI pipelines to validate and publish images, triggering a CD pipeline to generate environment-specific configurations and create pull requests against the GitOps repository, and finally, leveraging the GitOps operator to apply changes to the target environments based on the desired state defined in the repository.



## CI/CD Workflow for AKS

-  Declarative Method of Managing Kubernetes
-  Push-based CI/CD Workflow
-  Pull-based or GitOps Workflow

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Ok, we have now reached the culmination of this video course on Azure Kubernetes Service (AKS). Our journey began by delving into the fundamental aspects of Azure, encompassing essential concepts such as Compute, Storage, and Networking. With a some foundation in place, we transitioned towards the creation and containerization of a basic C# web application. Throughout that module, we focused on leveraging Docker desktop and crafting Dockerfiles to facilitate containerization.

Subsequently, we navigated to the Azure portal to establish an AKS cluster to deploy our world class application onto that cluster. Harnessing the power of kubectl, we effectively executed various imperative actions, including scaling of deployments. we explored the deployment of container images within the Azure Container Registry (ACR), unearthing its potential for application enhancement and version management. By harnessing the image hosted on ACR, we facilitated a seamless application upgrade, only to effortlessly revert to a previous version when needed. Concluding this module, we bestowed upon you a overview of the Kubernetes Fleet, empowering you with a holistic understanding of this powerful toolset.

The subsequent module, entirely dedicated to AKS networking options and policies,

showcased the indispensable importance of robust network management within AKS deployments. Furthermore, we delved into security-centric topics, delving into the intricacies of Azure AD integration, defender for AKS, and Azure policy for AKS. Empowered with this knowledge, you are poised to fortify your AKS infrastructure with optimal security measures and governance practices.

As we progressed, we ventured into the realm of CI/CD patterns within AKS, looking at the art of managing AKS using declarative techniques. By adopting these practices, you can streamline the continuous integration and delivery pipelines within your AKS environment.

An indispensable component of any infrastructure, observability received some attention in module seven. We explored observability within AKS, equipping you with the necessary tools and techniques to monitor and gain valuable insights into your AKS deployments.

Finally, before we brought this course to a close, we presented you with an overview of alternative platforms that embrace the power of containers. By expanding your horizons and exploring these platforms, you can unlock a world of possibilities and further enhance your containerization endeavors.

We sincerely hope that your journey throughout this course has been as fulfilling as our commitment to its production. It has been an absolute pleasure to guide you through the intricacies of AKS, empowering you with the knowledge and skills needed to excel in your containerization journey. This is your host Pranav, signing off for now.



# KodeKloud

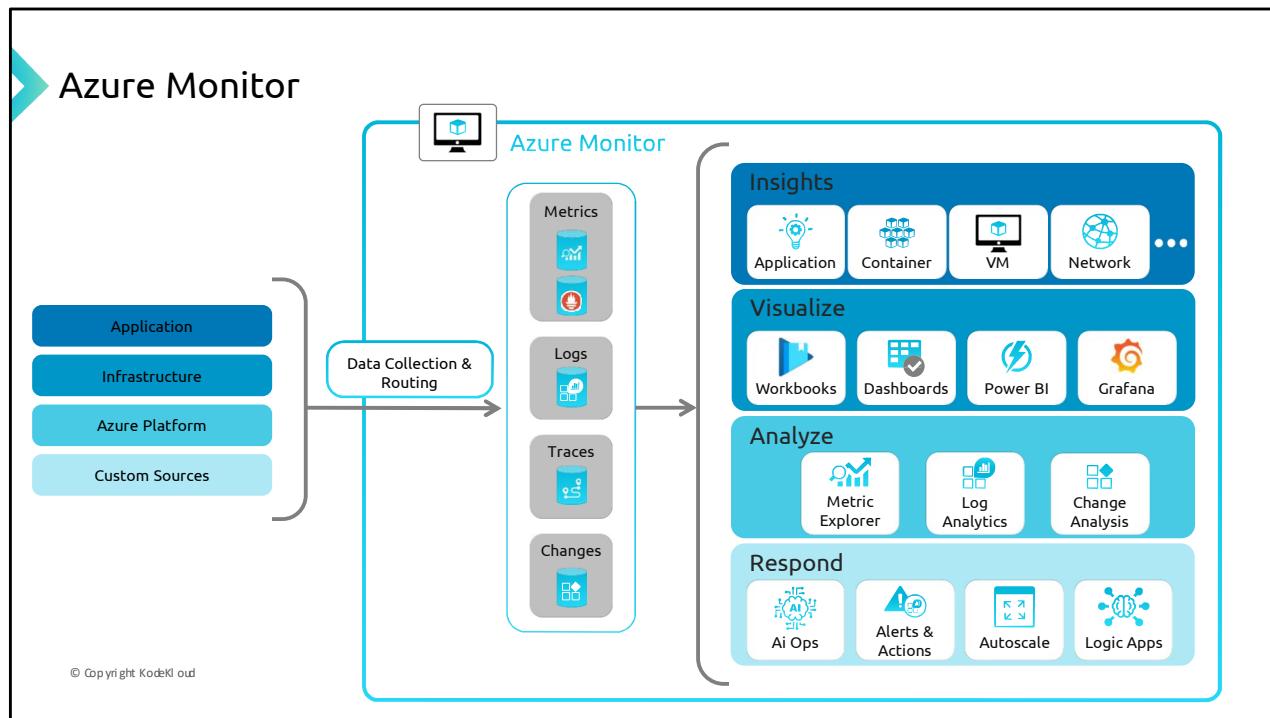
© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

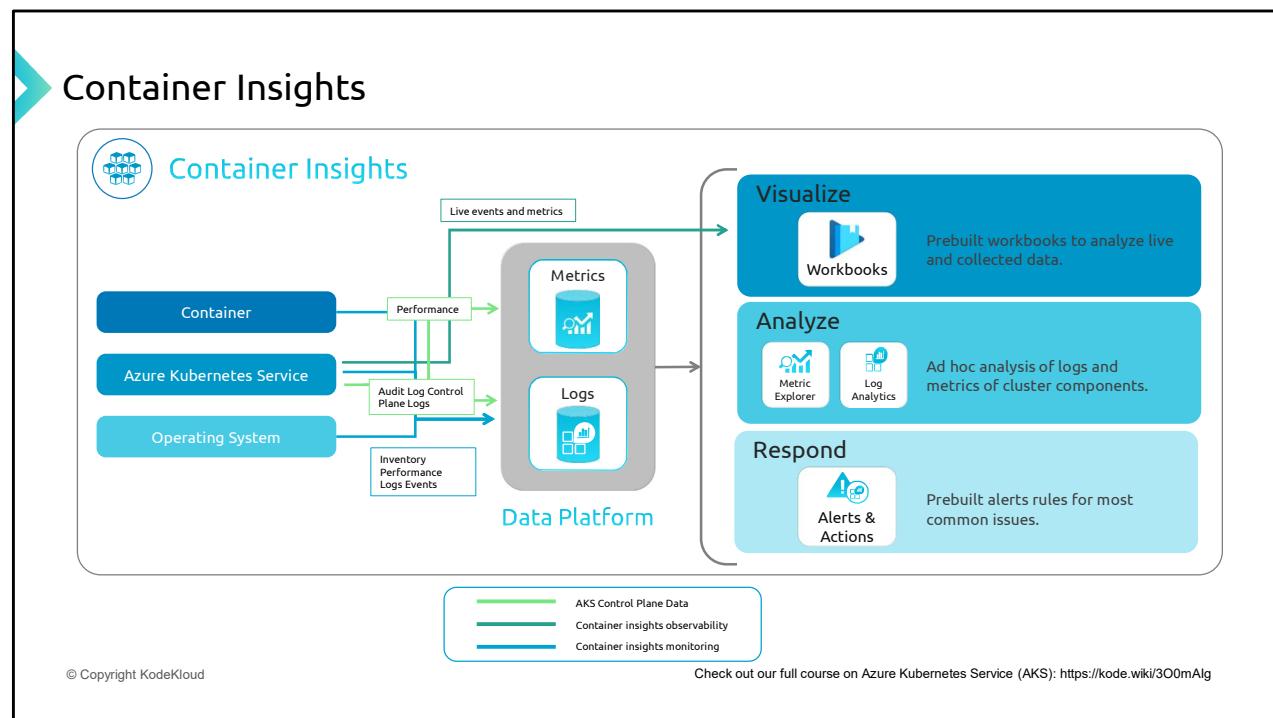
# Container Insights for AKS

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>



Azure Monitor is a monitoring solution provided by Azure that enables you to gain insights into the performance and health of your applications, infrastructure, and services running on the Azure platform. It offers a range of monitoring capabilities, including metrics, logs, alerts, and diagnostics, to help you effectively monitor and troubleshoot your Azure resources.



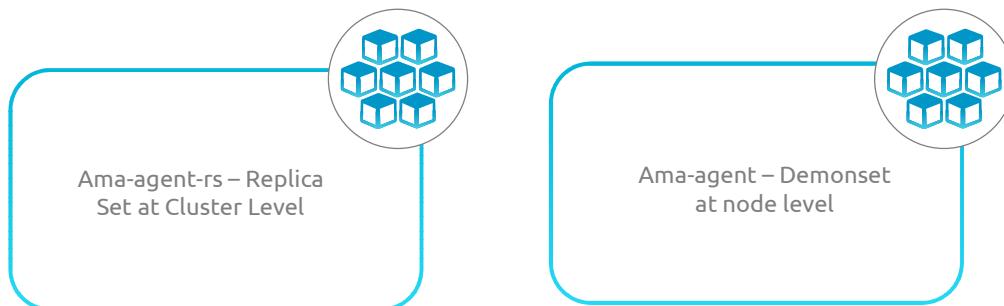
Container Insights, a feature within Azure Monitor, and is specifically designed for monitoring and managing containerized workloads deployed on Azure Kubernetes Service (AKS). It provides deep visibility into your AKS clusters and helps you gain insights into the performance, availability, and resource utilization of your containerized applications.

Azure Monitor collects and stores two primary types of data: metrics and logs.

1. **Metrics:** Metrics are numerical data that represent the performance and health of various resources. In the context of AKS and Container Insights, metrics can include information such as CPU usage, memory consumption, network traffic, and other resource-specific metrics. Azure Monitor aggregates and stores these metrics, allowing you to create charts, set up alerts, and perform analysis over time.

2. **Logs:** Logs contain structured or unstructured data that provide detailed information about the operations and activities within your AKS clusters. Container logs, application logs, and system logs are examples of log data that can be collected. Azure Monitor allows you to collect, store, and analyze these logs to gain insights into the behavior of your applications and diagnose issues.

## Agent Architecture



© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

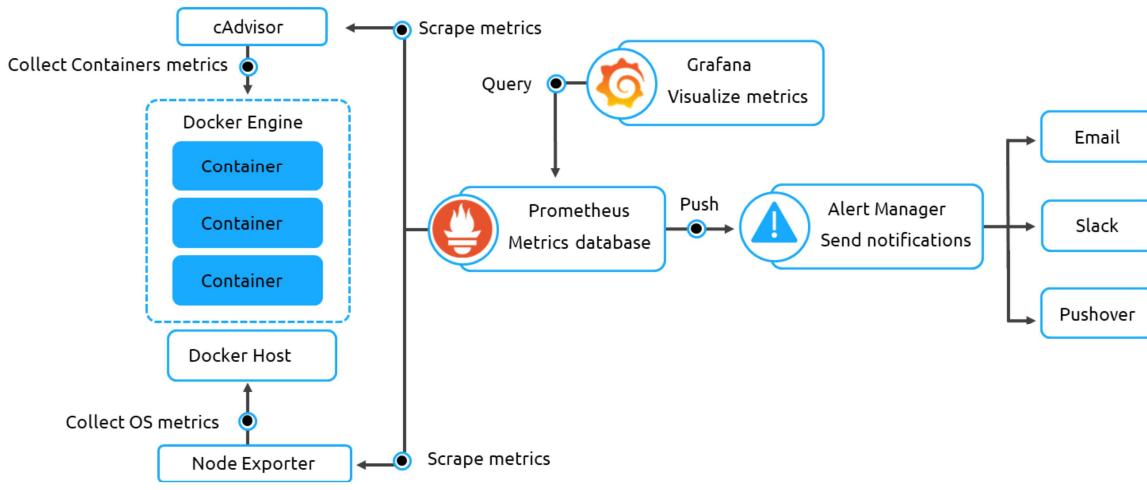
Let's take a moment to discuss the architecture of the AMA agent when enabling the monitoring add-on in AKS. There are two agents that are installed as part of this process.

The first agent is the AMA agent replica set, which is deployed on one of the nodes in the cluster. This agent's primary role is to gather metrics at the cluster level. It ensures that even if the agent running on the node where it is deployed goes down, it still collects the necessary information.

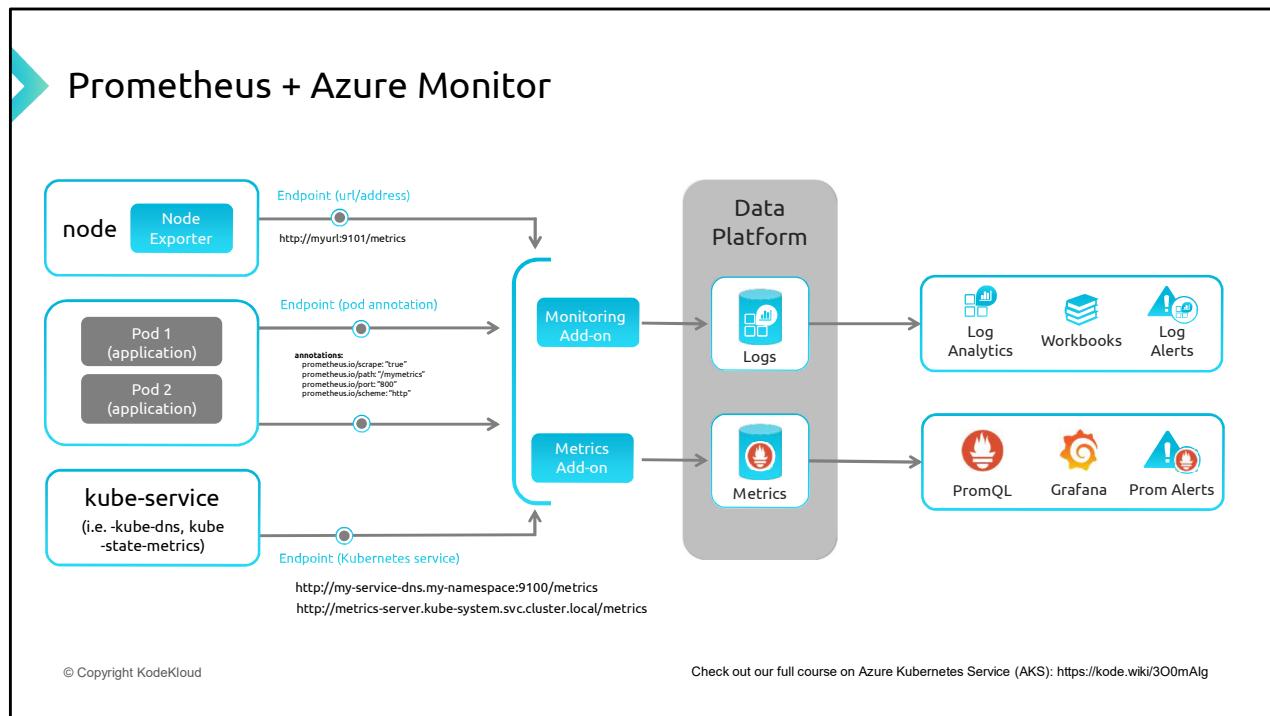
The second agent is installed on every node in the cluster. This agent is responsible for collecting metrics at the node level, as well as monitoring all the pods running on that particular node.

In summary, when enabling the monitoring add-on, two agents are installed: the AMA agent replica set for cluster-level metrics collection and fault tolerance, and the agent on each node for node-level metrics and pod monitoring.

## Prometheus + Grafana



- Prometheus is an open-source systems monitoring and alerting toolkit.
- Grafana allows you to query, visualize, alert on and understand your metrics no matter where they are stored.



Prometheus is a popular open source metric monitoring solution and is a part of the [Cloud Native Compute Foundation](#). Azure Monitor for containers provides a seamless onboarding experience to collect Prometheus metrics. Typically, to use Prometheus, you need to setup and manage a Prometheus server with a store. By integrating with Azure Monitor, a Prometheus server is not required. You just need to expose the Prometheus metrics endpoint through your exporters or pods (application), and the containerized agent for Azure Monitor for containers can scrape the metrics for you.



## AKS Observability-Summary



Container Insights



Monitoring using Prometheus and Grafana



Azure Portal to establish an AKS Cluster



AKS Networking options and Policies



We ventured into the realm of CI/CD Patterns within AKS



AKS Observability



Alternative platforms that embrace the power of containers

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Prometheus is a popular open source metric monitoring solution and is a part of the [Cloud Native Compute Foundation](#). Azure Monitor for containers provides a seamless onboarding experience to collect Prometheus metrics. Typically, to use Prometheus, you need to setup and manage a Prometheus server with a store. By integrating with Azure Monitor, a Prometheus server is not required. You just need to expose the Prometheus metrics endpoint through your exporters or pods (application), and the containerized agent for Azure Monitor for containers can scrape the metrics for you.



## Next Lesson

# Quick Summary of All Container Workload Hosts in Azure

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/3O0mAlg>

Ok, we have now reached the culmination of this video course on Azure Kubernetes Service (AKS). Our journey began by delving into the fundamental aspects of Azure, encompassing essential concepts such as Compute, Storage, and Networking. With a some foundation in place, we transitioned towards the creation and containerization of a basic C# web application. Throughout that module, we focused on leveraging Docker desktop and crafting Dockerfiles to facilitate containerization.

Subsequently, we navigated to the Azure portal to establish an AKS cluster to deploy our world class application onto that cluster. Harnessing the power of kubectl, we effectively executed various imperative actions, including scaling of deployments. we explored the deployment of container images within the Azure Container Registry (ACR), unearthing its potential for application enhancement and version management. By harnessing the image hosted on ACR, we facilitated a seamless application upgrade, only to effortlessly revert to a previous version when needed. Concluding this module, we bestowed upon you a overview of the Kubernetes Fleet, empowering you with a holistic understanding of this powerful toolset.

The subsequent module, entirely dedicated to AKS networking options and policies,

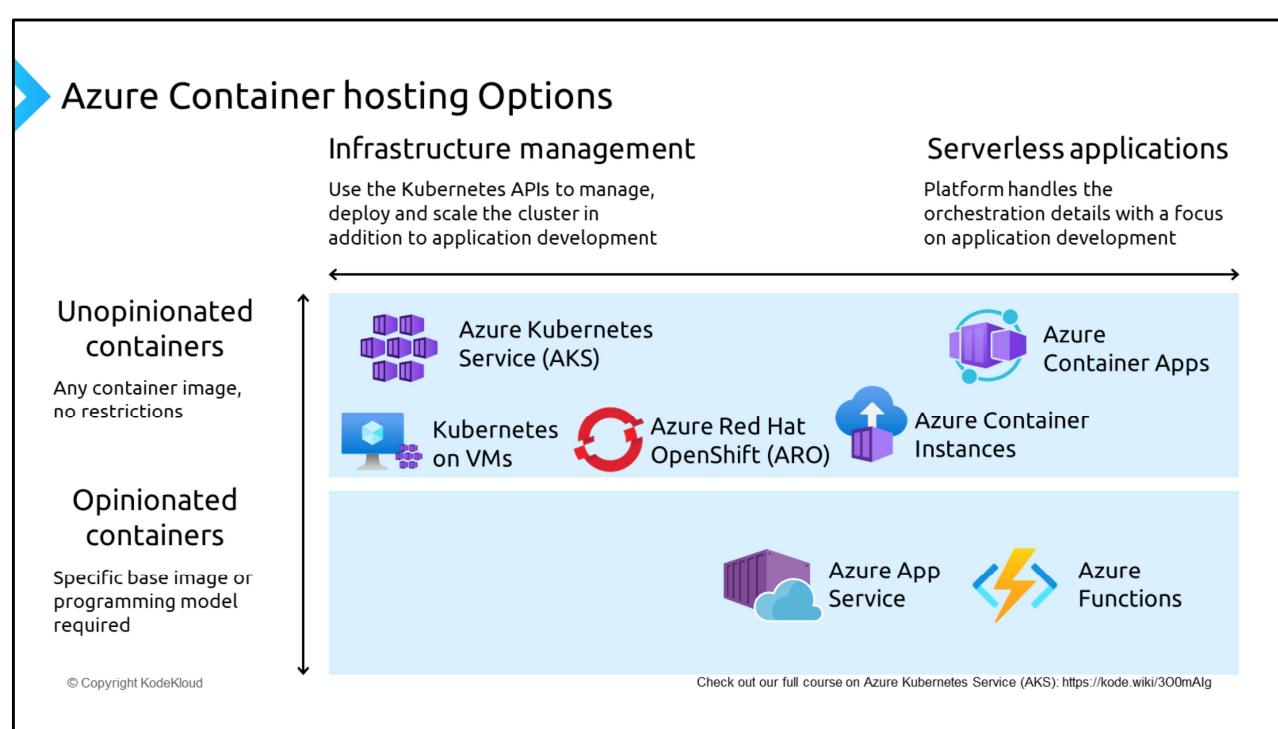
showcased the indispensable importance of robust network management within AKS deployments. Furthermore, we delved into security-centric topics, delving into the intricacies of Azure AD integration, defender for AKS, and Azure policy for AKS. Empowered with this knowledge, you are poised to fortify your AKS infrastructure with optimal security measures and governance practices.

As we progressed, we ventured into the realm of CI/CD patterns within AKS, looking at the art of managing AKS using declarative techniques. By adopting these practices, you can streamline the continuous integration and delivery pipelines within your AKS environment.

An indispensable component of any infrastructure, observability received some attention in module seven. We explored observability within AKS, equipping you with the necessary tools and techniques to monitor and gain valuable insights into your AKS deployments.

Finally, before we brought this course to a close, we presented you with an overview of alternative platforms that embrace the power of containers. By expanding your horizons and exploring these platforms, you can unlock a world of possibilities and further enhance your containerization endeavors.

We sincerely hope that your journey throughout this course has been as fulfilling as our commitment to its production. It has been an absolute pleasure to guide you through the intricacies of AKS, empowering you with the knowledge and skills needed to excel in your containerization journey. This is your host Pranav, signing off for now.



We are almost at the end of this course and AKS and I am glad you have made it this far. As you start to use AKS in your own environments, one common question you will get asked is if AKS is the only way to host containers on Azure and if it's the right platform to host your containers.

On the screen, you can see many native Azure services that can host containers. Depending on your use case, you can choose one of the option. For example, if you need full control of container orchestration and have a custom image to run, you will need to pick a service from top left segment i.e AKS, Kubernetes on a VM or Redhat open shift.

Lets quickly talk about various available options:

**AKS and Azure Red Hat OpenShift** provide access to a full Kubernetes or Red Hat OpenShift cluster, and you are in full control of all aspects of the cluster including **ingress, networking, maintenance, and scale**. If you want full control, you can also go old school and deploy k8s on VMs.

Azure Container Apps operates at the level above the cluster, where the developer only has to worry about the app itself, and the cluster management is handled by the service even though it uses k8s.. You don't have access to the apis, no Custom Resource Definition etc. It is also worth noting that only Linux based containers are supported on ACA.

**Azure Container Instances** provides a way to launch containers with hypervisor isolation. It provides a container as a unit in Azure. It's a low-level compute option, akin to provisioning a VM with a user experience designed to match a local container runtime. It's used to power experiences in a more generic fashion including Virtual Nodes as part of AKS. It is worth noting that (ACI) is primarily designed for running individual containers or microservices rather than supporting full-fledged container orchestrations like Kubernetes.

**Azure App Service** provides a purpose-built experience for hosting web applications and web APIs, including the ability to publish code directly without managing a container, and features like deployment slots and “test in production”.

Azure Container Apps is optimized for microservices that need to communicate to each other, which is difficult to do with Azure App Service.

Finally, **Azure Functions** is about executing event-driven serverless code with an end-to-end development experience. Code deployed to Azure Functions must either use the Functions SDK or use the Functions base container image and adhere to the Functions predefined programming model.

I encourage you to read through some of these options and decide which one works for your application better.



## Azure Kubernetes Services (AKS) - Summary

- ✓ Fundamental Aspects of Azure
- ✓ Creation and Containerization of a basic C# Web Application
- ✓ Azure Portal to establish an AKS Cluster
- ✓ AKS Networking options and Policies
- ✓ We ventured into the realm of CI/CD Patterns within AKS
- ✓ AKS Observability
- ✓ Alternative platforms that embrace the power of containers

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>

Ok, we have now reached the culmination of this video course on Azure Kubernetes Service (AKS). Our journey began by delving into the fundamental aspects of Azure, encompassing essential concepts such as Compute, Storage, and Networking. With a some foundation in place, we transitioned towards the creation and containerization of a basic C# web application. Throughout that module, we focused on leveraging Docker desktop and crafting Dockerfiles to facilitate containerization.

Subsequently, we navigated to the Azure portal to establish an AKS cluster to deploy our world class application onto that cluster. Harnessing the power of kubectl, we effectively executed various imperative actions, including scaling of deployments. we explored the deployment of container images within the Azure Container Registry (ACR), unearthing its potential for application enhancement and version management. By harnessing the image hosted on ACR, we facilitated a seamless application upgrade, only to effortlessly revert to a previous version when needed. Concluding this module, we bestowed upon you a overview of the Kubernetes Fleet, empowering you with a holistic understanding of this powerful toolset.

The subsequent module, entirely dedicated to AKS networking options and policies,

showcased the indispensable importance of robust network management within AKS deployments. Furthermore, we delved into security-centric topics, delving into the intricacies of Azure AD integration, defender for AKS, and Azure policy for AKS. Empowered with this knowledge, you are poised to fortify your AKS infrastructure with optimal security measures and governance practices.

As we progressed, we ventured into the realm of CI/CD patterns within AKS, looking at the art of managing AKS using declarative techniques. By adopting these practices, you can streamline the continuous integration and delivery pipelines within your AKS environment.

An indispensable component of any infrastructure, observability received some attention in module seven. We explored observability within AKS, equipping you with the necessary tools and techniques to monitor and gain valuable insights into your AKS deployments.

Finally, before we brought this course to a close, we presented you with an overview of alternative platforms that embrace the power of containers. By expanding your horizons and exploring these platforms, you can unlock a world of possibilities and further enhance your containerization endeavors.

We sincerely hope that your journey throughout this course has been as fulfilling as our commitment to its production. It has been an absolute pleasure to guide you through the intricacies of AKS, empowering you with the knowledge and skills needed to excel in your containerization journey. This is your host Pranav, signing off for now.



# KodeKloud

© Copyright KodeKloud

Check out our full course on Azure Kubernetes Service (AKS): <https://kode.wiki/300mAlg>